


BIG DATA PROCESSING

EGCI 466 – No SQL Databases

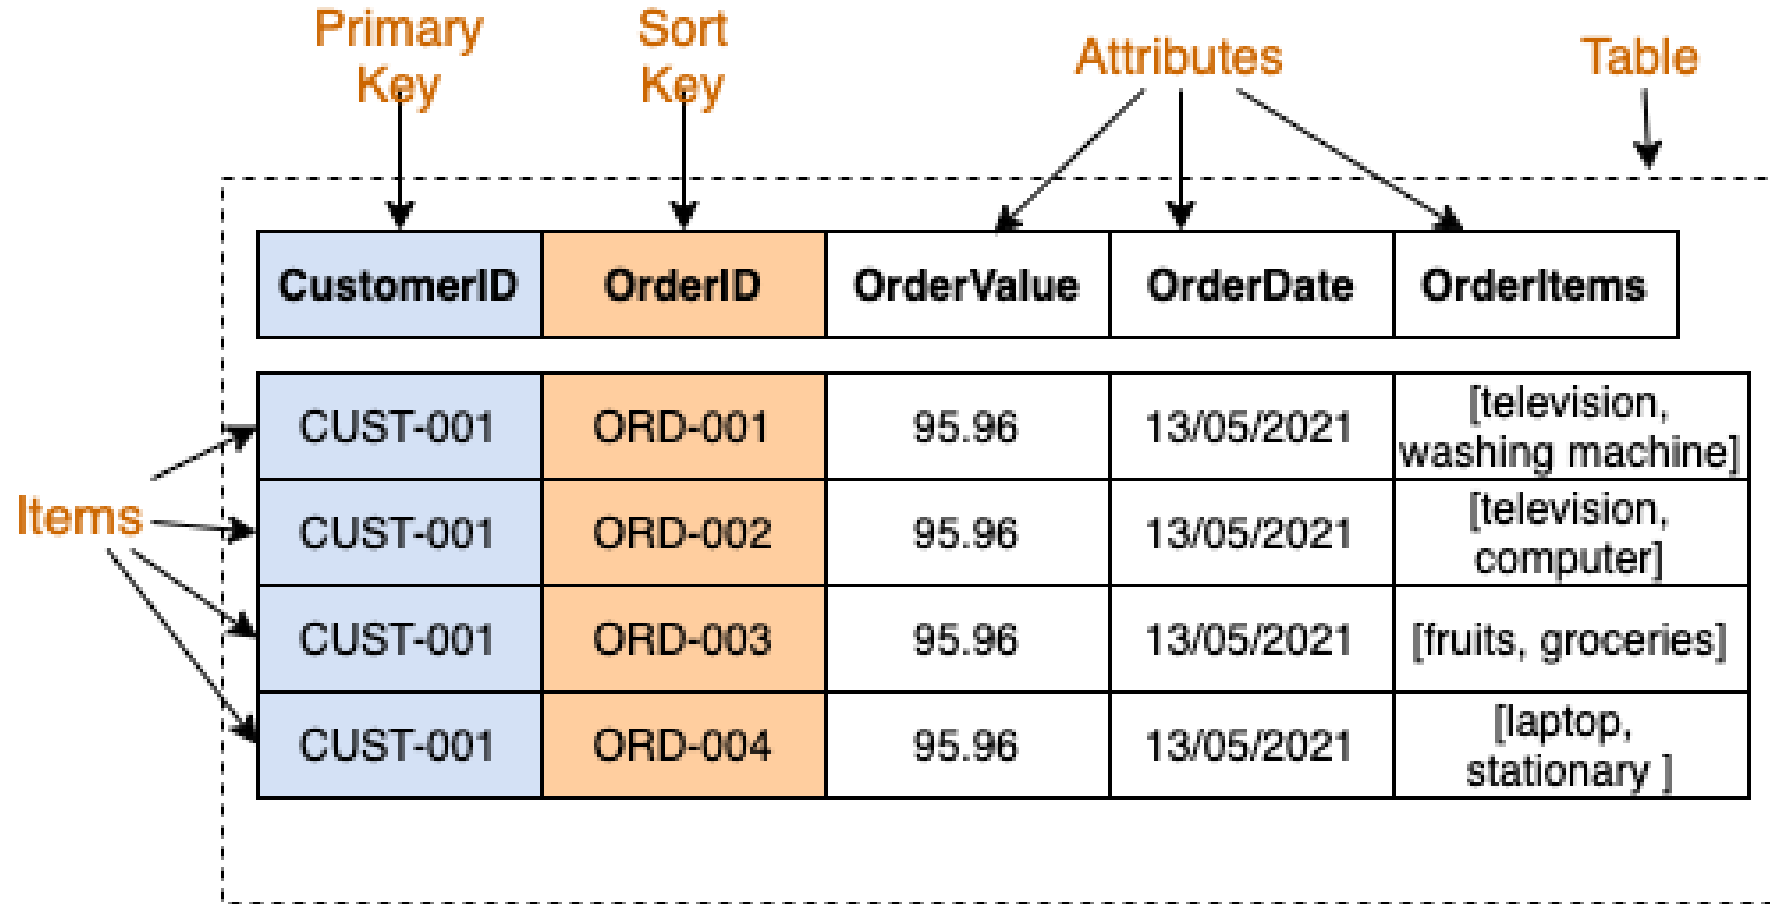
Key-value

Key-value

- Least Complex
- Each keys has it's own value
- Ideal for simple Create Read Update and Delete(CRUD)
- Scale well (by key)



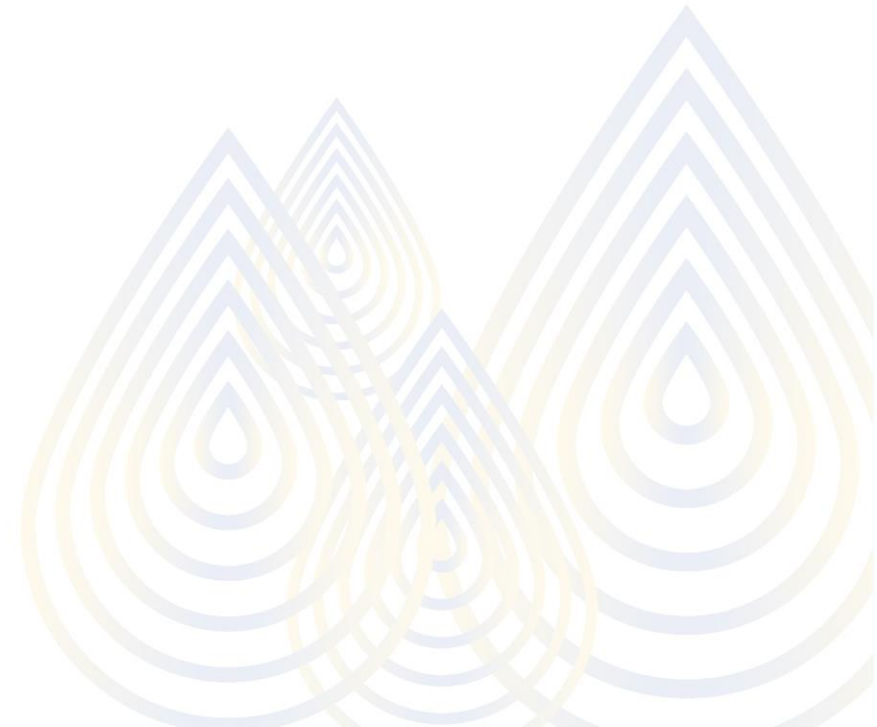
Key-value



Primary Key	Sort Key	Attributes			Table
CustomerID	OrderID	OrderValue	OrderDate	OrderItems	
CUST-001	ORD-001	95.96	13/05/2021	[television, washing machine]	
CUST-001	ORD-002	95.96	13/05/2021	[television, computer]	
CUST-001	ORD-003	95.96	13/05/2021	[fruits, groceries]	
CUST-001	ORD-004	95.96	13/05/2021	[laptop, stationary]	

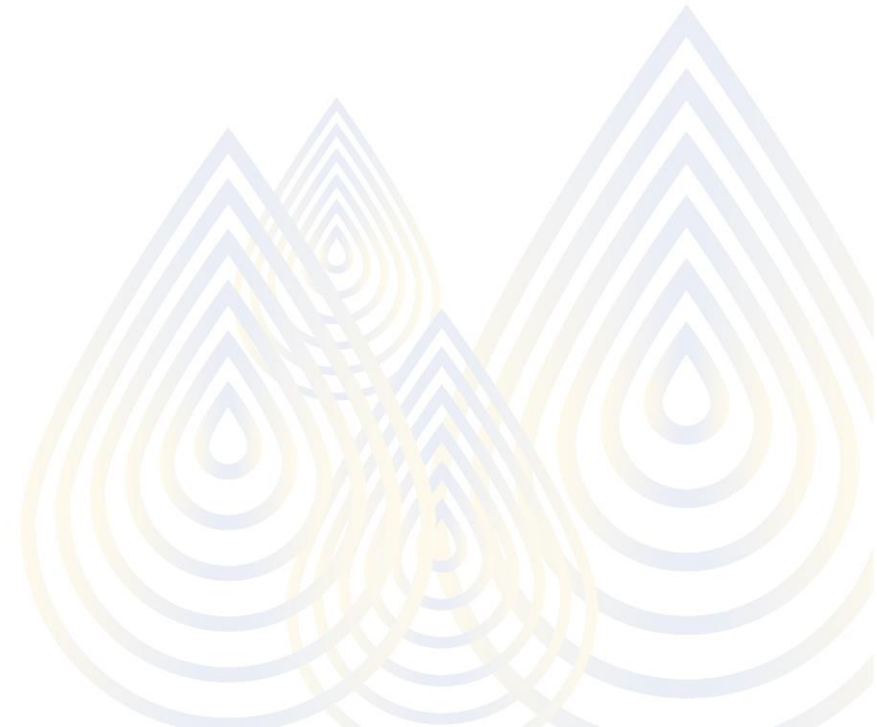
Key-value

- Quick CRUD (Create-Read-Update-Delete) operation
 - Storing/retrieving for web app
- in-app user profiles
- Shopping cart only



Key-value: Not suitable for

- Interconnected data
 - Social Media
 - Recommendation system
- Very high consistency
- Require a lot of queries




EXAMPLE

- [dbm](#) : unix system
- Sdbm
- GNU dbm
- [Berkeley DB](#)
- Amazon DynamoDB
 - Used By: Samsung Cloud, Zoom, Disney+, Nike
- Aerospike - An open-source ultra-low latency, reliability, and ability to handle large load.
- Redis - A multi-purpose database that also acts as memory cache and message broker.
- Riak - Made for developing apps, it works well with other databases and apps.

Document Based

Document-based

- Key-Values → value visible for query
- Each piece of data is considered a document
 - Typically JSON or XML format
- Each document offers a flexible schema
- Content of document can be indexed and queried





Mahidol University
Wisdom of the Land

JSON

```
{  
  "status": 200,  
  "photos":  
  [  
    {  
      "typeName": "Facebook",  
      "type": "facebook",  
      "typeId": "facebook",  
      "url": "http://graph.facebook.com/amoghnatu/picture?type=large",  
      "isPrimary": true  
    }  
  ],  
  "contactInfo": {  
    "familyName": "Natu",  
    "fullName": "Amogh Natu",  
    "givenName": "Amogh"  
  },  
  "demographics": {  
    "gender": "male"  
  },  
  "socialProfiles":  
  [  
    {  
      "id": "1839143973",  
      "typeName": "Facebook",  
      "username": "amoghnatu",  
      "type": "facebook",  
      "typeId": "facebook",  
      "url": "http://www.facebook.com/amoghnatu"  
    }  
  ]  
}
```

Key-value pair

Tuple


Square
brackets
indicate arrays

XML

```
- <experiments version="1.2" revision="100915" total="58" total-samples="4011" total-assays="3847">
  - <experiment>
    <releasedate>2007-11-22</releasedate>
    <species>Mus musculus</species>
    - <miamescores>
      <reportersequencescore>1</reportersequencescore>
      <factorvaluescore>1</factorvaluescore>
      <measuredbioassaydatascore>0</measuredbioassaydatascore>
      <protocolscore>0</protocolscore>
      <derivedbioassaydatascore>1</derivedbioassaydatascore>
      <overallscore>3</overallscore>
    </miamescores>
    <assays>18</assays>
    <samples>18</samples>
    <rawdatafiles>0</rawdatafiles>
    <fgemdatafiles>18</fgemdatafiles>
    - <sampleattribute>
      <category>CellType</category>
      <value>primary chondrocyte</value>
      <value>primary dermal fibroblast</value>
      <value>primary osteoblast</value>
    </sampleattribute>
    - <sampleattribute>
      <category>Organism</category>
      <value>Mus musculus</value>
    </sampleattribute>
    - <experimentalfactor>
      <name>CellType</name>
      <value>primary chondrocyte</value>
      <value>primary dermal fibroblast</value>
```


Document-based Query Property

```
<document>
  <report>
    <author>Video database</author>
    <date>June 12, 2000</date>
  </report>
  <paper>
    <title>XML query data model</title>
    <author>Don Robie</author>
    <source>W3C, June 2000</source>
  </paper>
</document>
```



Collection

- A group of stored document
 - All student records → collection
 - Staff records in Employee section



EXAMPLE: MONGO DB

- A document and a NoSQL database
- Document

Key:value

```
{  
  "firstName": "John",  
  "lastName": "Doe",  
  "email": "john.doe@email.com",  
  "studentId": 20217484  
}
```


Mongo DB

- [Get mongodb account → https://cloud.mongodb.com/](https://cloud.mongodb.com/)
- https://github.com/AjMing/BigData/blob/main/MongoDB_PyMongo_Example.ipynb

More reference:

https://www.cloudskillsboost.google/course_templates/731

(some are depreciated)



Document in Detail

```
{  
  "firstName": "Jaidee",  
  "lastName": "Deejung",  
  "email": "Jaidee@gmail.com",  
  "studentId": 6522111  
}
```

```
{  
  "firstName": "Jaidee",  
  "lastName": "Deejung",  
  "email": "Jaidee@gmail.com",  
  "studentId": 6522111  
  "address": {  
    "province": "Nakhon Pathom",  
    "district": "Salaya"  
  },  
  "Interests": ["football", "reading", "café hopping"]  
}
```

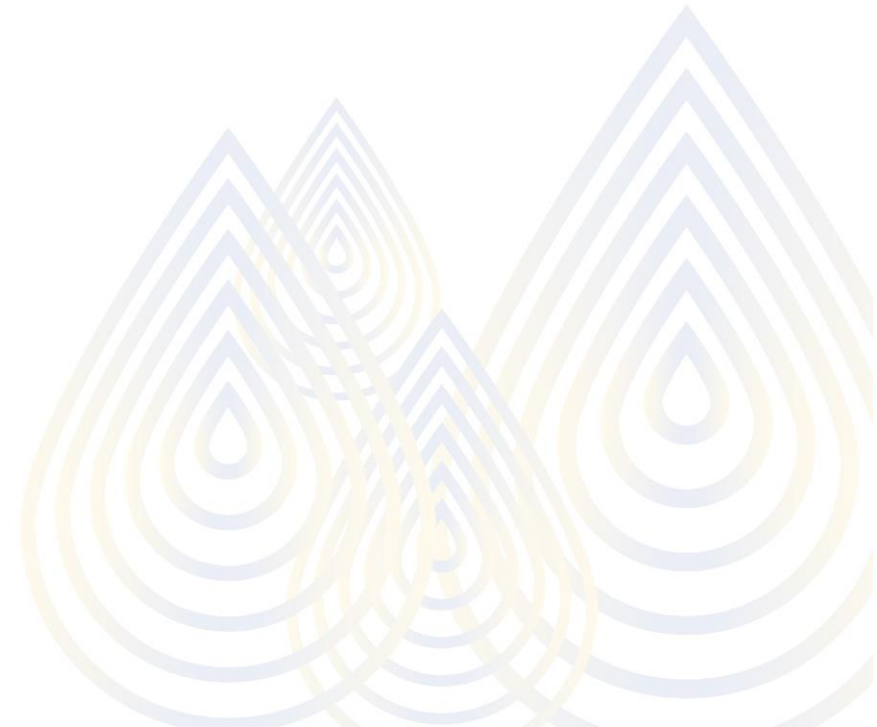

Flexible Schema

```
{  
  "street": "10 High St",  
  "city": "London",  
  "postcode": "W1 1SU"  
}
```

```
{  
  "street": "8717 West St",  
  "city": "New York",  
  "zip": "10940"  
}
```

No Schema required

- In RDBMS
 - Create table
 - Insert into Table
- in MongoDB
 - Insert directly
 - easily add field in the database



Combine different type of data

```
{  
  "symbol": "IBM",  
  "open": 235.9,  
  "high": 237.47,  
  "low": 233.17  
}
```

```
{  
  "stockName": "IBM",  
  "pricing": {  
    "o": 235.9,  
    "h": 237.47,  
    "l": 233.17  
  }  
  date: "2021-03-01T00:00:00+0000"  
}
```

Query and Analytics

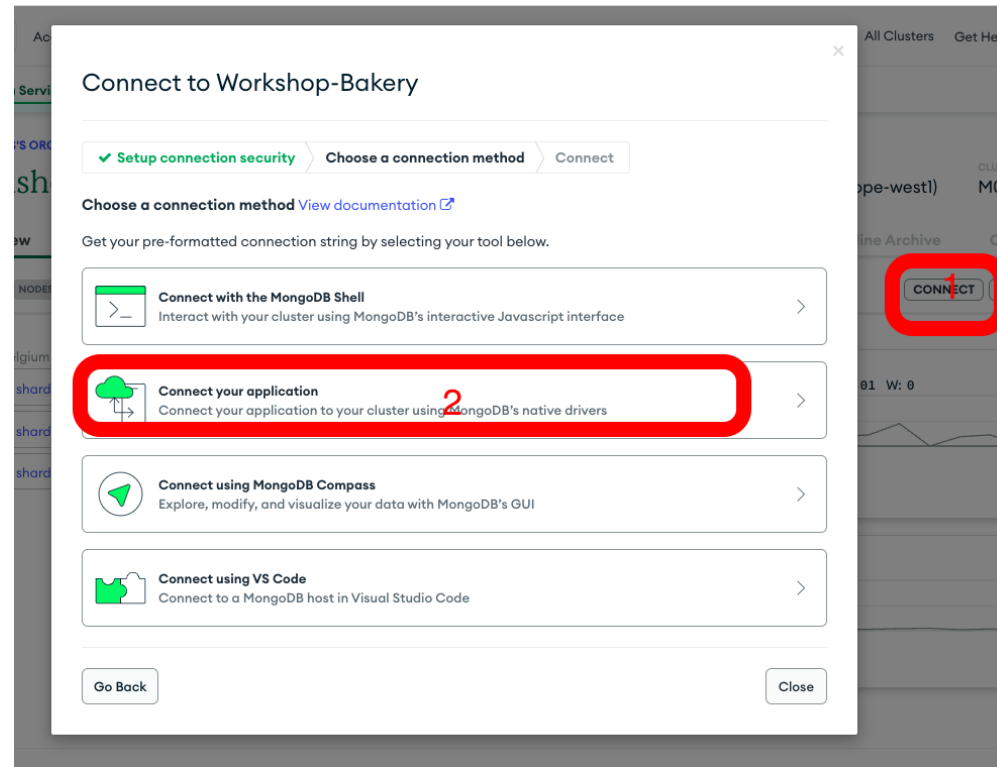
- MongoDB querying system MQL
 - wide range of operations

```
db.orders.aggregate( [  
  { $match: { status: "urgent" } },  
  { $group: { _id: "$productName", sumQuantity: { $sum: "$quantity" } } }  
] )
```



EXAMPLE of pymongo

- Setting up



- white list (remove afterward)

allow all ip address



SECURITY

Get password

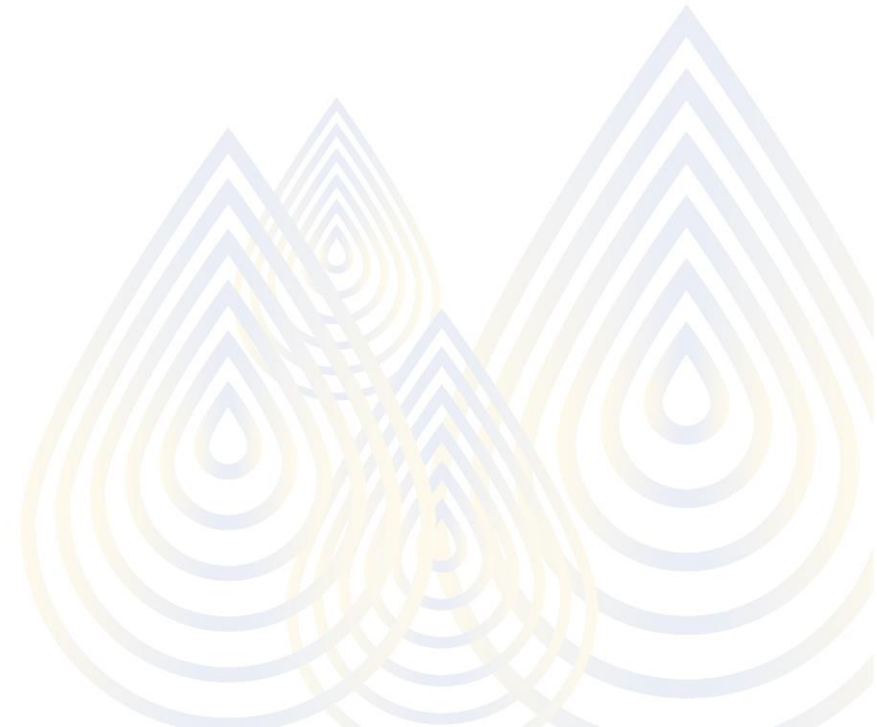
Database Access

Network Access

Advanced

EXAMPLE of pymongo

- [Pymongo in colab](#)
- [MongoDB with Big Query](#)




CRUD operation

- Create

- use sample
- `db.user.insertOne({"firstName": "John",
"lastName": "Doe",
"email": "john.doe@mahidol.edu",
"id": 6082111 }
)`
- `db.user.insertMany(userlist)`

- Read


- `user.findOne({"email": john.doe@mahidol.edu})`
- `user.find({"lastName": Doe})`
- `user.countDocuments({"lastName": Doe})`



CRUD operation


- Replace

- `user=db.user.findOne({"lastName":"Doe"})`
- `db.user["onlineOnly"]=true`
- `db.user["email"]="john.doe@mahidol.ac.th"`
- `db.user.replaceOne({"lastName":"Doe"},student)`



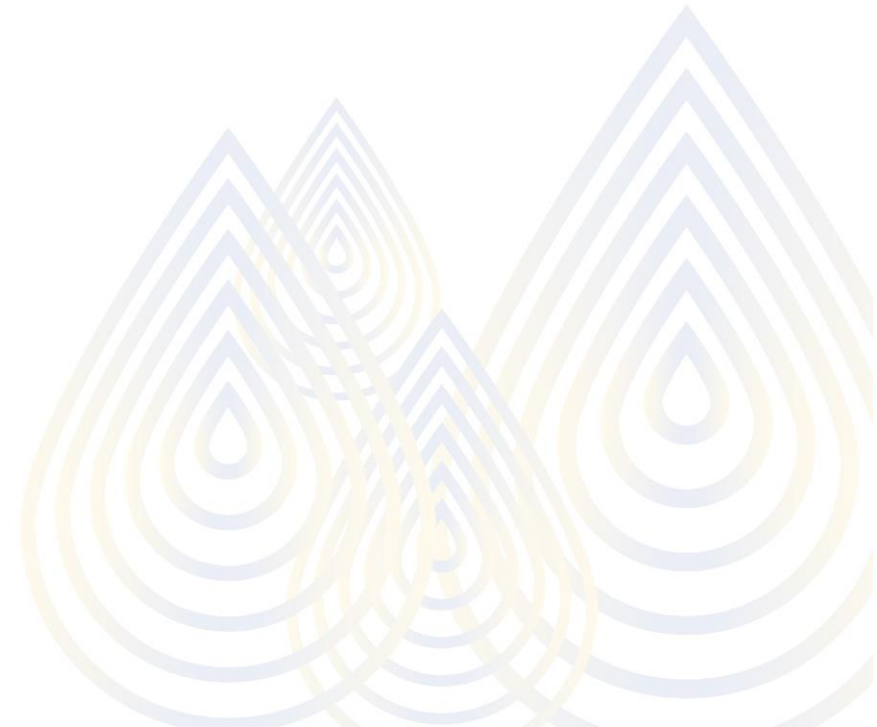
Update

- change= {"\$set": {"onlineOnly":true, email: johnd@campus.edu }}
- db.user.updateOne({"lastName":"Doe"},changes)
- db.user.updateMany({ }, {"\$set": {"onlineOnly":true}})



Delete

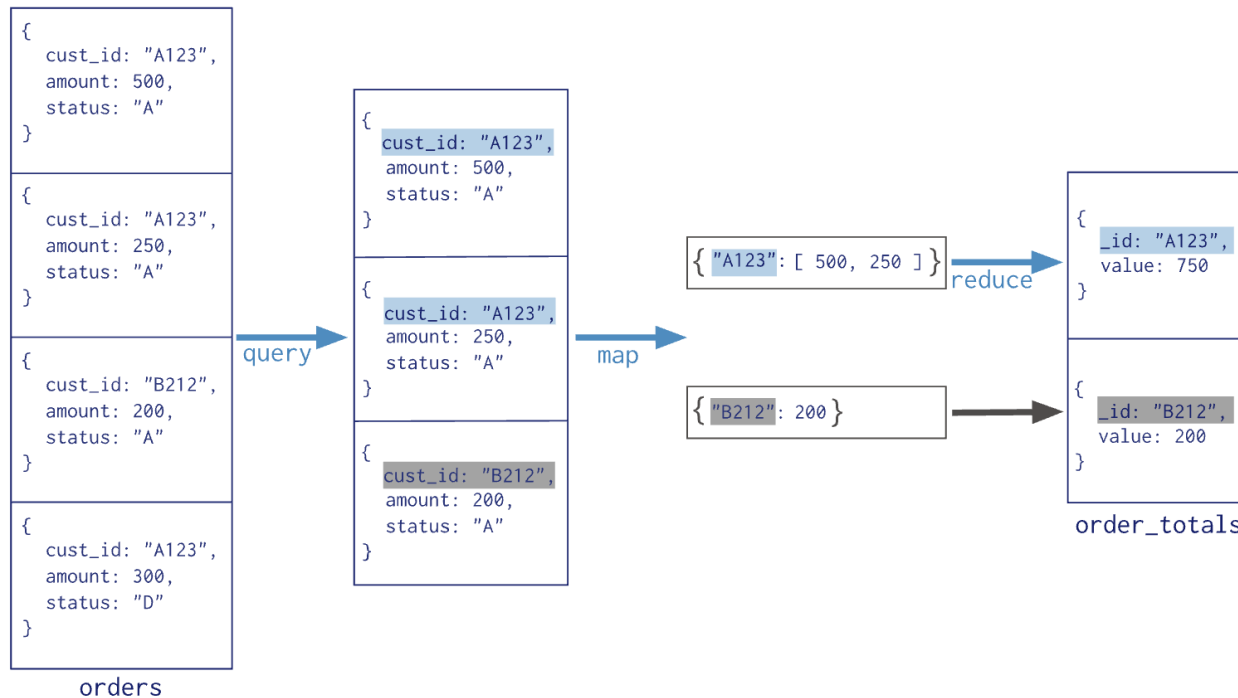
- `db.user.deleteOne({"lastName":"Doe"})`
- `db.user.deleteMany({"lastName":"Doe"})`



Map-reduce(deprecated)

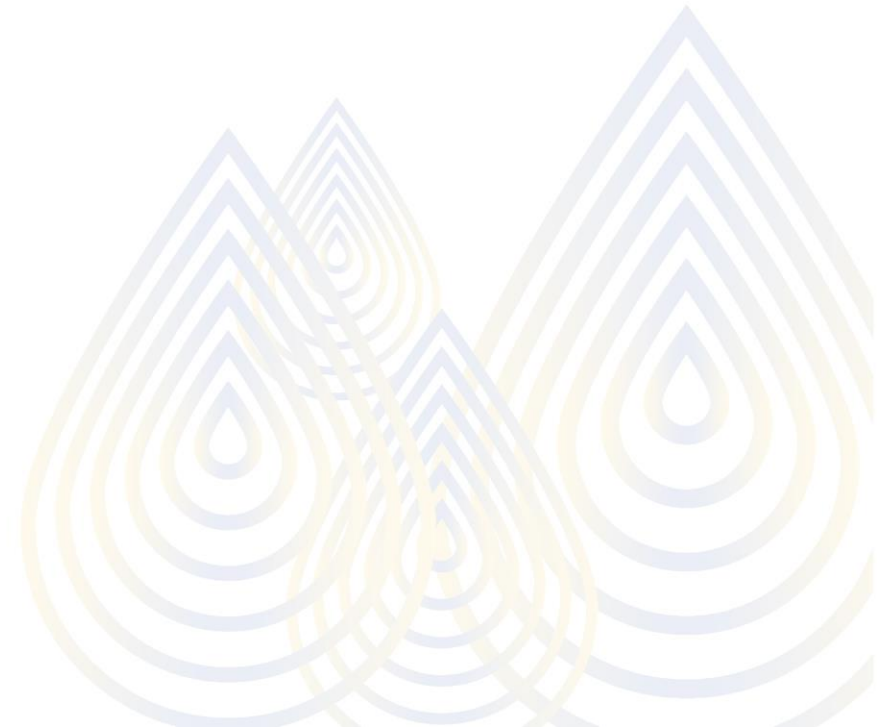
```

Collection
↓
db.orders.mapReduce(
  map   → function() { emit( this.cust_id, this.amount ); },
  reduce → function(key, values) { return Array.sum( values ) },
  {
    query → { status: "A" },
    output → "order_totals"
  }
)
  
```



High Availability

- It highly available
- Highly redundancy
- Typically 3 node replica sets, 1 primary



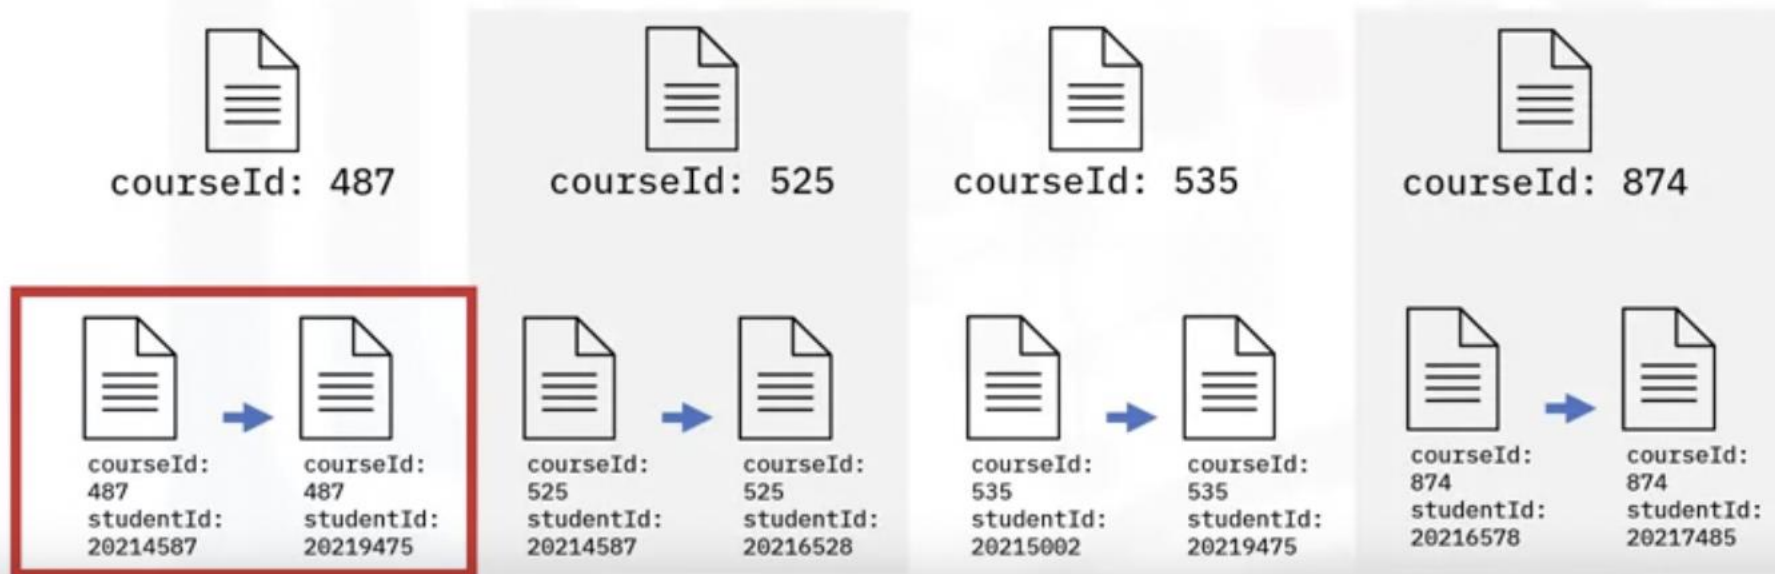
Indexes

- Quickly locate data without looking for it everywhere
- Create for frequent queries
 - `db.courseEnrollment.createIndex({"courseId":1})`
- For sorting
 - `db.courseEnrollment.createIndex({"courseId":1 , "studentId":1})`
- Index. → location
 - using balance tree



Tree index

```
db.courseEnrollment.createIndex({"courseId": 1, "studentId":1})
```




Aggregation Framework

```
> db.courseResults.aggregate([  
  { $match: { "year": 2020 } },  
  { $group: { "_id": "$courseId", "avgScore": { $avg: "$score" } } }  
])
```

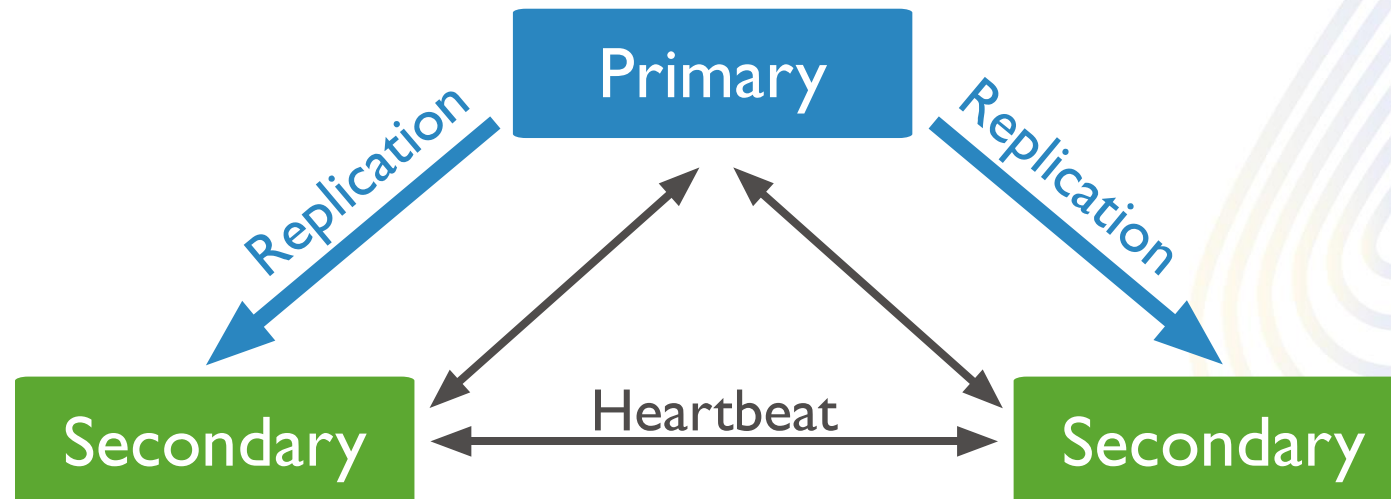
Aggregation stages

- \$project
 - change the project
- \$Sort
- \$ Count
- \$merge →into



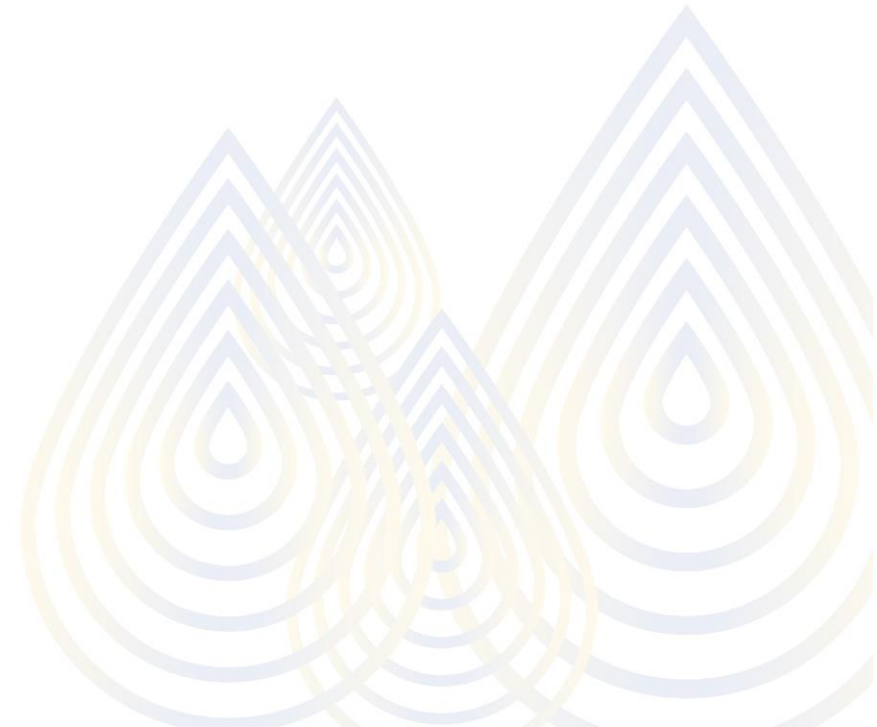
Replication & Sharding

- A MongoDB cluster is made of data bearing nodes
- All nodes contain the same data
- Write is done on Primary node → Replicate to Secondary nodes

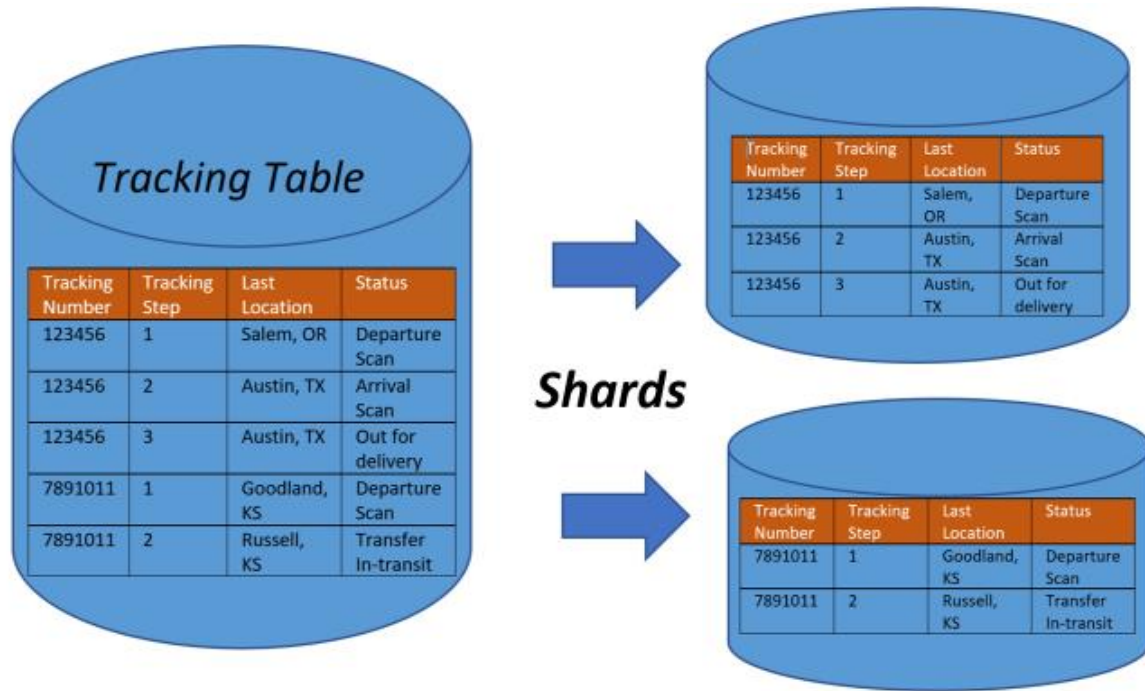


Benefits of Replication

- Redundancy
- High availability
- Fault tolerance
- Rely more on back up on stores



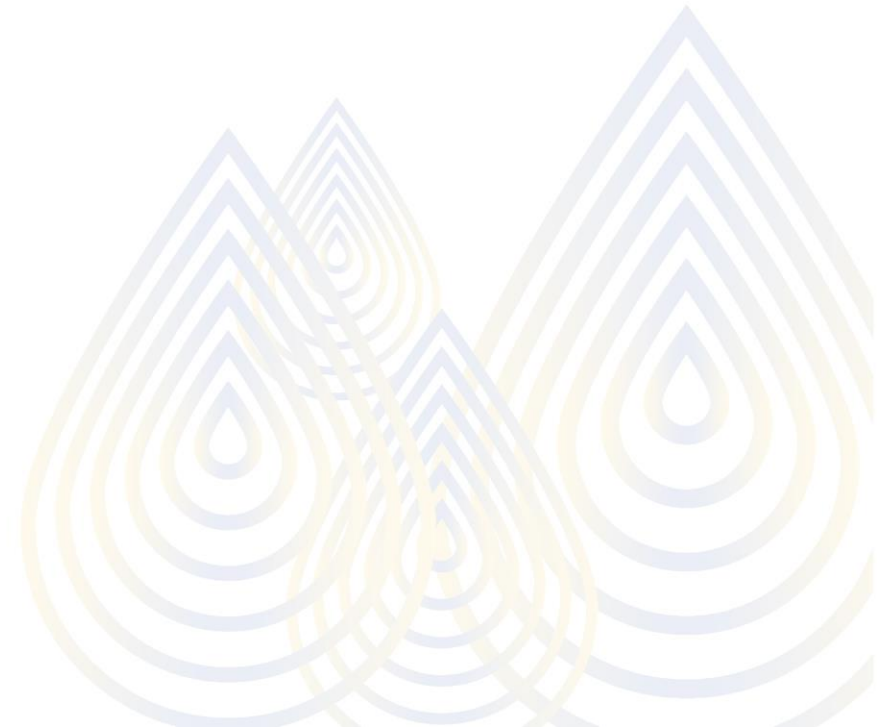
Sharding (Horizontal Scaling)



- Increase your throughput by directing to only to relevant shards.
- Store more data that couldn't fit on a single node.
- Split data across shards based on regions

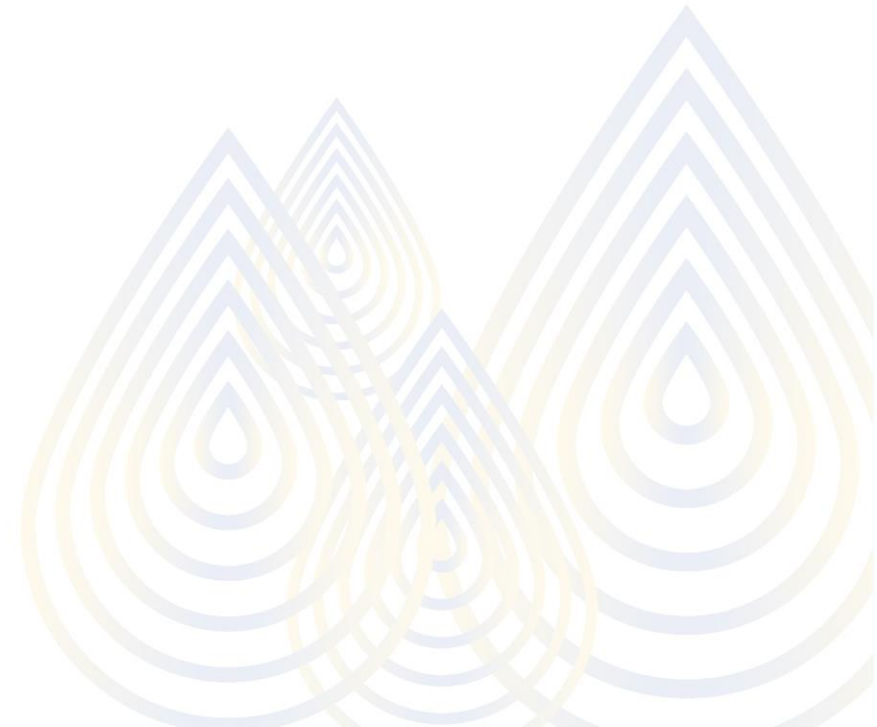
Examples: IoT devices

- Different types of data
 - different sensors: temp sensor, wind sensor
 - Vast amount of data
 - Large scale



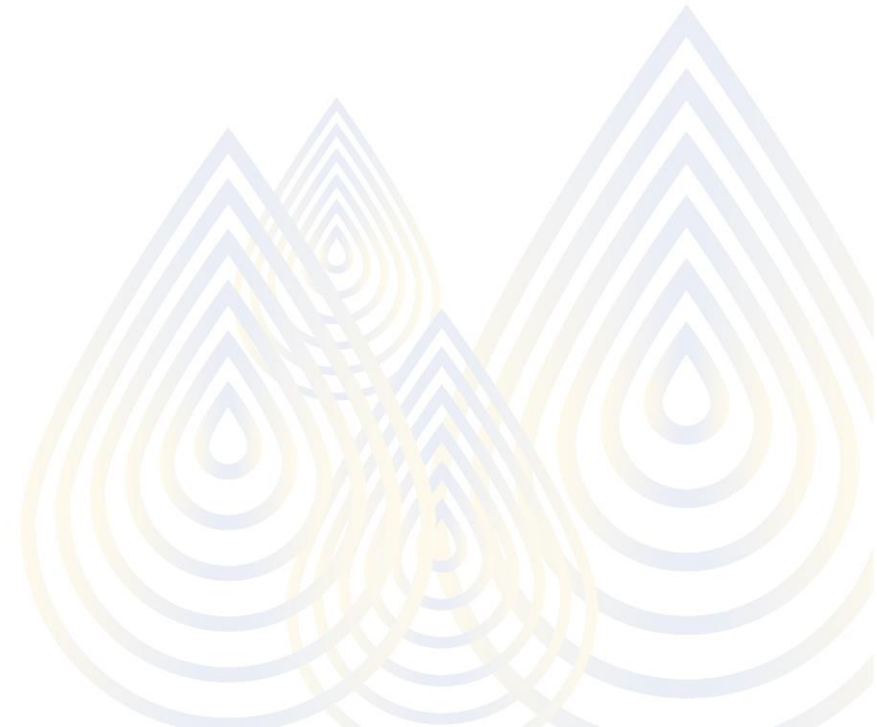
Examples: E-commerce

- Different attributes
 - Clothes: colour , size
 - Shoe: Colour, size
 - Phone: storage, network, color,brand
 - Book: Publisher, Author, pages, year, title
- Optimise for read
 - they can be stored together




Real-time Analytics

- Quick response to changes
- Simplified ETL
- **Real time**, along with operational data



Document-based Unsuitable

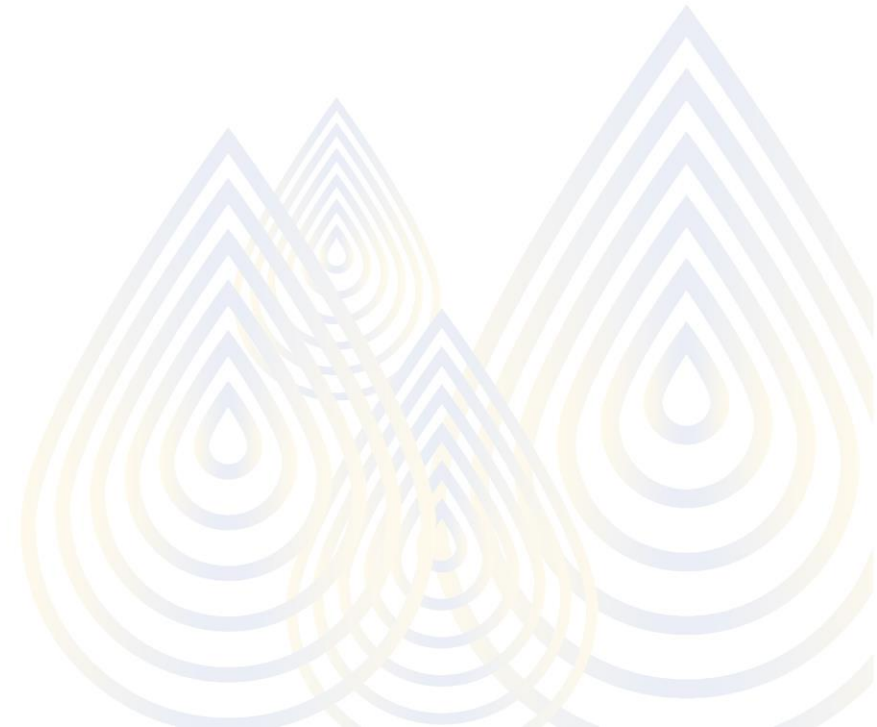
- Cannot handle transaction with multiple documents
- Not for aggregated-oriented model, handling complex join operations
- Not suited for high consistency is required




Column-Based

Column-Based Databases

- Columnar or Wide-Column databases
- Spawned from Google's 'Bigtable'
- Store data in columns
- Each row can has different column



Column-Based Databases



Columns have three values

- Key or columns name
- Timestamp
- Value

Example:

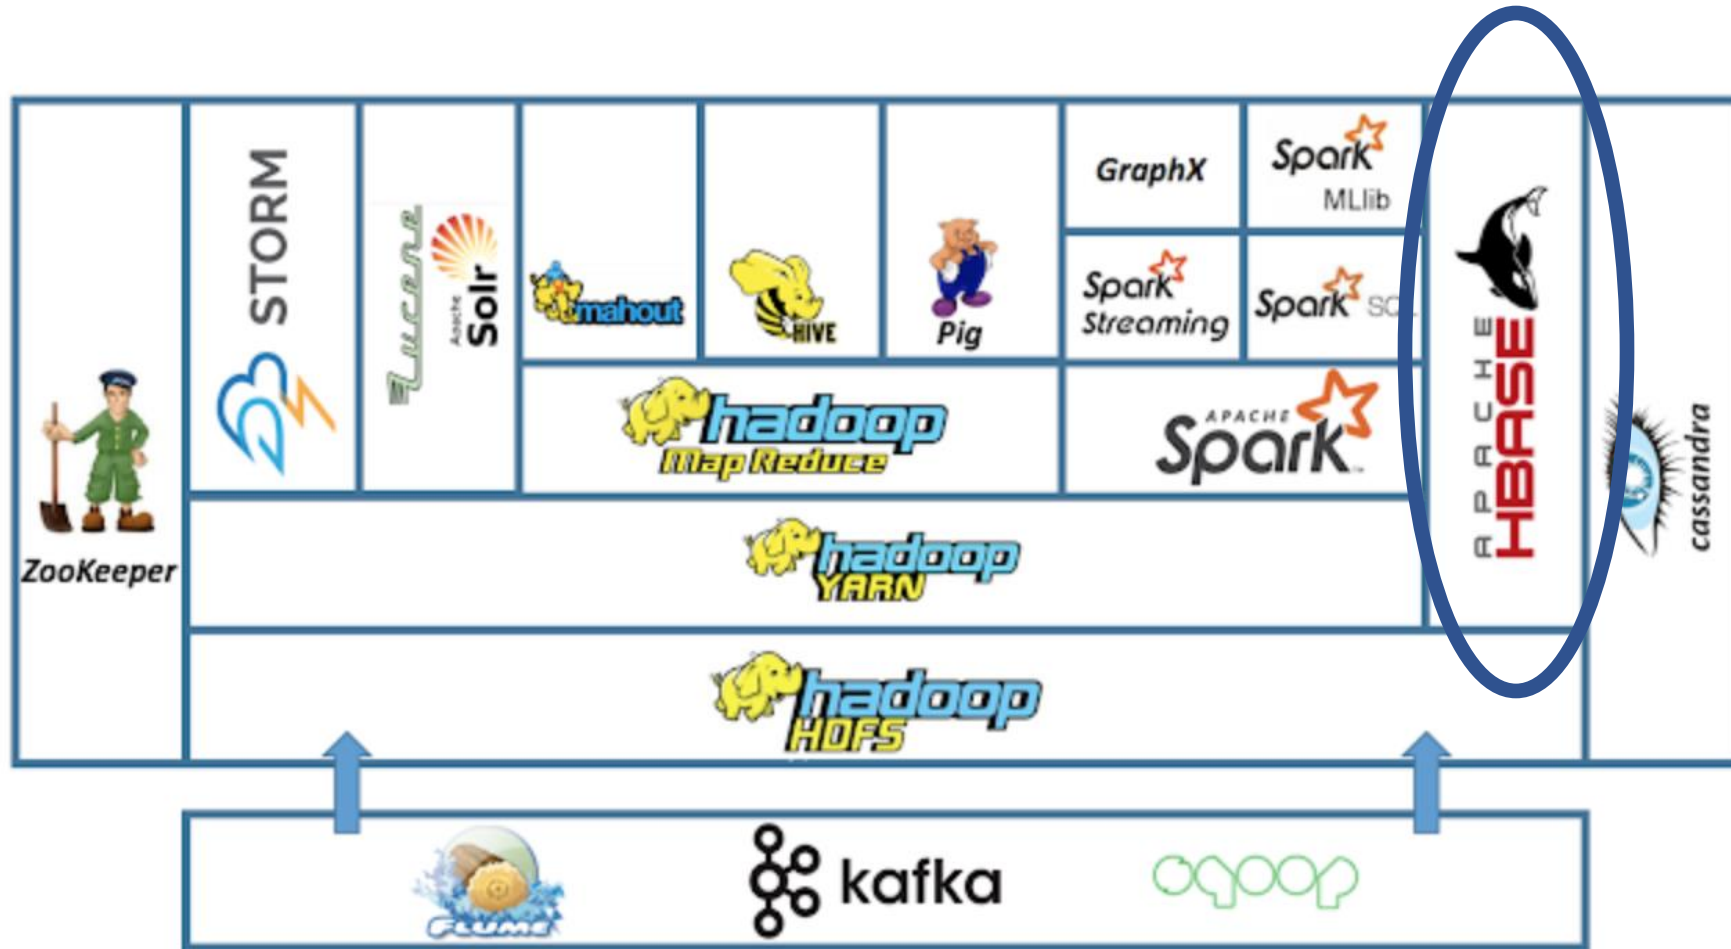
Column Family
User Profile

UserProfile

Bob	emailAddress	gender	age
	bob@example.com	male	35
	1465676582	1465676582	1465676582
Britney	emailAddress	gender	
	brit@example.com	female	
	1465676432	1465676432	
Tori	emailAddress	country	hairColor
	tori@example.com	Sweden	Blue
	1435636158	1435636158	1465633654



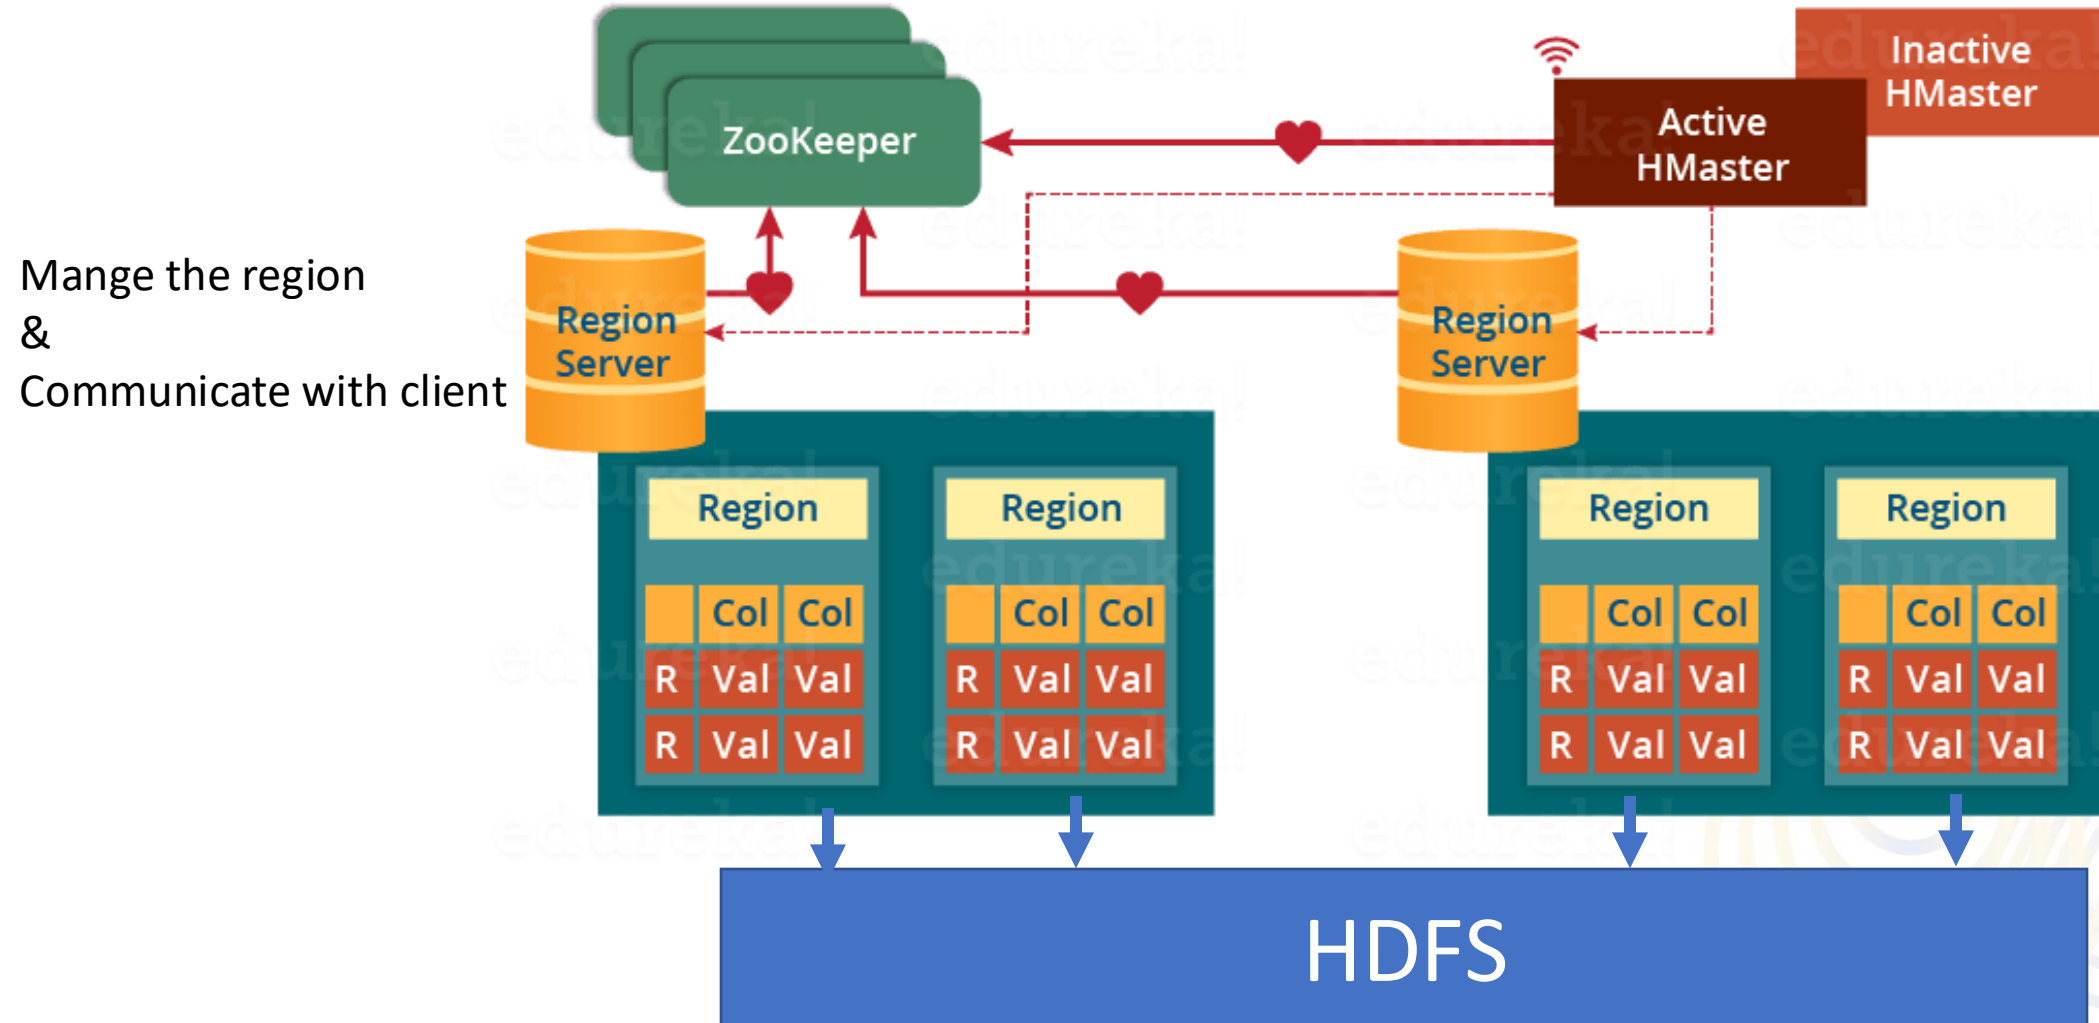
Hadoop Layer



HBASE Architecture

maintain configuration
Track node failure

Distribute/monitor services to region server





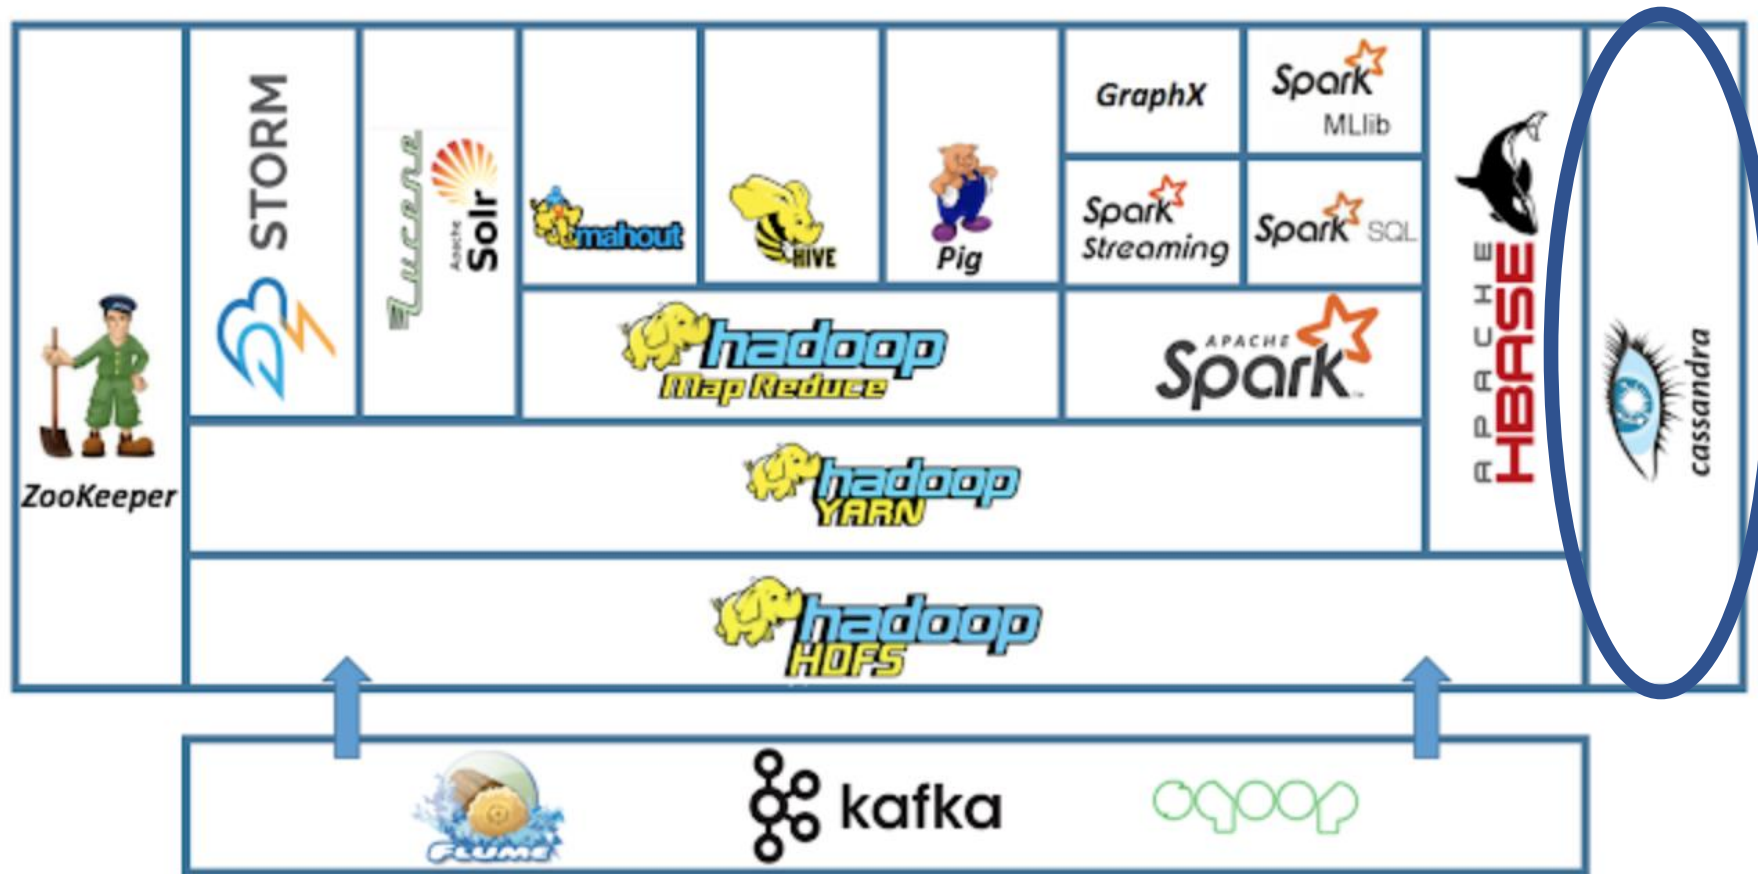
HBASE Column Family

Row Key	Column Family			
Row Key	Customers		Products	
Customer ID	Customer Name	City & Country	Product Name	Price
1	Sam Smith	California, US	Mike	\$500
2	Arijit Singh	Goa, India	Speakers	\$1000
3	Ellie Goulding	London, UK	Headphones	\$800
4	Wiz Khalifa	North Dakota, US	Guitar	\$2500

Figure: HBase Table

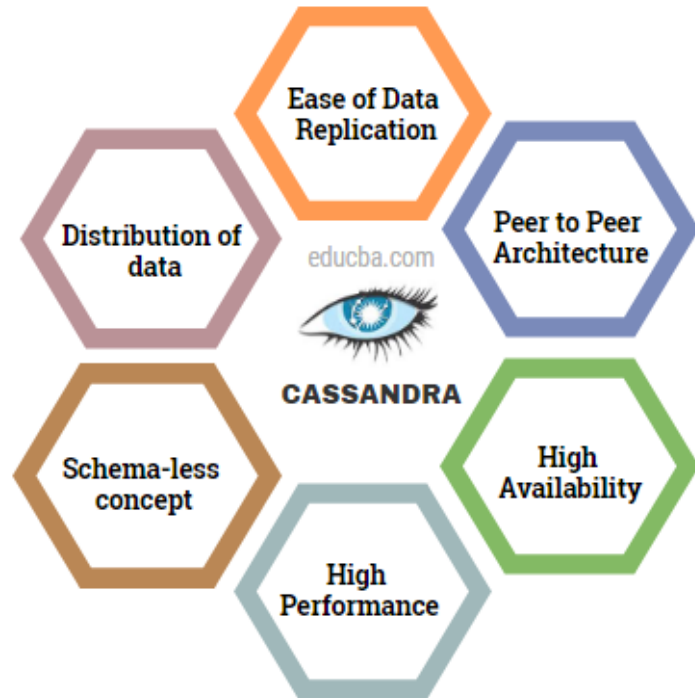
Cassandra

Hadoop Layer



Apache Cassandra

What is Cassandra



The CQL/Cassandra Mapping


```
CREATE TABLE example (
  A text,
  B text,
  C text,
  D text,
  E text,
  F text,
  PRIMARY KEY ((A,B),C,D)
);
```

A	B	C	D	E	F
a	b	c	d	e	f
a	b	c	g	h	i
a	b	j	k	l	m
a	n	o	p	q	r
s	t	u	v	w	x

	c:d:E	c:d:F	c:g:E	c:g:F	j:k:E	j:k:F
a:b	e	f	h	i	l	m
	o:p:E	o:p:F				
a:n	q	r				
			u:v:E	u:v:F		
			s:t	w	x	

Apache Cassandra

- Highly available
- Create peer-to-peer architecture
- Not good for a lot of update/delete Write to main memory and update periodically
- Good for fast write (append)
- Netflix, Uber, Spotify, Time series applications





Cassandra Architecture

Node



- The smallest physical item of a cluster
- No Primary/Secondary
- Equal to all nodes in the cluster
- Also Coordinator

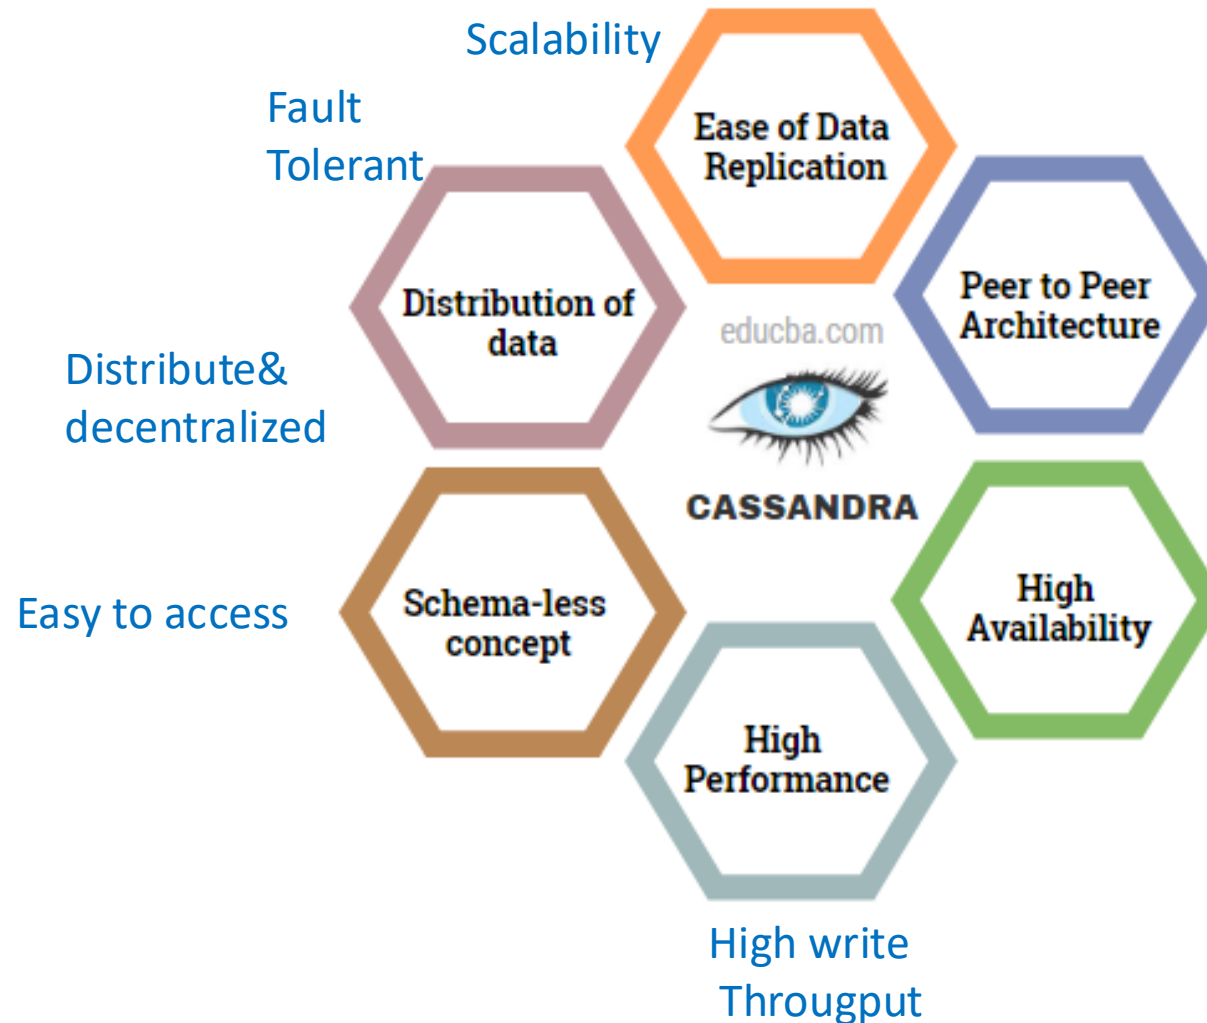
Ring/Cluster





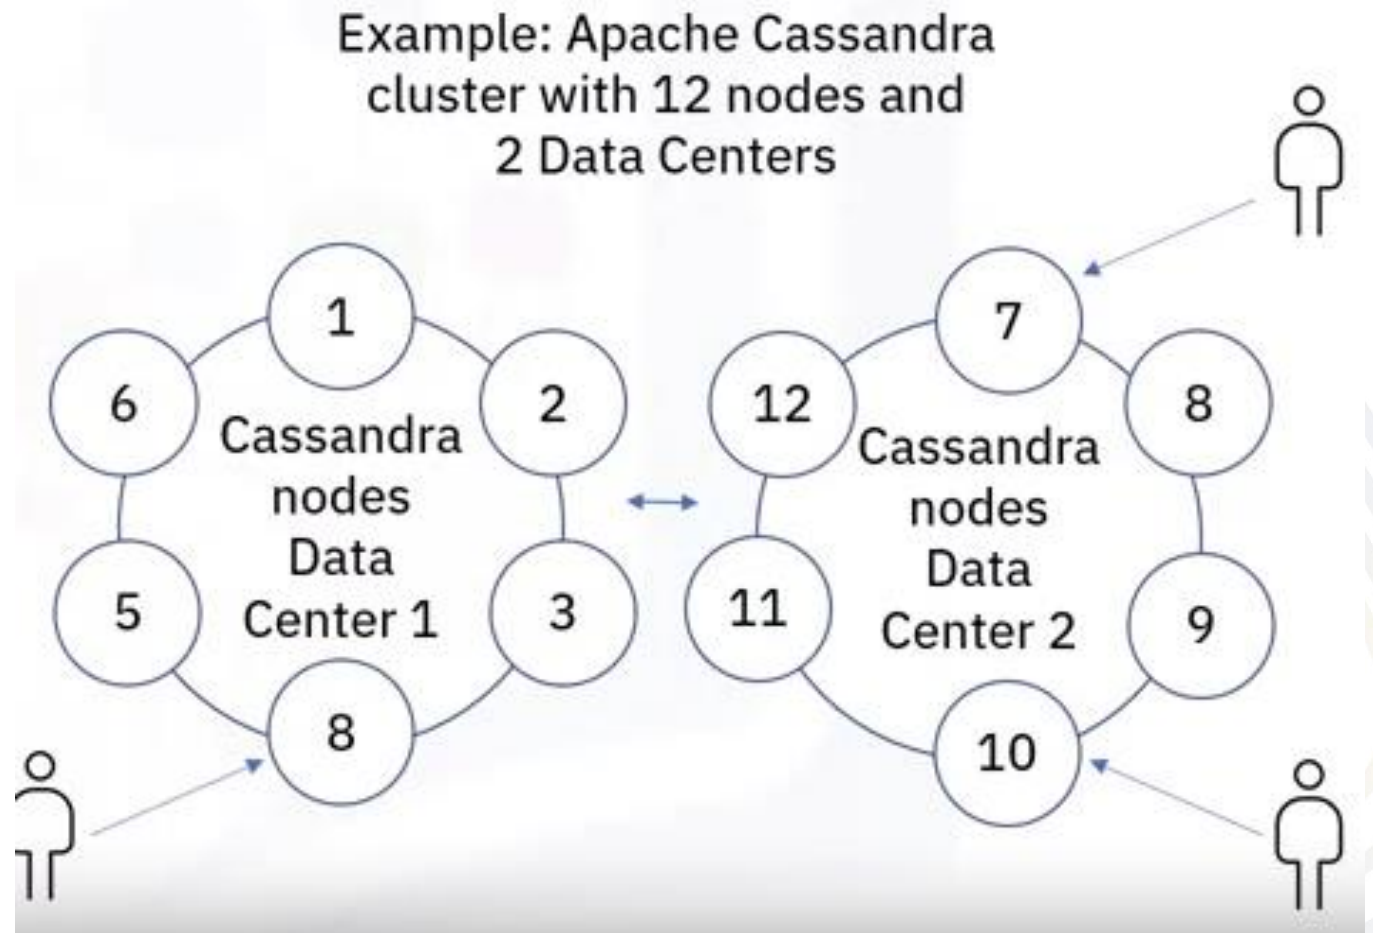
Mahidol University
Wisdom of the Land

Cassandra Key Feature



Distributed and decentralized

- Run on multiple distributed machines
- Application can route the user request optimally
- Use peer-to-peer architecture (gossip)
- All cluster are identical



Data Queries

Data Distribution




Table (Partition key= State)

Initial data → Partitions

UserID	Name	State
1	John	TX
2	Elaine	CA
3	Alex	NY
4	Jay	CA
5	Julio	NY
6	Elen	CA

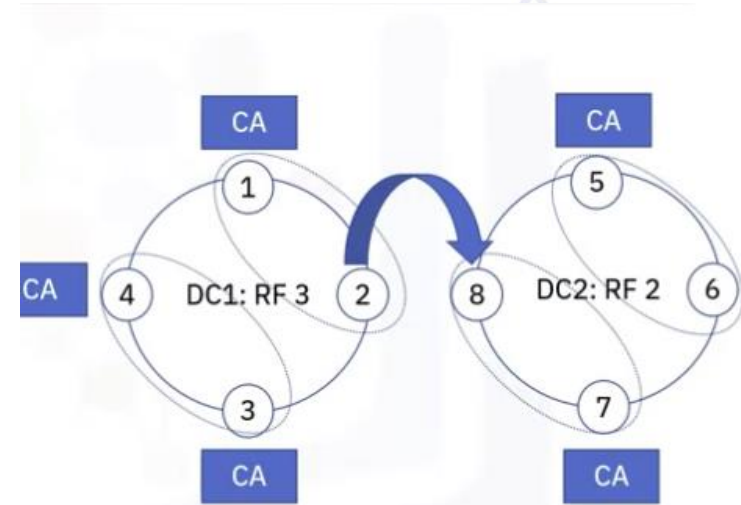
State	Name	UserID
TX	John	1
CA	Elaine	2
CA	Elen	6
CA	Jay	4
NY	Alex	3
NY	Julio	5

Query: All users in a state =>

PartitionKey = State Tokens

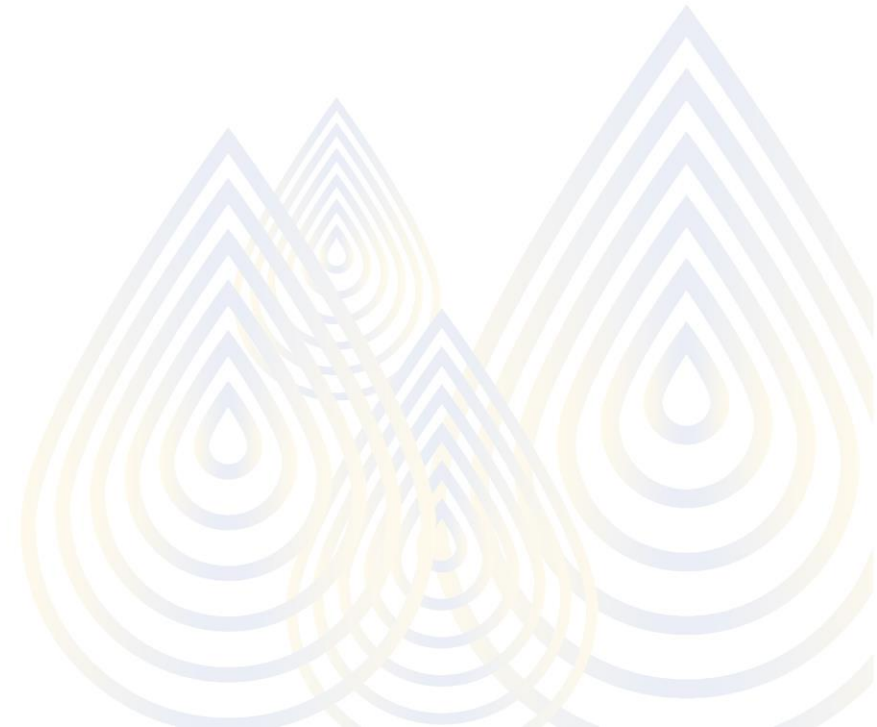
Data replication

- Replicas
 - Tells how many nodes contain the data(partition)
- Done clockwise – based on placement of rack/ data center
- Replication Factor: determine the number nodes to whole the replication of replicas



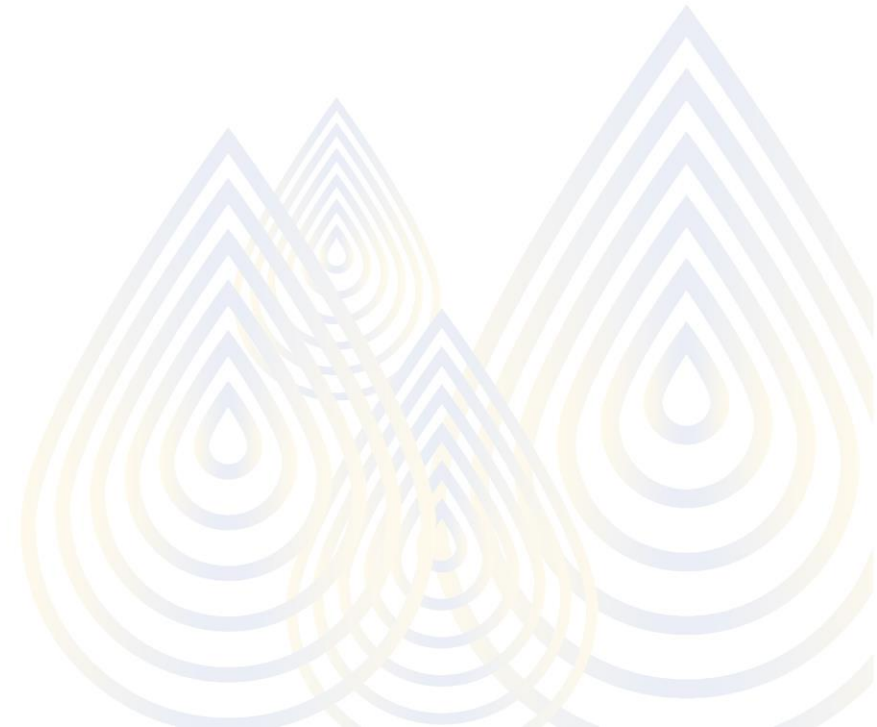
Availability vs Consistency

- Always available***
- **Tunable consistency**
 - Per operation set consistency(read/write)
 - Tune → Strong or Eventual
 - Conflict is solved during READ operations



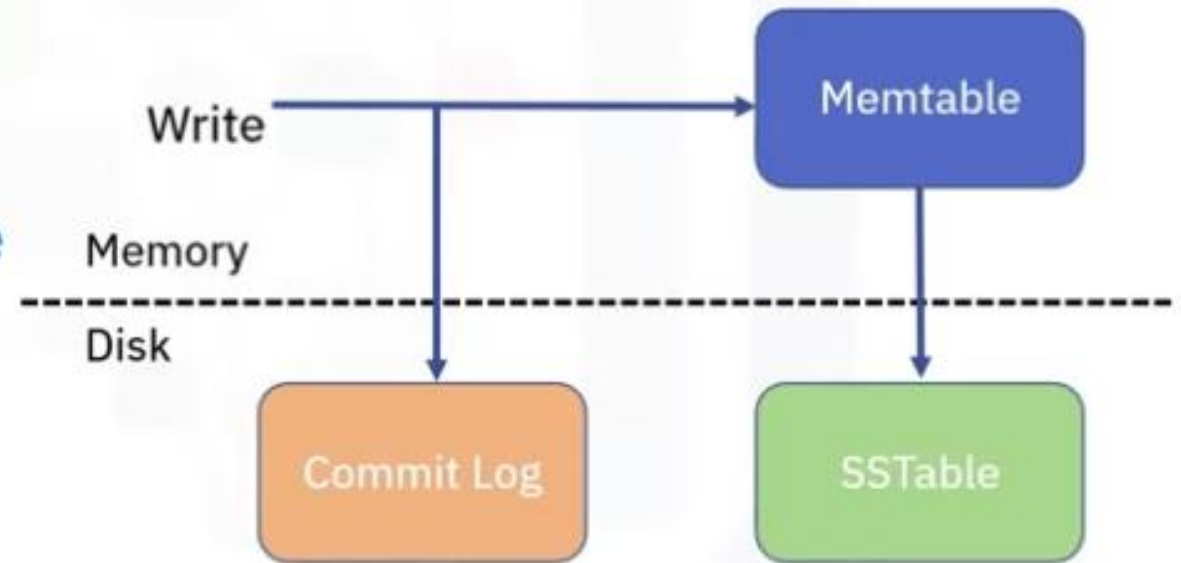
Fast and Linear Scalability

- Scale horizontally by adding new nodes in the cluster
- Performance increases linearly
- New nodes are automatically assigned tokens
- Addition/Removing nodes is done seamlessly



High write throughput

- Write are done in memory nodes
- Write in memory → Flush on disk
- All disk are append sequentially




CQL syntax

```
CREATE TABLE test (  
    groupid uuid,  
    name text,  
    occupation text,  
    age int,  
    PRIMARY KEY ((groupid), name));  
  
INSERT INTO test (groupid, name, occupation, age)  
    VALUES (1001, 'Thomas', 'engineer', 24), (1001, 'James', 'designer',  
30, (1002, 'Lily', 'writer', 35));  
  
SELECT * FROM test WHERE groupid = 1001;
```

Columnar Databases

Suitable for


- Large amounts of sparse data – Good for compression
- Can handle across cluster of nodes
- Best on column-wise data analytic
- Popular for counter type database
- Column can have expiry date as a parameter
 - Good for trial period
 - Anything with TTL(Time to live)



Columnar Databases

NOT Suitable for

- **Difficult query**— require changes to the column to make it easier
- More difficult updates




Use cases

- Log file
- E-commerce


Example The browsing history data might include the following columns for each user interaction:

- user_id (integer)
- timestamp (datetime)
- product_id (integer)
- category (string)
- price (float)
- device (string)
- duration_seconds (integer)
- action (e.g., 'view', 'add_to_cart', 'purchase')



Other types of database

- Redis
 - Enhanced key-value store
 - Store different types of data– hash, bitmap, ziplist
 - [Twitter usage](#)
- AsterixDB
 - Semi-structured data
- Solr
 - Text management
 - Search Engine

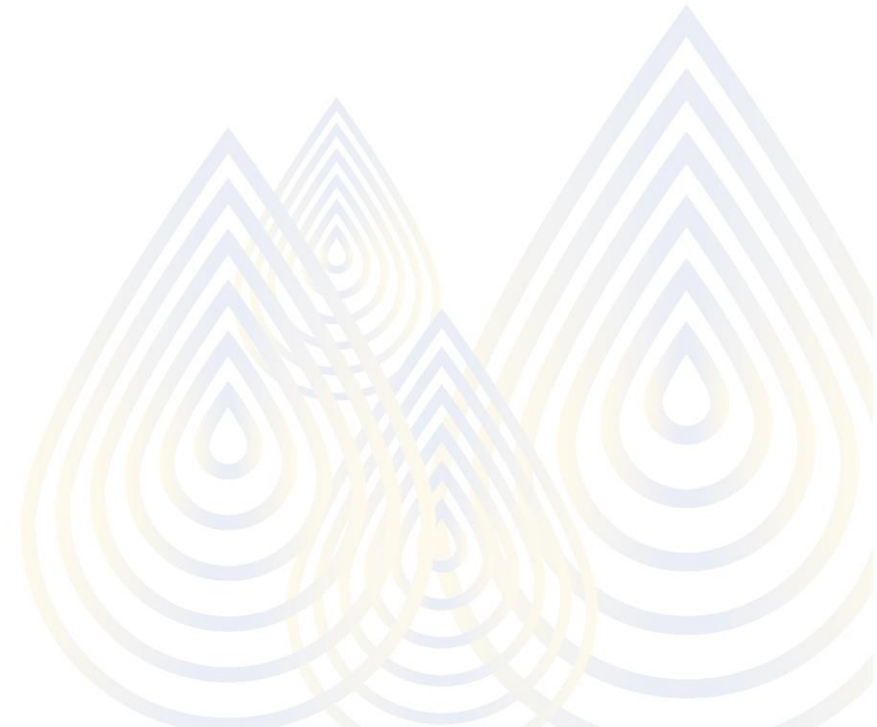


Cassandra

- [Deploy Cassandra using GCP Market place](#)
- [Intro to cassandra](#)
- [Insert/Update/delete](#)

More Examples

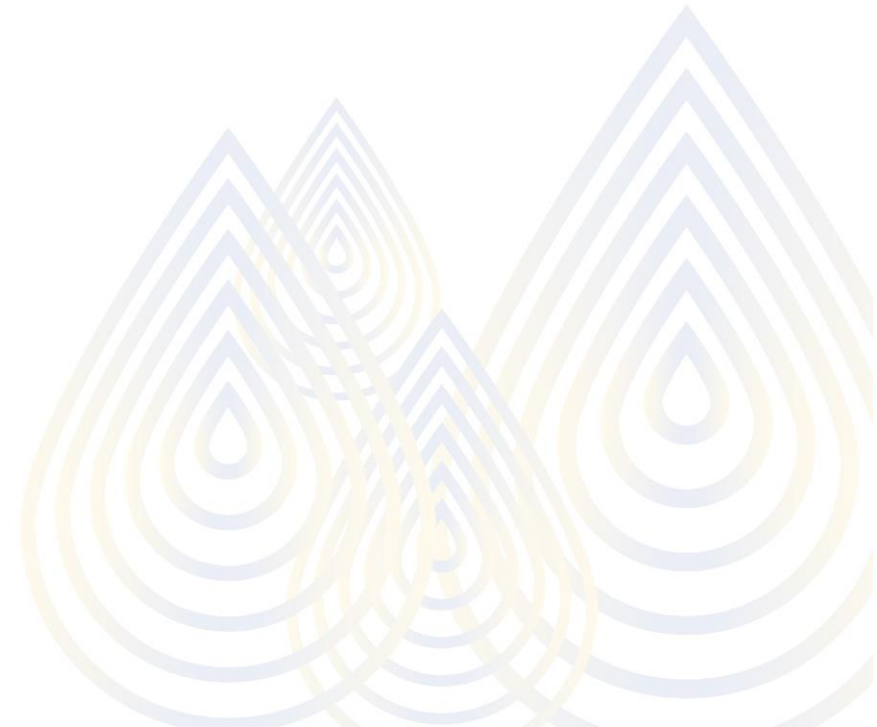
- [Order Management](#)
- [Sensor data](#)



Graph Database

Graph Database

- Store information in entities(node) and relationship(edges)
- Good for dataset with graph-like structure
- Not scale well horizontally




Graph Database

Highly connected and related

Example

- Social networking
- Routing
- Map application
- Recommendation engines



Summary

RDBMS

- Consistency
- Structured data(Fixed Schema)
- Transaction
- Join operations

- Non- RDBMS

- High Performance
- Unstructured data(Flexible Schema)
- Availability
- Easy Scalability

