

# BIG DATA PROCESSING

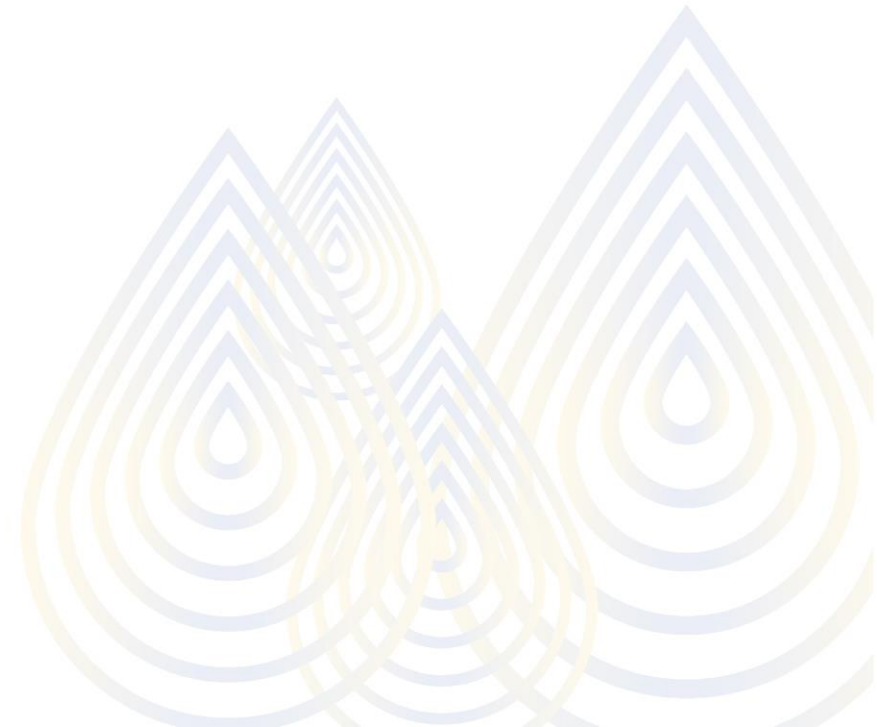
EGCI 466

Week5

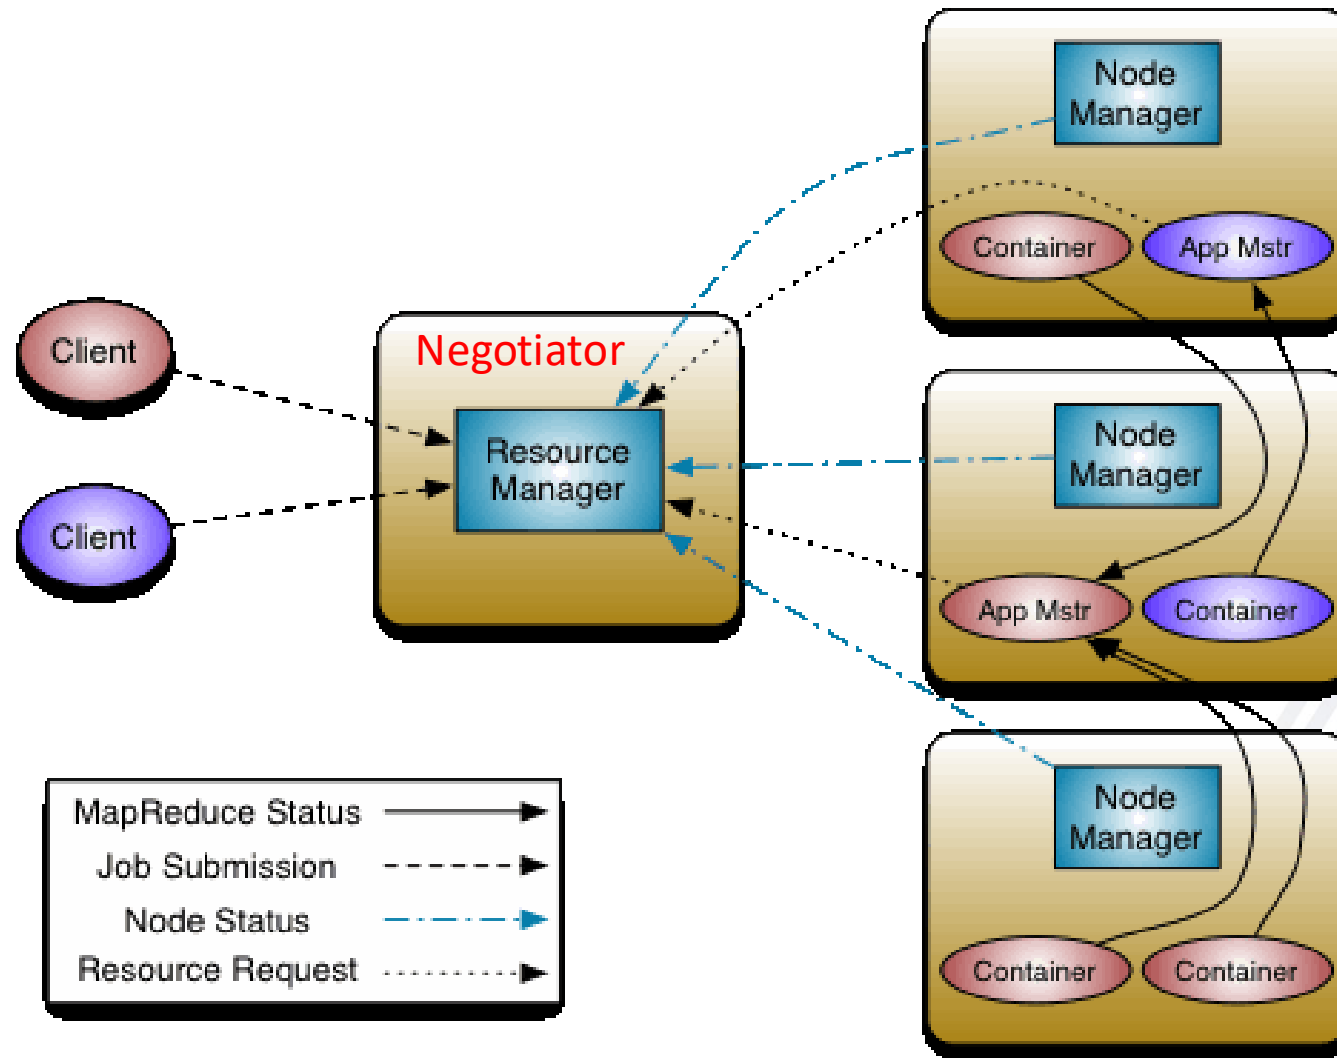


# Early Start

- Cloud Storage
  - Delete all previous dataproc storage (make sure they're deleted)
  - Create your cloud bucket (if you not have one)
- Create a dataproc
  - 1 Master (E2)
  - 2 worker nodes (E2)
  - Storage 50GB/clusters

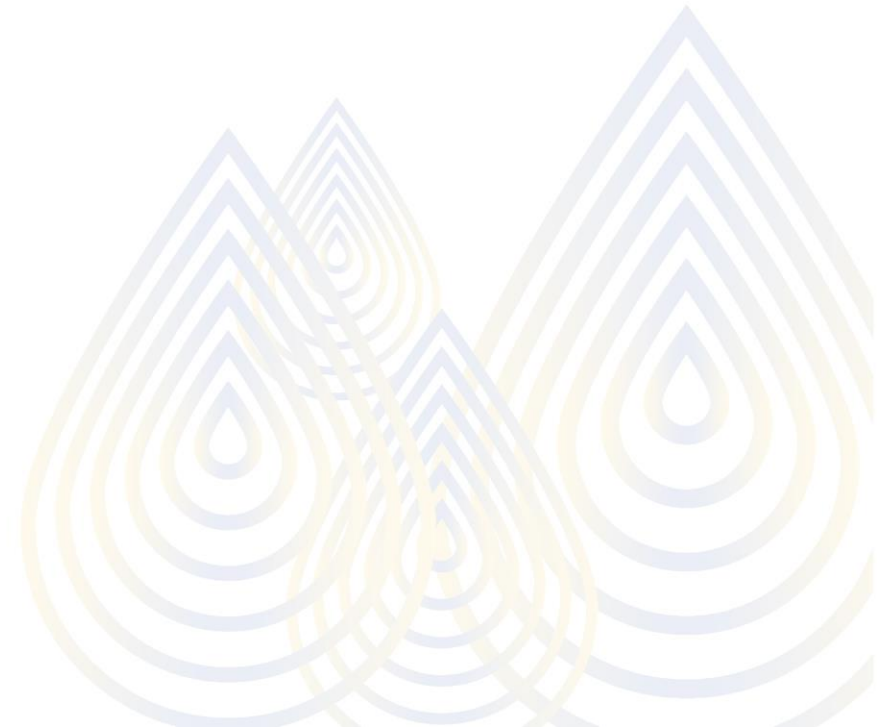


# Yarn Architecture



# Yarn:Benefit

- High resource utilization
- Data → Value
- One data set → Many applications



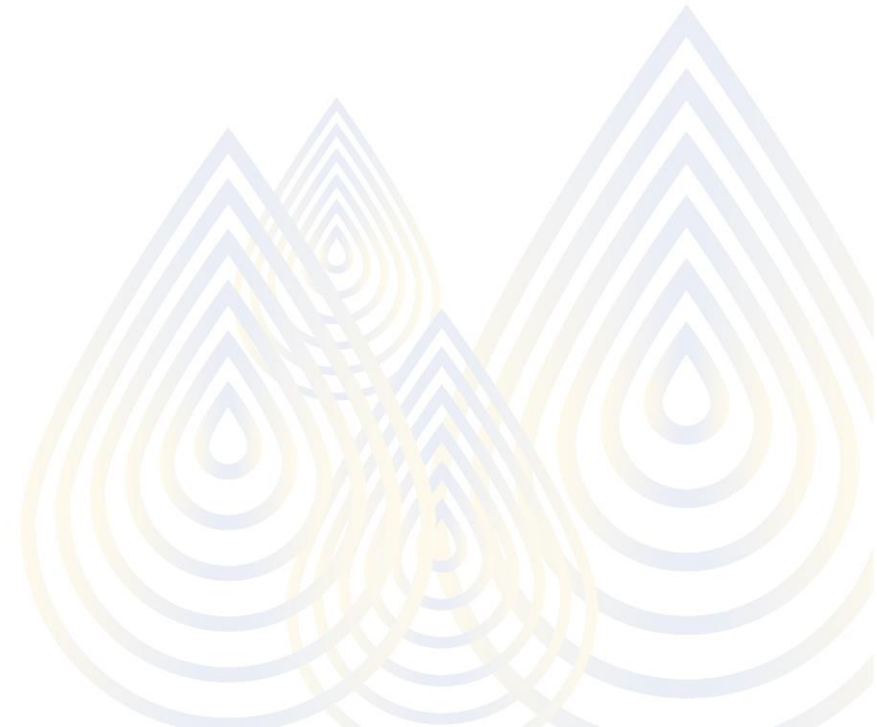
# Yarn

- Yet Another Resource Negotiator
- Prepares Hadoop RAM/CPU for
  - Batch
  - Interactive
  - Stream
  - Graph processing



# HIVE

- Data warehouse software
  - Design for managing tabular dataset
  - Data analysis
- 
- Scalable and fast and easy to use
  - Hive Query Language(HQL) is similar to SQL
  - It supports data clasnsing and filtering






# Hive and traditional RDBMS

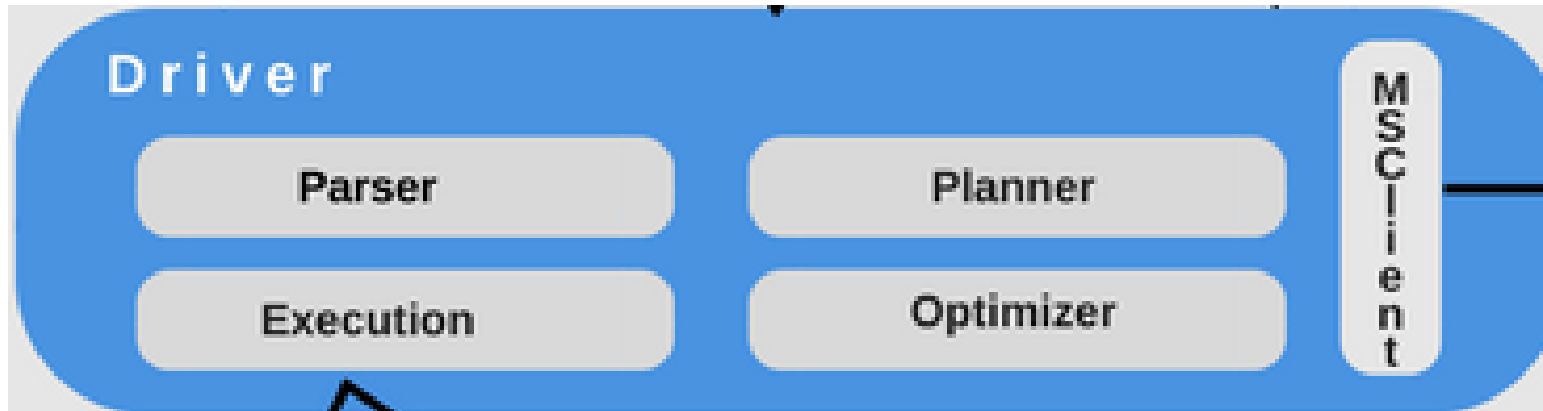
Traditional RDBMS	Hive
Used to maintain a Database and uses SQL	Used to maintain a data warehouse using Hive query language
Suited for real-time/dynamic data analysis like data from sensors	Suited for static data analysis like a text file containing names
Designed to read and write as many times as it needs	Designed on the methodology of write once, read many
Maximum data size it can handle is terabytes	Maximum data size it can handle is petabytes
Enforces that the schema must verify loading data before it can proceed	Doesn't enforce the schema to verify loading data
May not always have built-in for support data partitioning	Supports partitioning

# Hive architecture





# Hive Driver



**optimizer** transformations on the execution plan and splits the tasks to help speed up and improve efficiency.

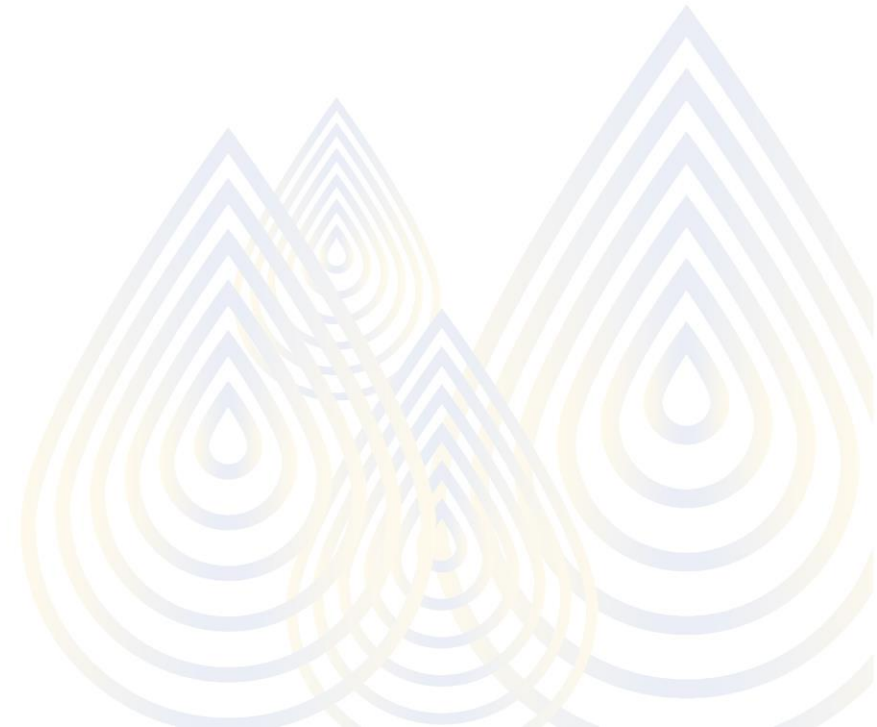
**executor** executes tasks after the optimizer has split the tasks.

# HBASE

- Columnar database
- Non-relational database system
- Runs on top of HDFS
- Provides a fault-tolerant way of sparse dataset
- Work well with real-time data and random read and write data access

# HBASE

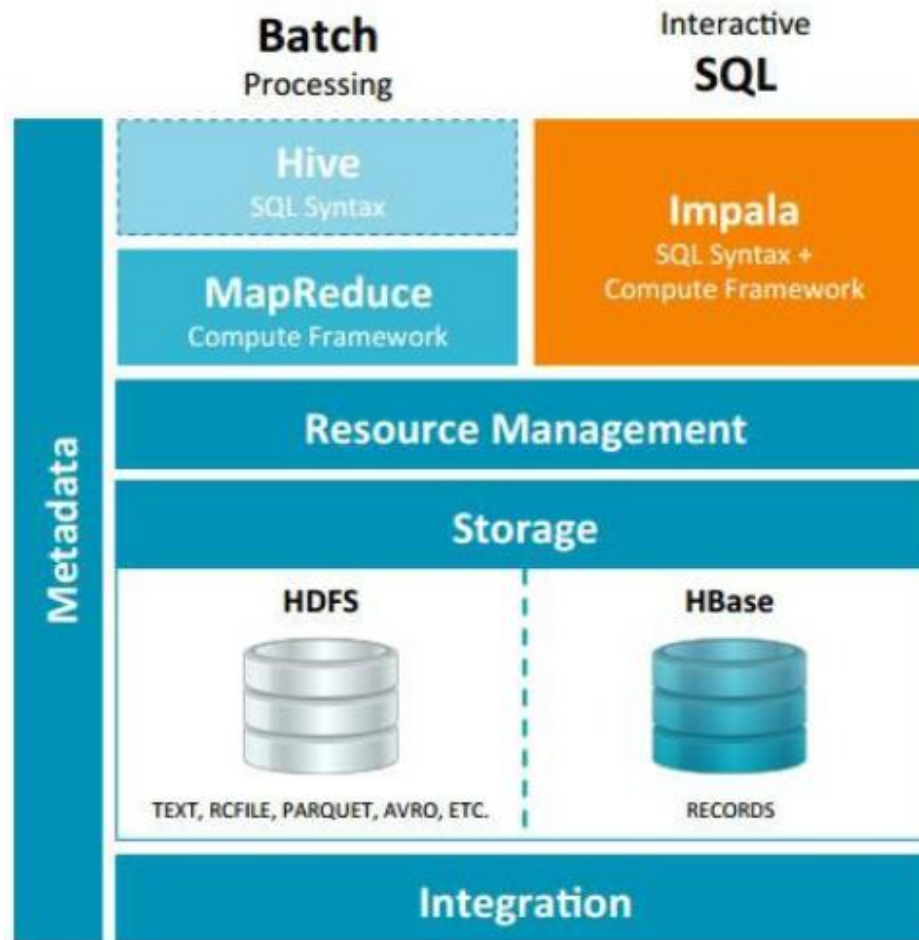
- Write heavy application
- Backup support for MapReduce jobs
- Consistent read & Write
- No fixed schema
- Easy to use with JAVA API



# HDFS and HBASE

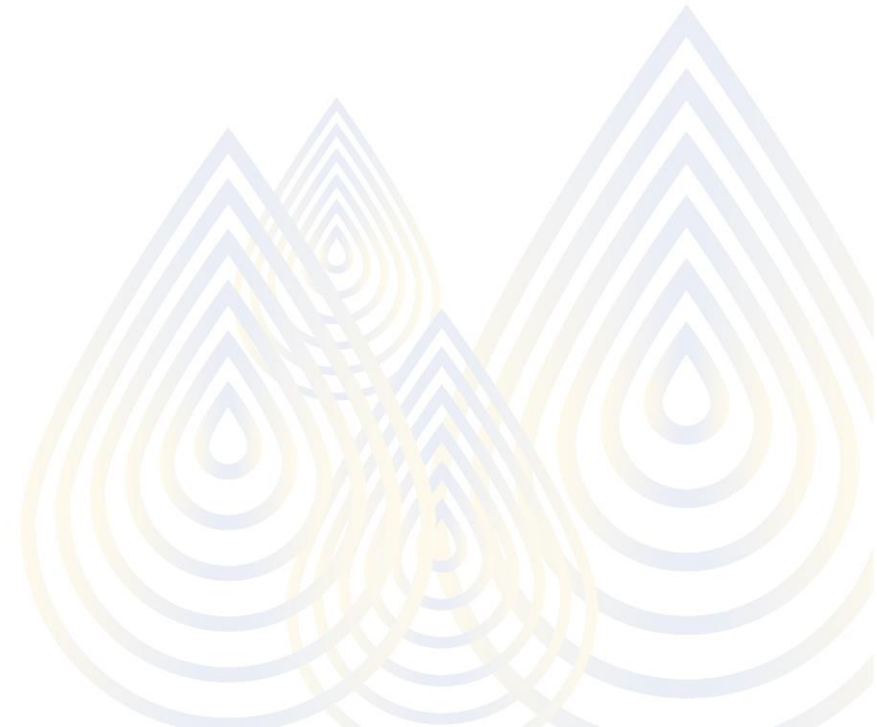
HDFS	HBASE
Stores data in distributed manner across different nodes	Store data in form of columns and rows in a table
Does not allow changes	Allow dynamic changes
Write once read many	Random read and write
Storing only	Store and processing big data

# Simple architecture




# Programming models

- Support Big data operations
  - Split volume data
  - Fast data access
  - Distribute computations to nodes
- Handle Fault Tolerance
  - Replicate data partition
  - Recover files when needed



# Traditional Parallel programming

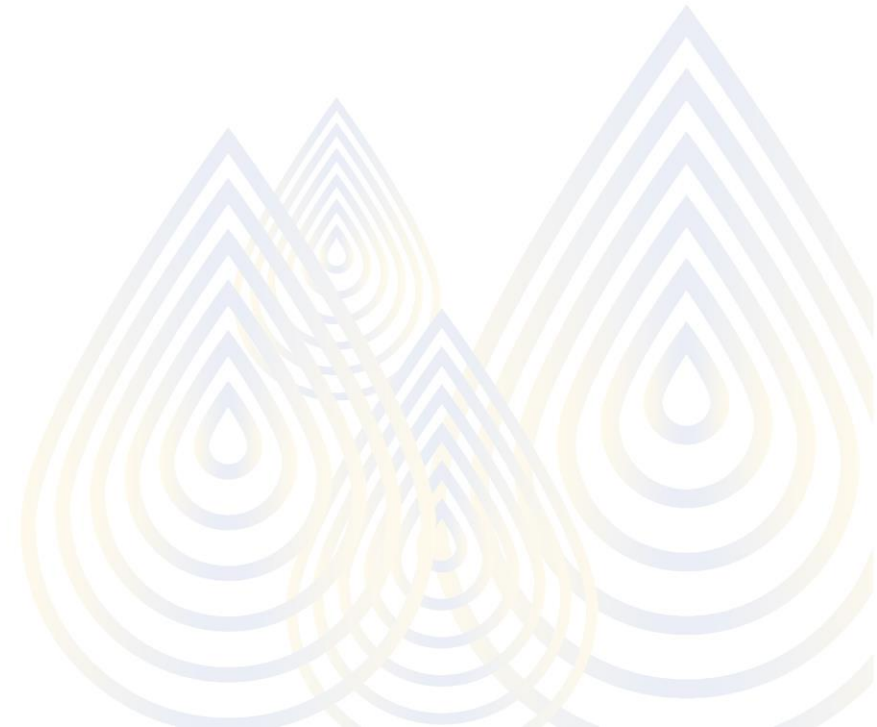
- Semaphore
- Threads
- Monitors
- Message Passing
- Shared Memory
- Locks





# Programming model for Big Data

- Enable Adding more racks
  - Add more resources
  - For increase data storage faster data access
  - **Scaling out**
- Optimized for specific data types
  - Table
  - Stream
  - Image
  - Document
  - Key-Model



# Functional Programming Example

- **Function creator**

$$f(x) = x + 1$$


**apply(f(x),**

**)=**

2
3
4
5
6

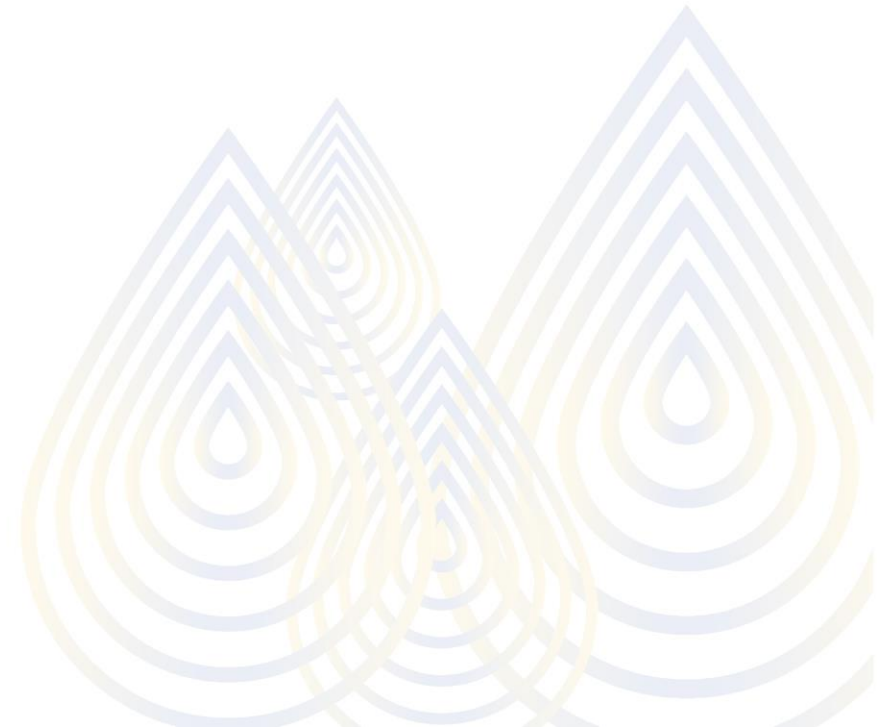
3
4
5
6
7

# Parallelization



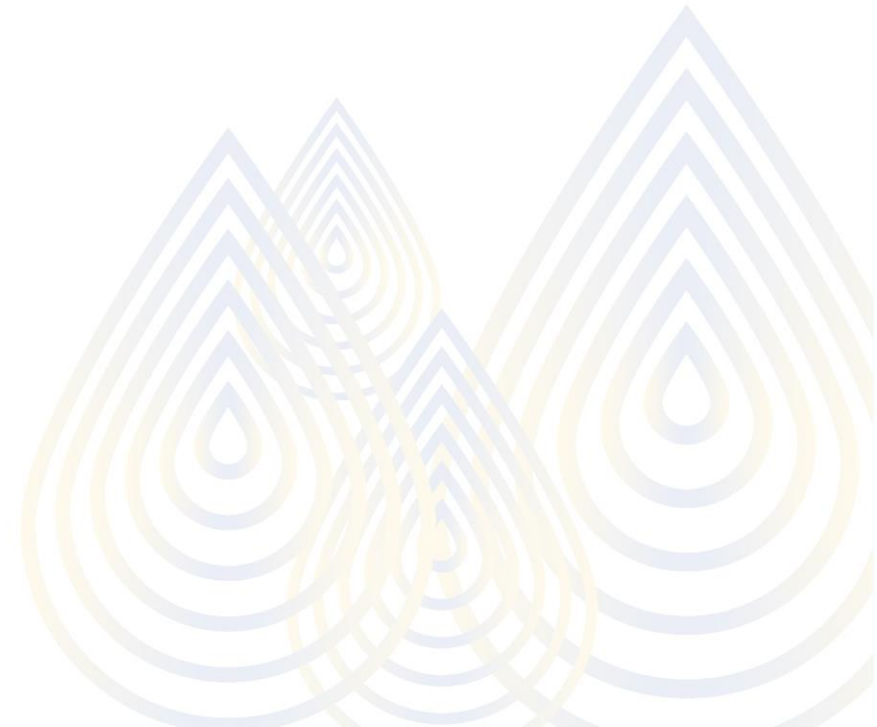
# MapReduce

- Programming model used in Hadoop for Big Data Processing
- Processing technique for distributed computing
- Consist of **Map** and Reduce **task**
- Can be coded in **Java**, C++, Python, Ruby and R



# Map Reduce

- Map
  - Apply operation to all elements
  - $f(x)=y$
  - Provide format of the operation for each data element
- Reduce
  - Summarize operation on elements



# Map and Reduce steps

File

- input File

Map

- Process data map to
- Key-Value pair

Rearrange

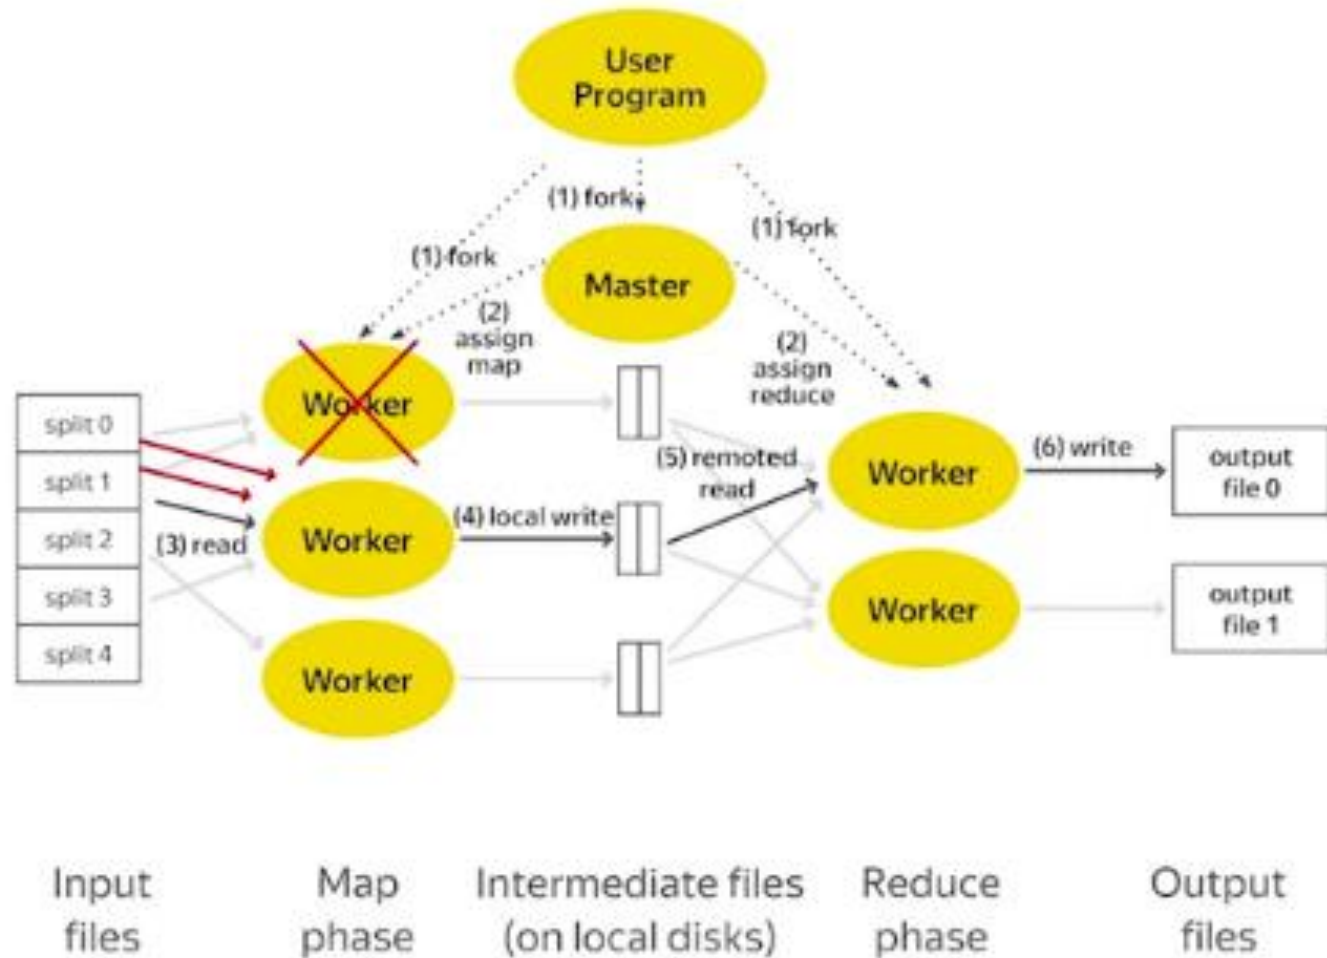
- sort
- reorganize

Reducer

- Aggregate and compute result
- Produce Final Output

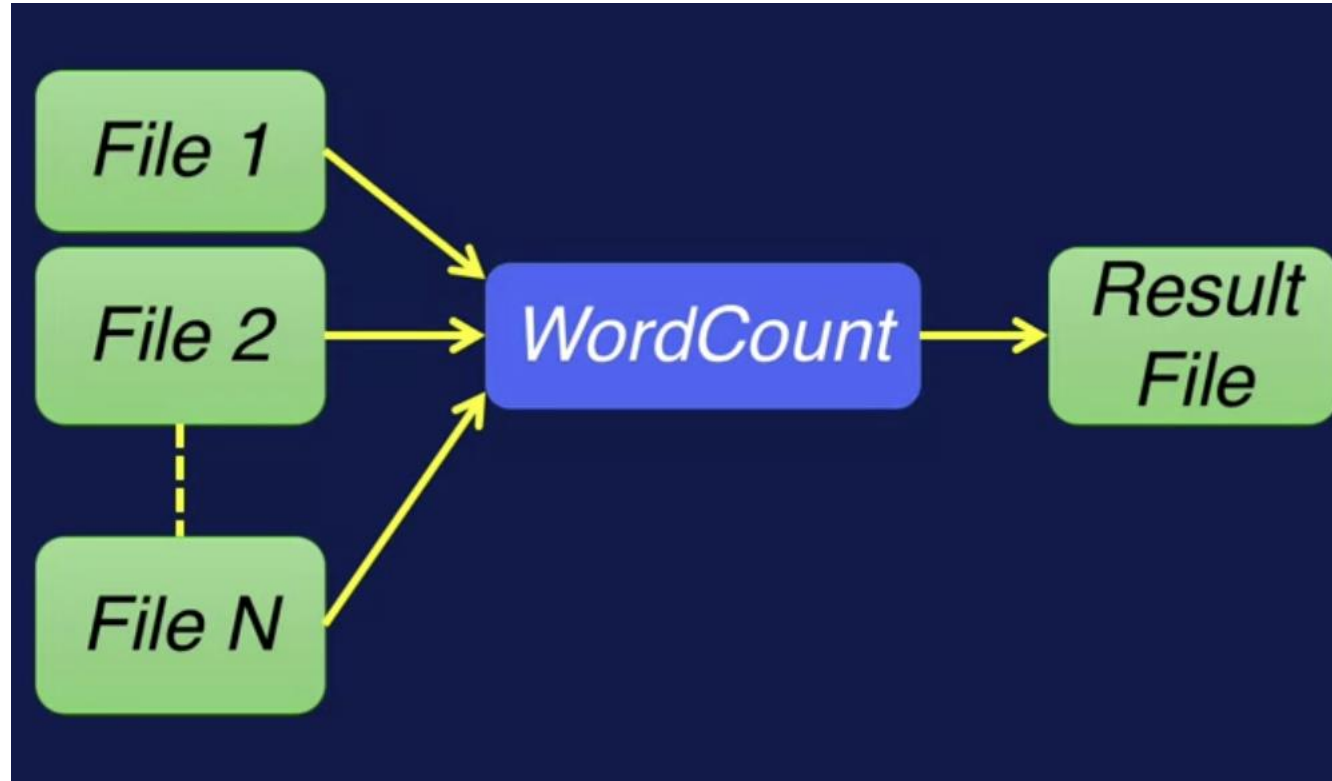
To keep track of its tasks → create a unique key

# Map Reduce





# WordCount Example



# Mapper

```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(Object key, Text value, Context context
                    ) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

# Key process

- Extends **mapper** – **defined** in Mapreduce Framework


## Input

- *key* Unique. words
- *Value* whole line of text

## Output

- map function: **key, value**
- IntWritable is always 1

(my,1) (bigger,1)  
(car,1) (than,1)  
(is,1) (your,1)  
(car,1)



# Reducer

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
                      Context context
                      ) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

# Key process

- Extends **mapper** – **defined** in Mapreduce Framework

## Input


- *key* Unique words
- *Value* =1

(my,1) (bigger,1)  
(car,1) (than,1)  
(is,1) (your,1)  
(car,1)

## Output

- *key* Unique words
- *Value* sum of all words

(my,1) (bigger,1)  
(car,2) (than,1)  
(is,1) (your,1)

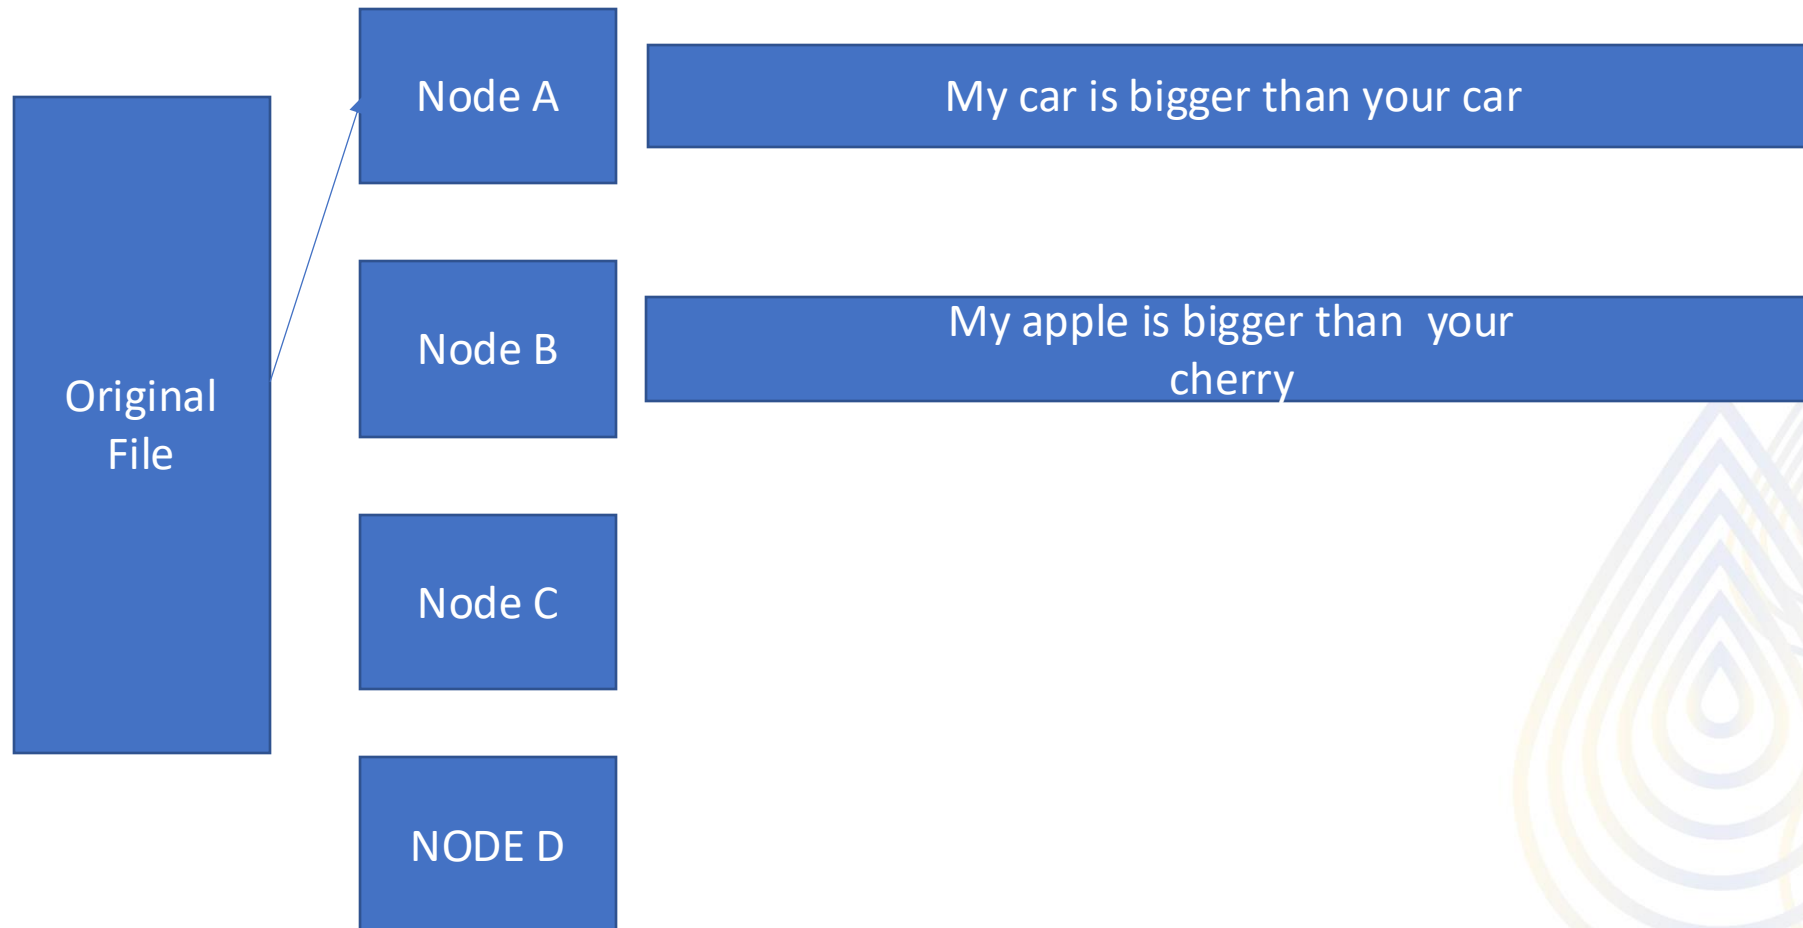


# Driver Code

```
import java.io.IOException;
import java.util.StringTokenizer;

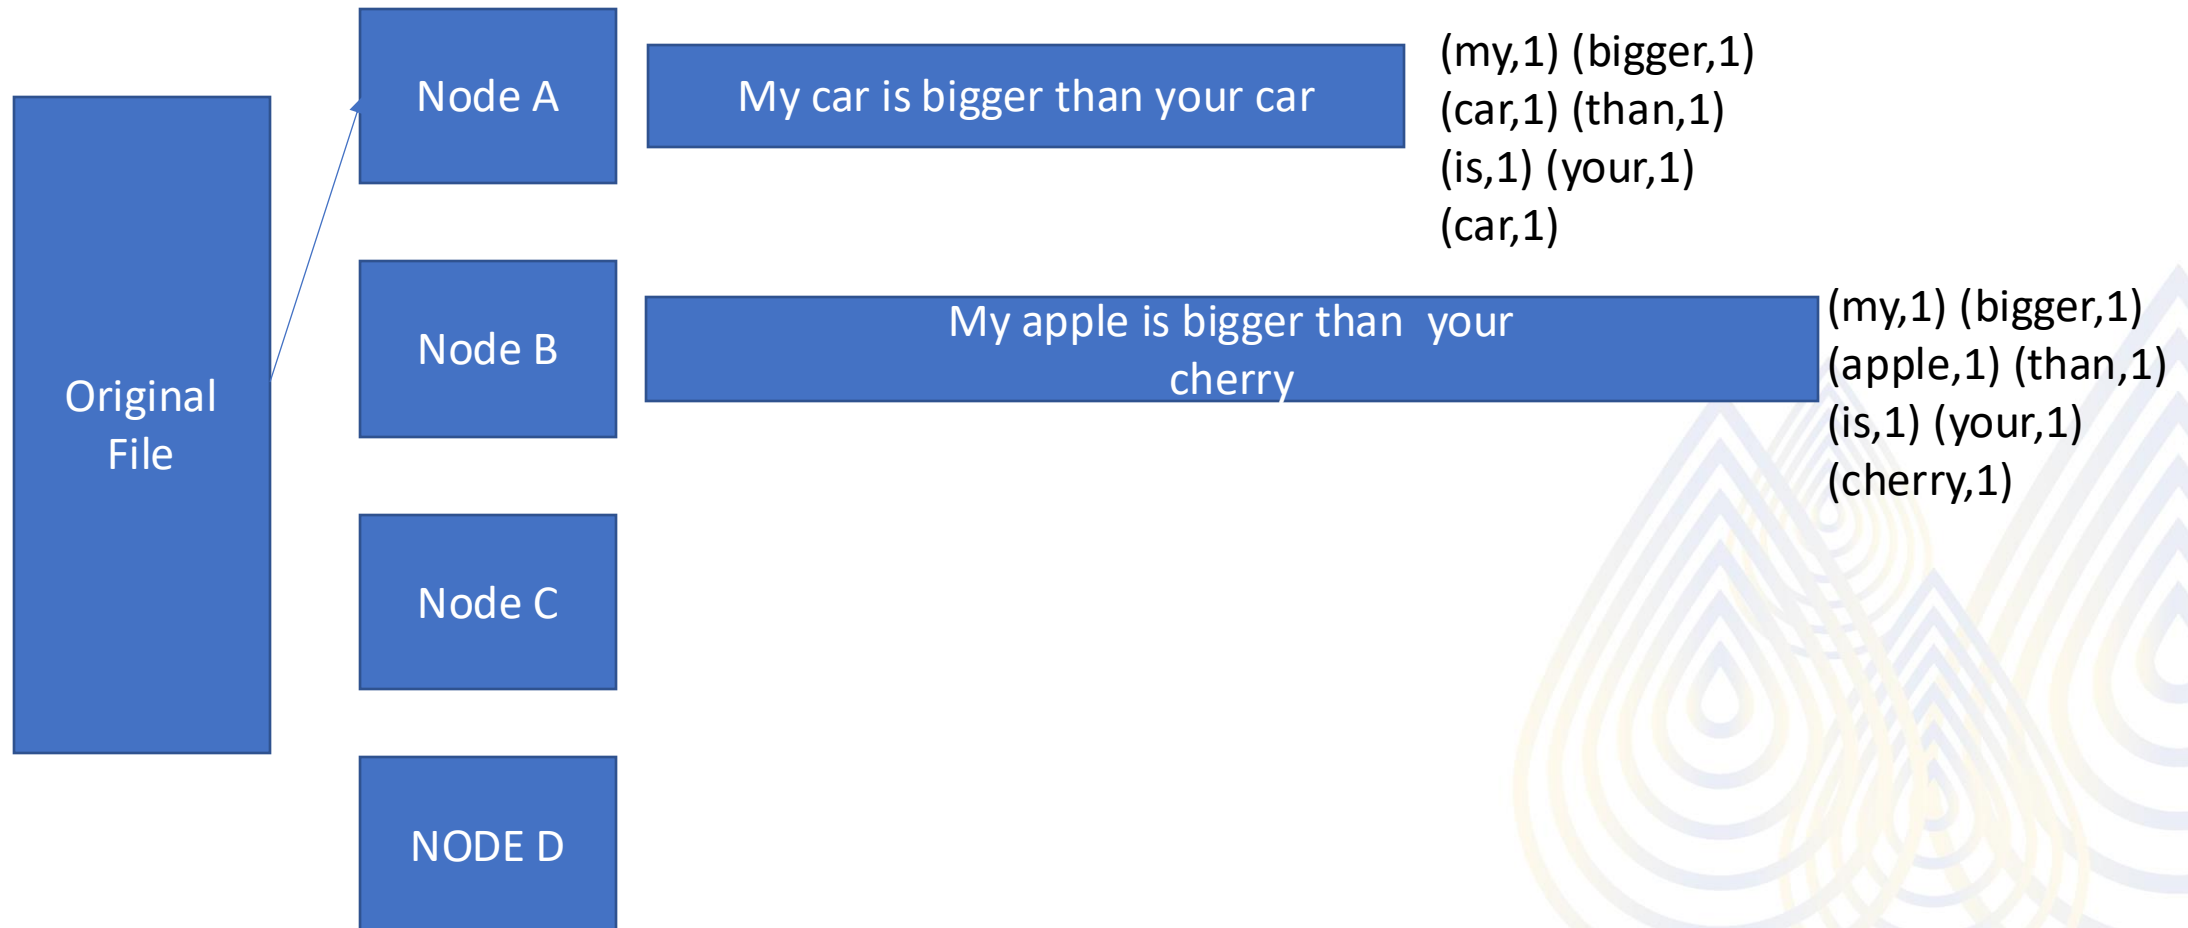
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

# How does it work?



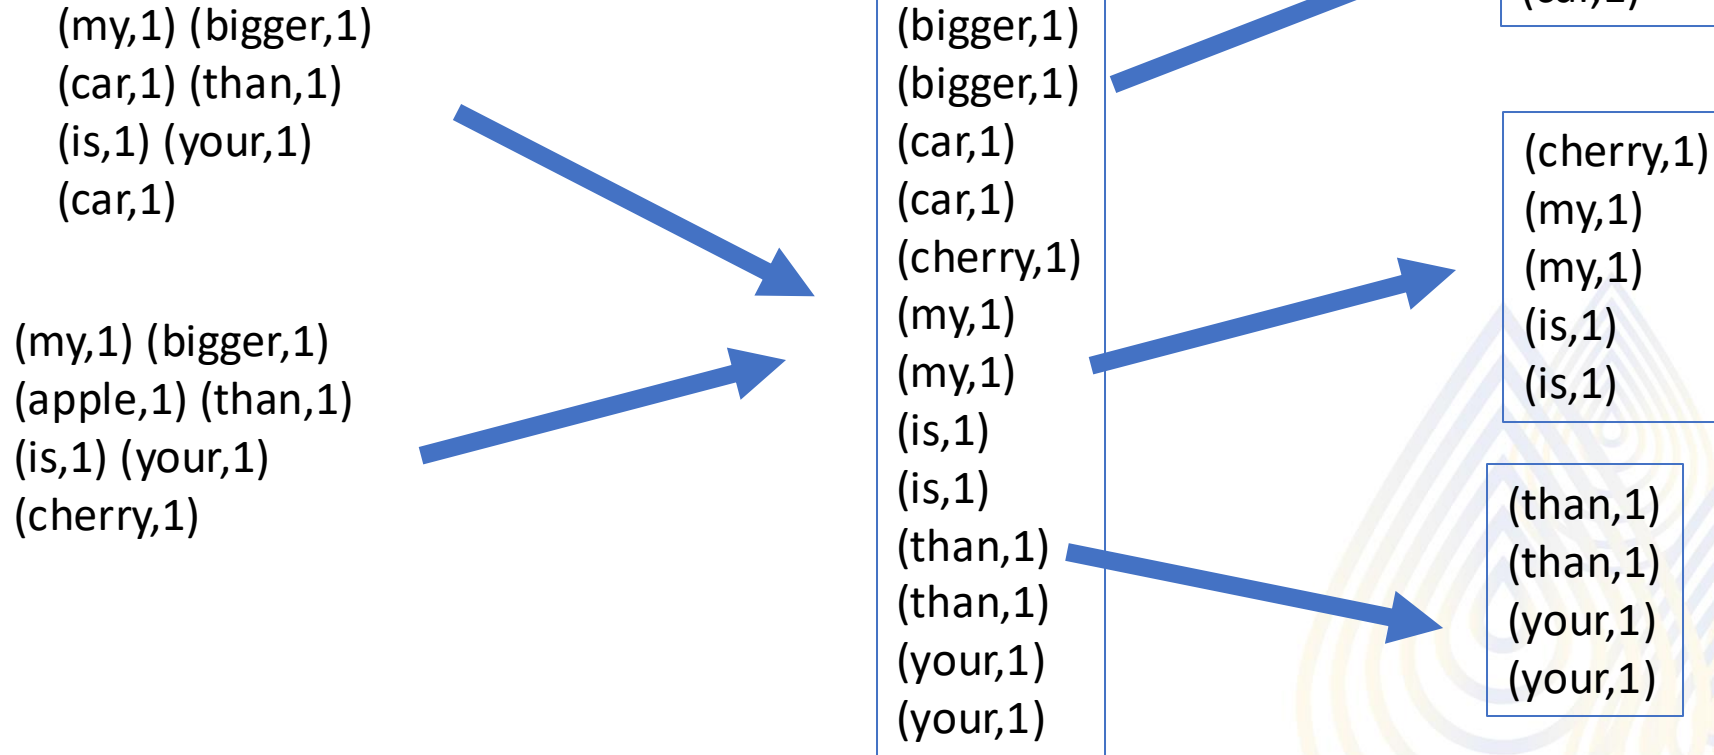


# Map step



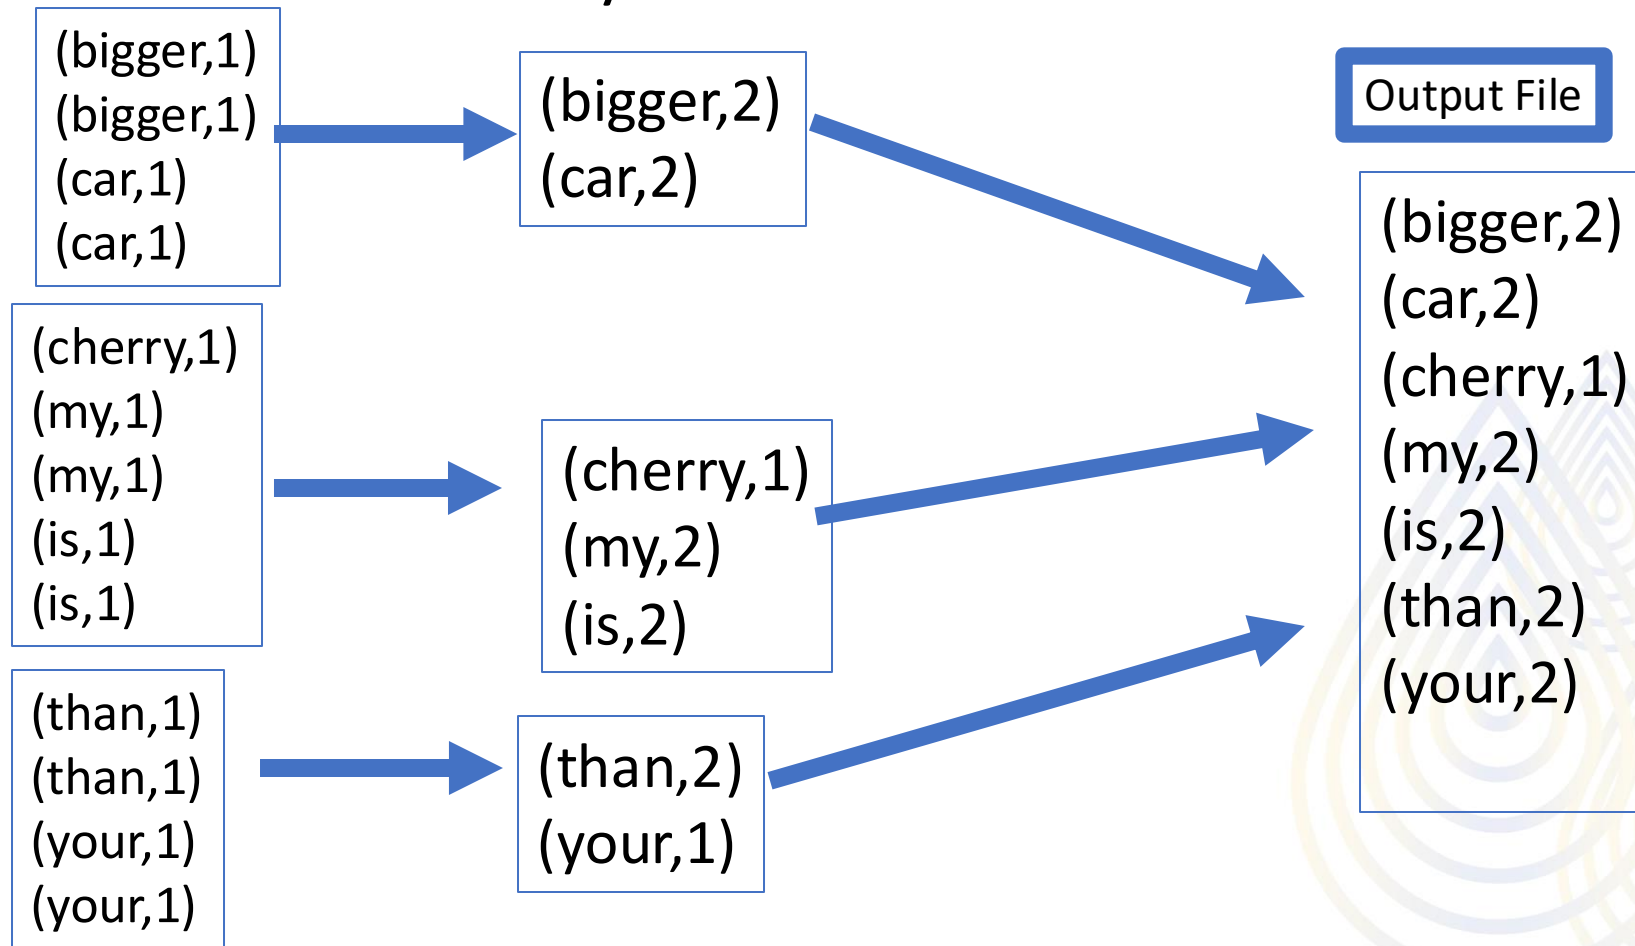
# sort and shuffle

- Pairs with same key  $\rightarrow$  move to the same node



# Reduce

- Add value to the same key



# Experiment 2: Manual Job creation

## Setup the environment

- `hadoop fs -mkdir /user/<user_name>`
- `export PATH=${JAVA_HOME}/bin:${PATH}`
- `export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar`

# Manual Job creation(2)

## Upload a file

**Wordcount.java** into your ssh (not hdfs)

## Compile a file

- Hadoop `com.sun.tools.javac.Main./WordCount.java`
- `jar cfwordcount.jar WordCount*.class`

## Run

- `hadoop jar WordCount.jar WordCount covid.txt out`

## Explore the new directory

## Download (for further use)

Put to your cloud storage (to make it safe)

# Explore output: Run the map-reduce

- `hadoop fs -ls out/`
- `hadoop fs -cat out/partxx`
- `hadoop fs -copyToLocal out/partxxxx`  
*Or to your bucket (hadoop fs -cp file gs:///)*
- → *Check the output*

# Why MapReduce

- Parallel Computing
- Divide-> Run Tasks -> Combine -> DONE
- Process data in structured/un-structured file formats
- Support different programming languages
- Support platform for analysis





# Submit job

[←](#) Submit a job

Job ID \*

job-a5a9ed1f

Region \*

us-central1

Cluster \*

cluster-6a4e

Job type \*

Hadoop

Main class or JAR \*

hdfs:///user/mingmanas\_siv/wordcount.jar

Jar files

Archive files

Arguments

WordCount

gs://mingmanas123/words.txt

gs://mingmanas123/test

Max. restarts per hour

Specifies the Cloud Dataproc regional service, which determines what clusters are available.

The fully qualified name of a class in a provided or standard JAR file, for example, com.example.wordcount, or a provided JAR file to use the main class of that JAR file

JAR files are included in the CLASSPATH. Can be a GCS file with the gs:// prefix, an HDFS file on the cluster with the hdfs:// prefix, or a local file on the cluster with the file:// prefix.

Archive files are extracted in the Spark working directory. Can be a GCS file with the gs:// prefix, an HDFS file on the cluster with the hdfs:// prefix, or a local file on the cluster with the file:// prefix. Supported file types: .jar, .tar, .tar.gz, .tgz, .zip.

Additional arguments to pass to the main class. Press Return after each argument.

I leave blank if you don't want to allow automatic restarts on job failure. [I learn more](#)

Jar files

Related Arguments

## Succeeded job

Job ID	job-e4816a17
Job UUID	56cb4381-1980-4ec4-a02d-b341c1b6e158
Type	Dataproc job
Status	✔ Succeeded

MONITORING

CONFIGURATION

The charts below represent the metrics from the cluster that this job ran on, scoped to the time that the job ran. The charts for a job may lag behind the job run by several minutes.

SAVE AS DASHBOARD

RESET ZOOM

YARN memory

YARN per

15GiB

Output

LINE WRAP: OFF

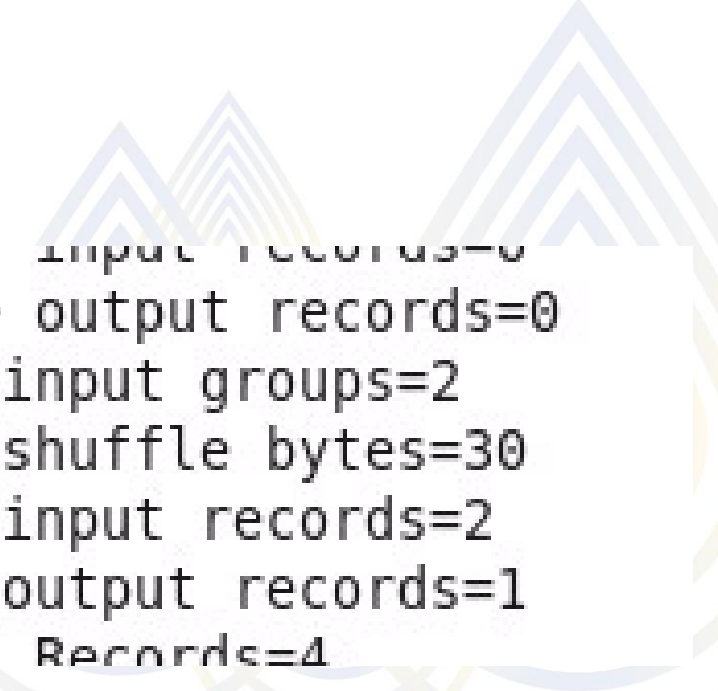
```

Total committed heap usage (bytes)=1452277760
Peak Map Physical memory (bytes)=727236608
Peak Map Virtual memory (bytes)=4723675136
Peak Reduce Physical memory (bytes)=436551680
Peak Reduce Virtual memory (bytes)=4730257408
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=5458199
File Output Format Counters
  Bytes Written=1435536
    
```

# Not familiar with JAVA?

- Python streamers

- `hadoop jar ///usr/lib/hadoop/hadoop-streaming.jar \`  
`-files mapper.py,reducer.py \`  
`-mapper 'python mapper.py' \`  
`-numReduceTasks 1 \`  
`-reducer 'python reducer.py' \`  
`-input covid.txt \`  
`-output wc1/`



```
Combine Input Records=0  
Combine output records=0  
Reduce input groups=2  
Reduce shuffle bytes=30  
Reduce input records=2  
Reduce output records=1  
Spilled Records=4
```

# mapper and reducer files

```
import sys
line_count=0
for line in sys.stdin:
    line_count+=1
print '%d' % line_count
```

mapper.py

```
import sys
line_count=0
for line in sys.stdin:
    line_count+=int(line)
print '%d' % line_count
```

reducer.py

# Standard Program

- **Use this file**

`:///usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples-3.3.6.jar`

- **Explore its usage**

- `hadoop jar [your file]`

- **Run the example**

- `hadoop jar [your file] wordcount [your gs://file] [your folder]`

- **Explore output**

- **Check Status in web interface**

# Submit job

iproc

rs

rs

lows

scaling policies

vices

store

ation

onent exchange

ench

Submit a job

Job ID \*

job-e4816a17

Region \*

us-central1

Specifies the Cloud Dataproc regional service, which determines what clusters are available.

Cluster \*

cluster-c483

Job type \*

Hadoop

Main class or JAR \*

file:///usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar

The fully qualified name of a class in a provided or standard JAR file, for example, com.example.wordcount, or a provided JAR file to use the main class of that JAR file

Jar files

JAR files are included in the CLASSPATH. Can be a GCS file with the gs:// prefix, an HDFS file on the cluster with the hdfs:// prefix, or a local file on the cluster with the file:// prefix.

Archive files

Archive files are extracted in the Spark working directory. Can be a GCS file with the gs:// prefix, an HDFS file on the cluster with the hdfs:// prefix, or a local file on the cluster with the file:// prefix. Supported file types: jar, tar, tar.gz, tgz, zip.

Arguments

wordcount gs://mingmanas123/words.txt gs://mingmanas123/wordcountOut/

Additional arguments to pass to the main class. Press Return after each argument.

Max. restarts per hour

Leave blank if you don't want to allow automatic restarts on job failure. [Learn more](#)

Jar files

Related Arguments

## Succeeded job

Job ID	job-e4816a17
Job UUID	56cb4381-1980-4ec4-a02d-b341c1b6e158
Type	Dataproc job
Status	<span>✓</span> Succeeded

MONITORING

CONFIGURATION

The charts below represent the metrics from the cluster that this job ran on, scoped to the time that a job may lag behind the job run by several minutes.

SAVE AS DASHBOARD

RESET ZOOM

YARN memory

YARN per

Output


LINE WRAP: OFF

```

Total committed heap usage (bytes)=1452277760
Peak Map Physical memory (bytes)=727236608
Peak Map Virtual memory (bytes)=4723675136
Peak Reduce Physical memory (bytes)=436551680
Peak Reduce Virtual memory (bytes)=4730257408
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=5458199
File Output Format Counters
  Bytes Written=1435536

```

# MORE EXAMPLE:K MEAN



# K-mean Examples

**Classify:** Assign observations to closest cluster center


$$z_i \leftarrow \arg \min_j \|\mu_j - \mathbf{x}_i\|_2^2$$

**Map:** For each data point, given  $(\{\mu_j\}, \mathbf{x}_i)$ , emit  $(z_i, \mathbf{x}_i)$

**Recenter:** Revise cluster centers as mean of assigned observations

$$\mu_j = \frac{1}{n_j} \sum_{i: z_i=k} \mathbf{x}_i$$

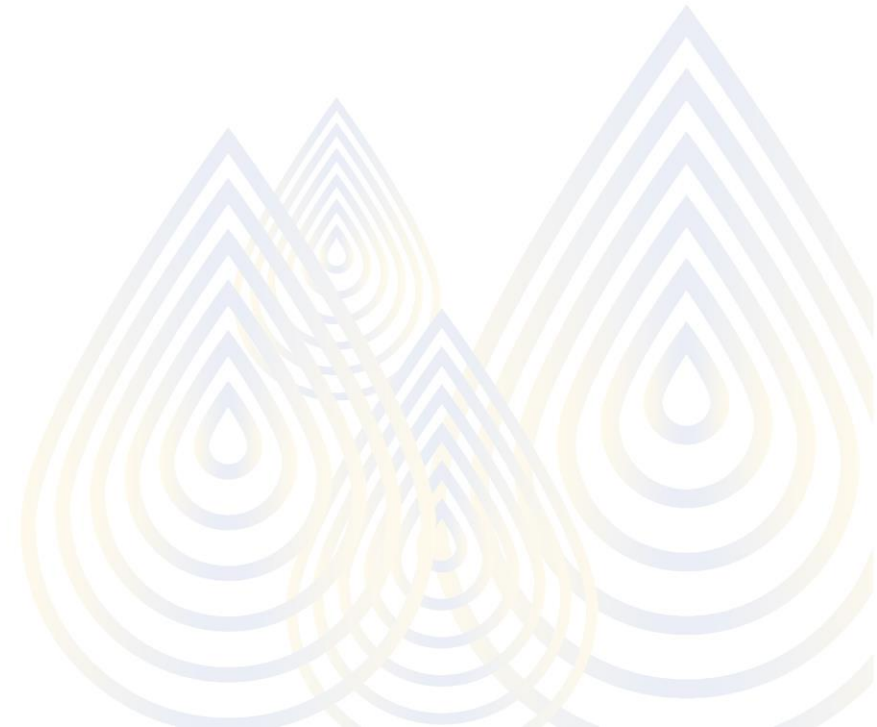
**Reduce:** Average over all points in cluster  $j$  ( $z_i=k$ )





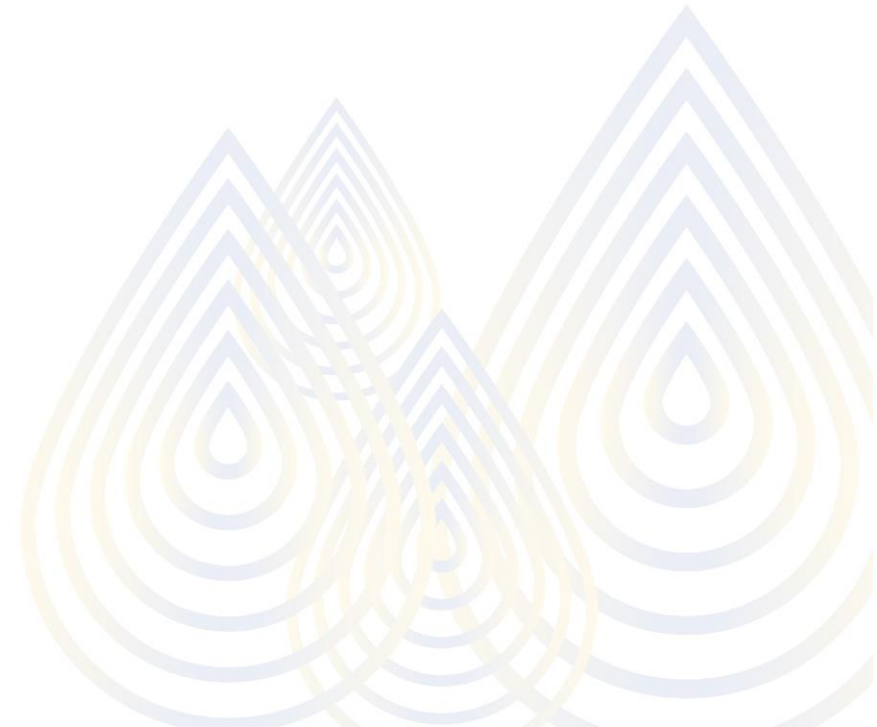
# Examples

- Distributed grep
- Count of URL Access Frequency
- Reverse Web-link graph:
  - $(\text{target}, \text{source}) \rightarrow (\text{target}, \text{list}(\text{source}))$
- Term-vector per Host
  - $\langle \text{word}, \text{frequency} \rangle \rightarrow \langle \text{hostname}, \text{term vector} \rangle$
- Inverted index
- Distributed Sort



# More Examples

- Joining data
- Vector multiplication
  - $\langle \text{index}, \text{value} \rangle$
- Feature selection
- Clustering
- word frequency



# Additional Reading

[Google](#): Summary

[Google White Paper: https://research.google.com/archive/mapreduce-osdi04.pdf](https://research.google.com/archive/mapreduce-osdi04.pdf)

