



Mahidol University  
*Wisdom of the Land*

# BIG DATA PROCESSING

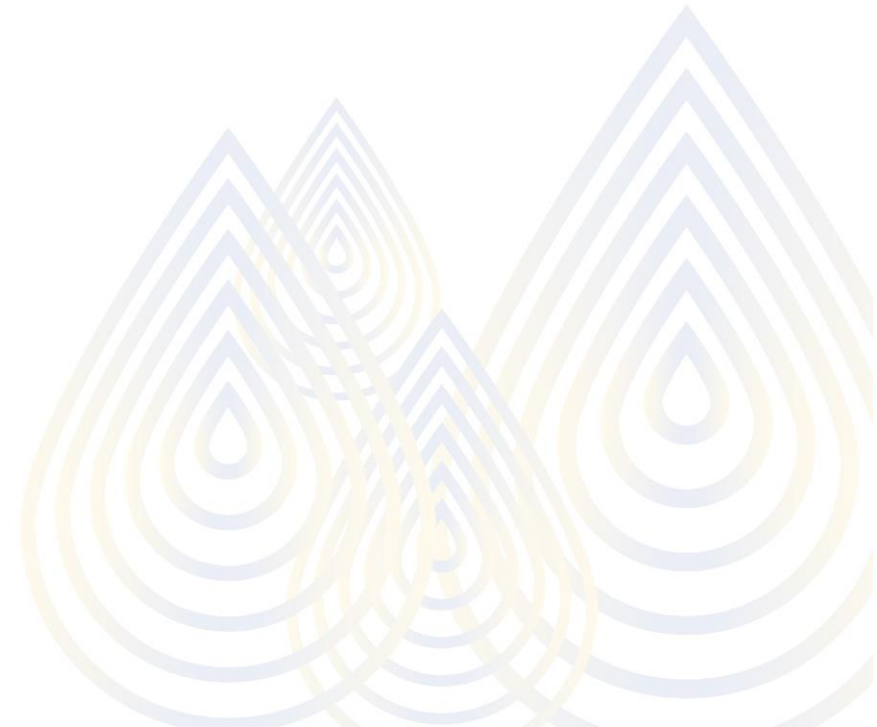
EGCI 466

SPARK



# Apache SPARK

- Supports a computing framework for large scale data
- Provides parallel and distributed processing
- Provide scalability
- Provide fault tolerance on **comodity hardware**

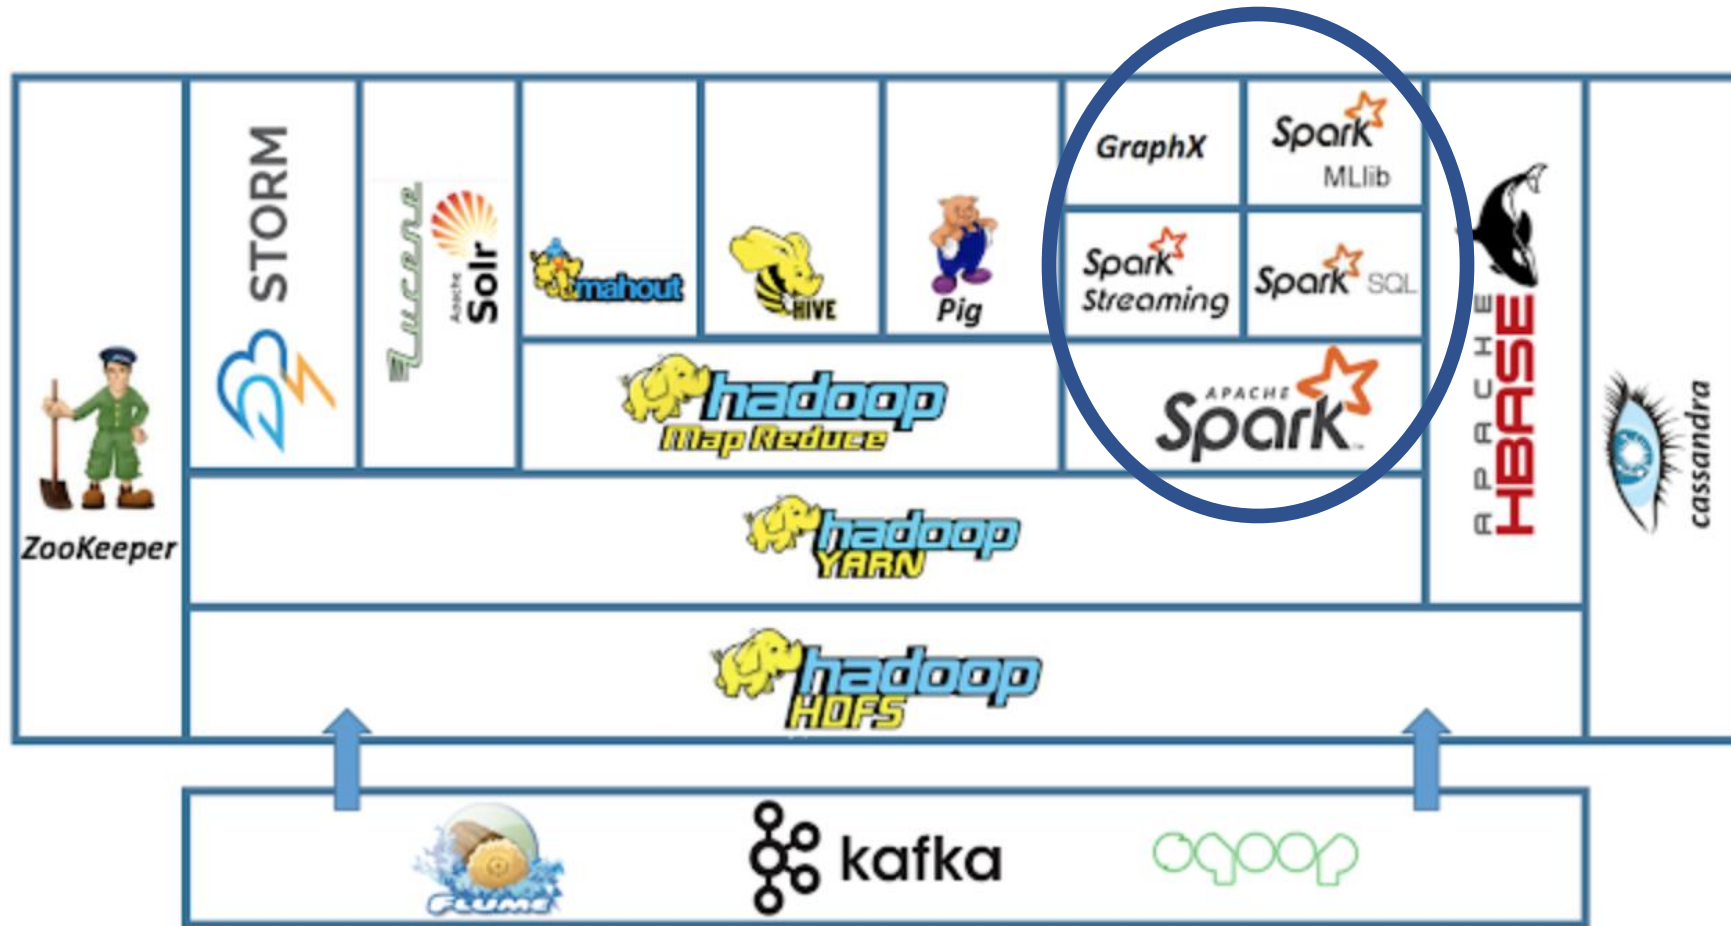


# Apache SPARK

- Is open sourced in-memory application framework
- Supports a computing framework for large-scale data processing and analysis.
- Provides parallel distributed data processing capabilities, scalability, and fault-tolerance on commodity hardware.
- Enable programming— [Scala](#), python and Java APIs



# Hadoop Layer



# Spark vs Hadoop Mapreduce

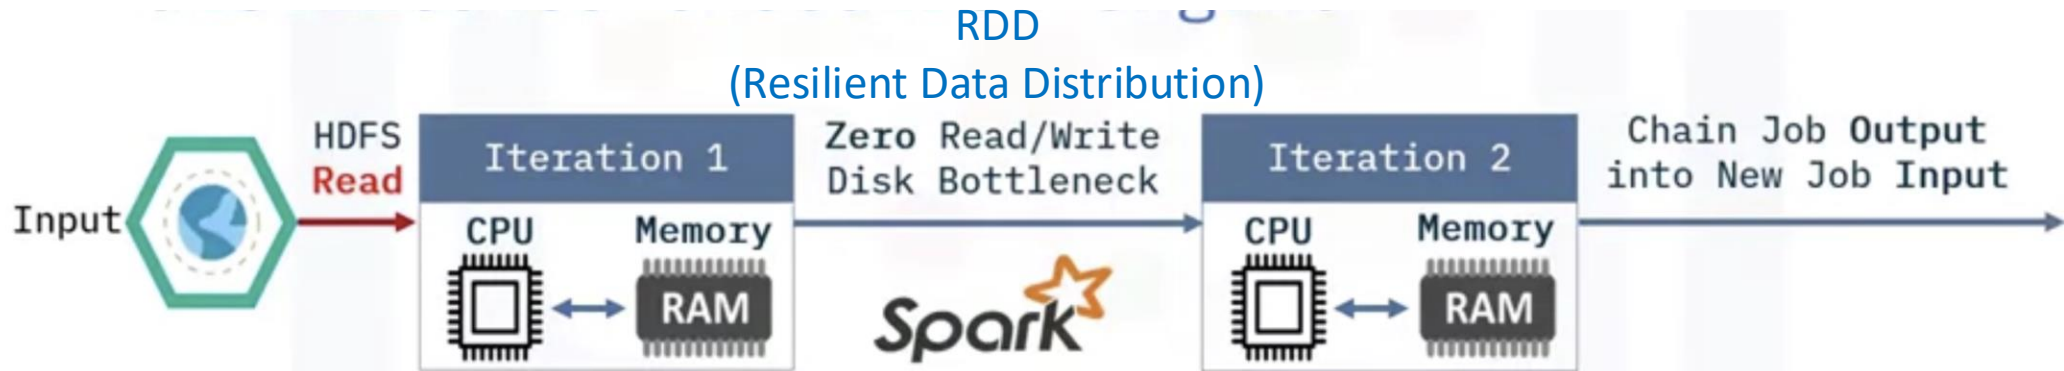
MapReduce involves a lot of I/O (**slow**)



- Limit only to map and reduce
- Difficult for more complicated data
- Native only for Java
- No interactive shell support (interface)
- No support for streaming

# Spark vs Hadoop Mapreduce

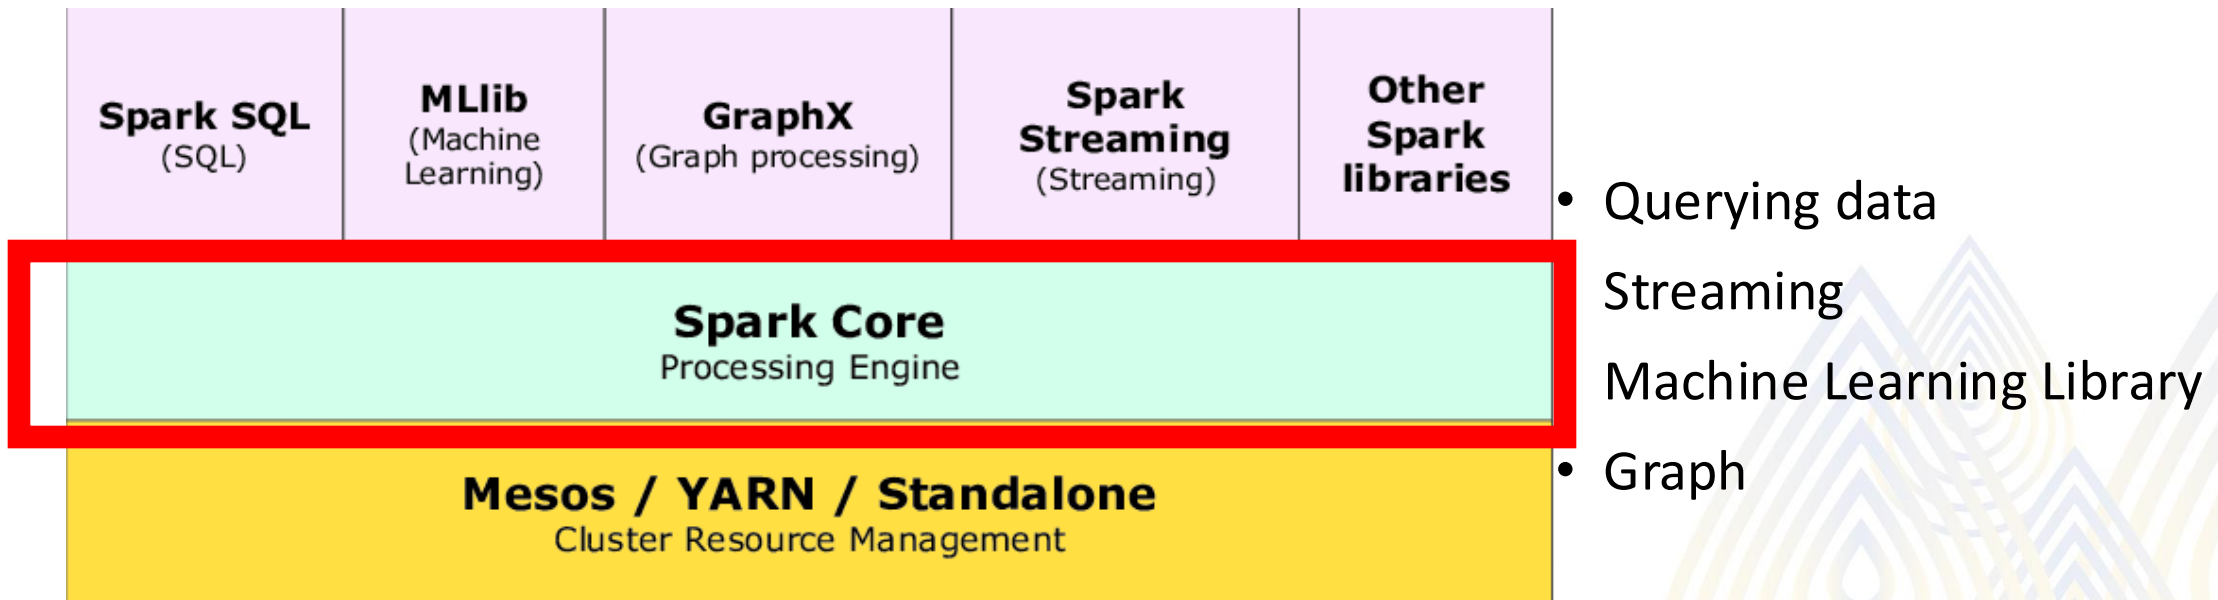
- In memory process
  - factor of 10-100 for some operations



- Provide lots of modules for distributed system
- Simple APIs
- Interactive shell
- Support multiple workloads: batch & streaming



# Spark Stack




Use RDD for programming abstraction  
(Resilient Data Distribution)

Carry data across many computing nodes in parallel, and transform it.

# Functional Programming

- FP: mathematical programming style
- Use **expressions** instead of **statements**
- First implementation LISP (LISt programming Language)
- Most recent: **Scala** ➔ **Spark**
- Others: Java, Python, R

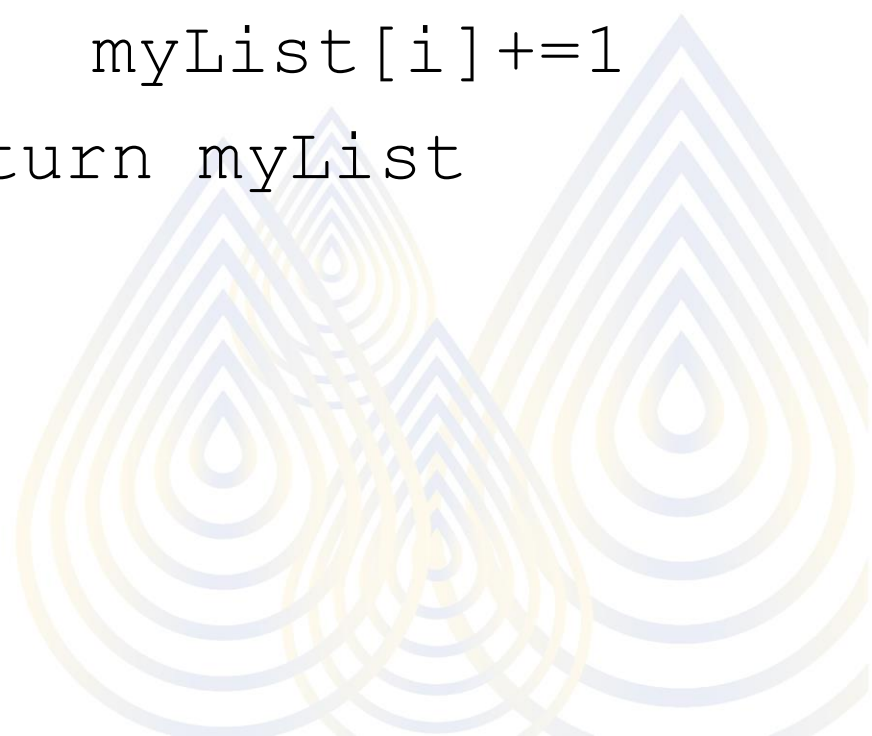




# Traditionally

```
int inc(int x[],n) {  
    int i;  
    for(i=0;i<n;i++)  
        x[i]++;  
}
```

```
def inc(myList) {  
    N=size(myList)  
    for i in range(N):  
        myList[i]+=1  
    return myList  
}
```



# Functional Programming Example

- **Function creator**

$$f(x) = x + 1$$


**apply(f(x),**

**)=**

2
3
4
5
6

3
4
5
6
7

# Parallelization



# lambda

Lambda: function or operation with simple operations

```
val add= (x:Int, y:Int)  
⇒x+y
```

```
println(add(1,2))
```

Scala

```
add=lambda x,y:  
    x+y
```

```
print(add(1,2))
```

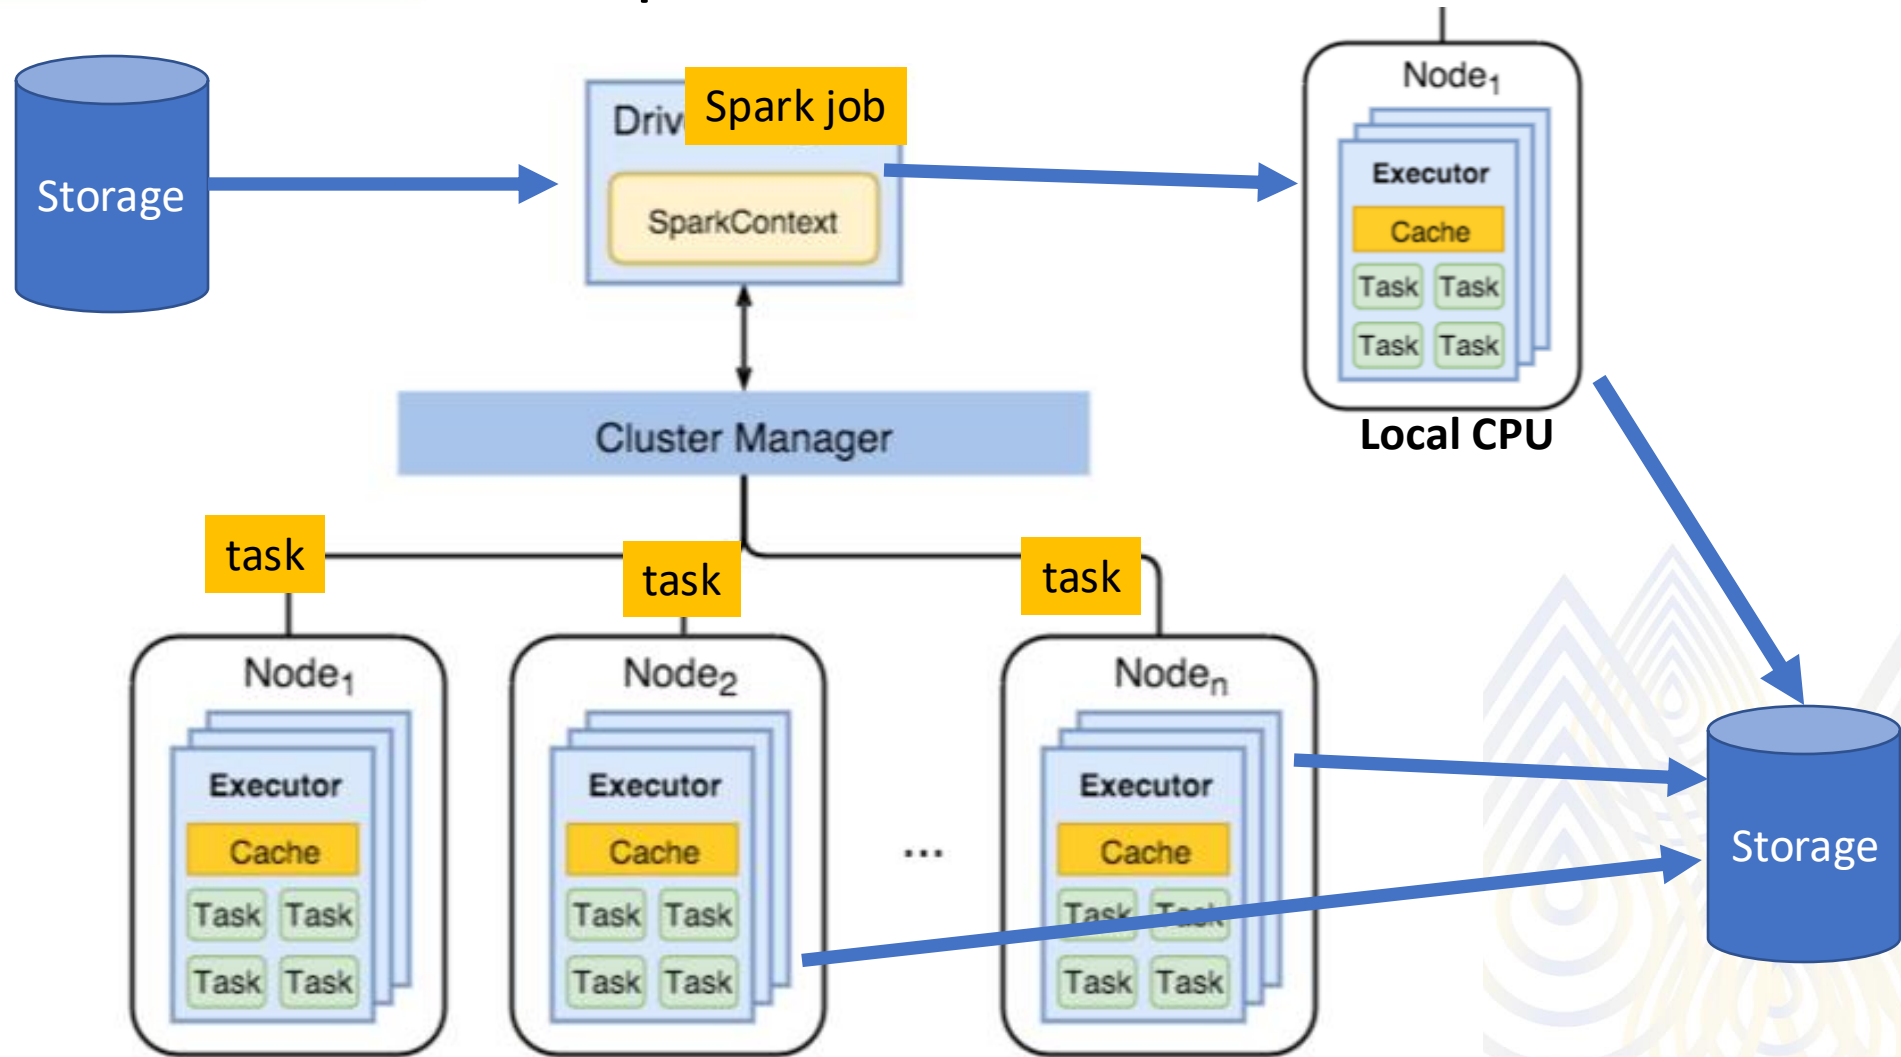
Python

# Spark Core

- Base engine, fault tolerance
- Perform scheduling, memory management
- Provide APIs to define RDD and other data type
- Use resilient distributed data sets (RDD) as the main programming abstraction
- Carry data across many the cluster nodes in parallel (distribute and parallelize)



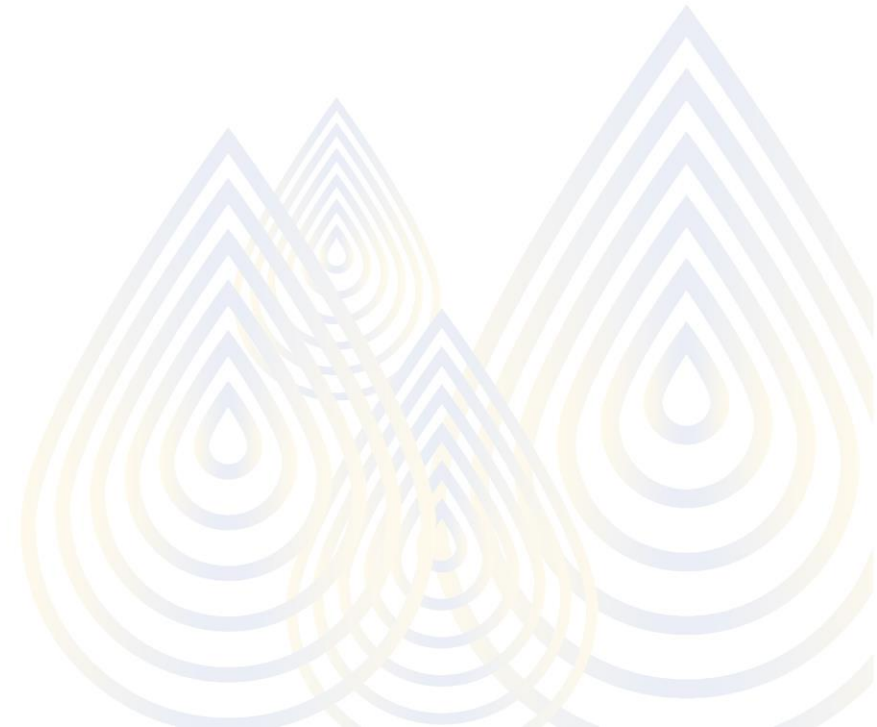
# Spark Core





# Resilient Distributed Dataset (RDD)

- Spark's primary data abstraction
- Collections of fault tolerant elements partitioned across the cluster's nodes
- Partitioned across the nodes of the cluster
- Capable of accepting parallel operations
- Immutable, cannot be changed once created




# Spark Applications

- Main components
  - Driver program → run user's main function + control worker node
  - worker nodes → Executor program → Running JVM driven by the driver
- Spark supports three main interfaces for cluster management
  - Spark's standalone cluster manager
  - the Apache Mesos,
  - Hadoop YARN
- Spark operation can Create, transform and Action

# Processing RDD

- Create
  - From Many Type of data
- Transformations
  - Create a new version of RDD
  - Lazy evaluation
  - Not immediately executed
  - wait for an action to be performed
- Actions
  - Converted and saved in a persistent storage -- HDFS or local drive
  - Run time error usually happen here

## + Transformations / Actions



# Creating RDD

- Import from input files

- `line = sc.textfile("hdfs://filename")`

- Create RDD from a list

- `data = [1, 2, 3, 4, 5] //python example`  
`distData = sc.parallelize(data, 3)`

- `data = Array(1, 2, 3, 4, 5) //Scala example`  
`val distData = sc.parallelize(data, 3)`

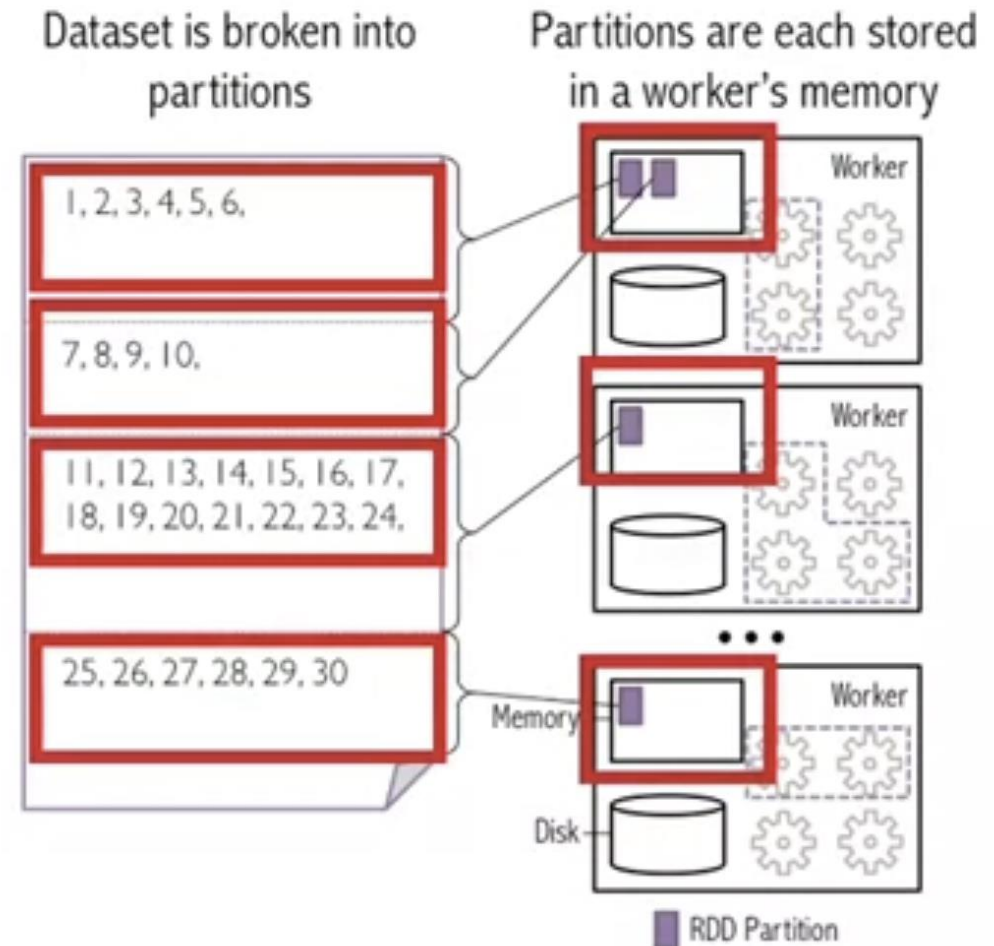
- Transform from existing RDD

parallelize into 3 partitions

[1,3], [2,5], [4]

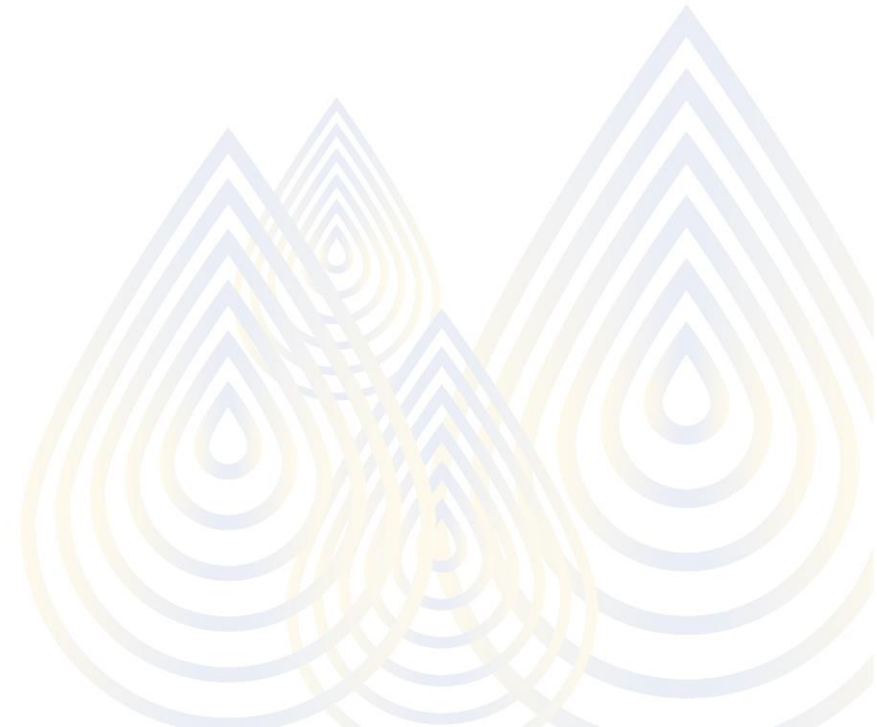
# RDD supported file types

- Text
- SequenceFiles
- Avro
- Parquet
- Hadoop input format



# RDD supported file formats

- Local
- Cassandra
- Hbase
- HDFS
- Amazon S3
- etc





# Resilience in RDD

- The persisting or caching of a data set in memory across operations.
- The cache is fault tolerant and always recoverable
- RDDs are immutable and Hadoop provide fault tolerant.
- Each node stores the partitions and subsequent action
- Persisting or caching is used as a key tool for iterative algorithms and fast interactive use

# Spark Cheat Sheet

Method	Description	Usage
<code>appName ()</code>	A name for your job to display on the cluster web UI.	<pre>from pyspark.sql import SparkSession spark = SparkSession.builder.appName("MyApp") . getOrCreate()</pre>
<code>cache ()</code>	caches the specified RDD in the memory of your cluster's workers. Caching operation takes place only when a Spark <b><u>action</u></b>	<pre>df = spark.read.csv("customer.csv") df.cache()</pre>
<code>pip install()</code>	Find the latest version and install	<code>pip install pyspark</code>
<code>sc.parallelize ()</code>	Creates a parallelized collection.	<code>rdd = sc.parallelize([1, 2, 3, 4, 5])</code>

# Advantages

- **Increases Manageability**

- small operations
- group operations

- **Reduction of Complexity**

- **Time** is saved. The action gets triggered only when data is required
- **Space** is used only when necessary, which saves space.

- **Saves Computational Power and increases Speed:**

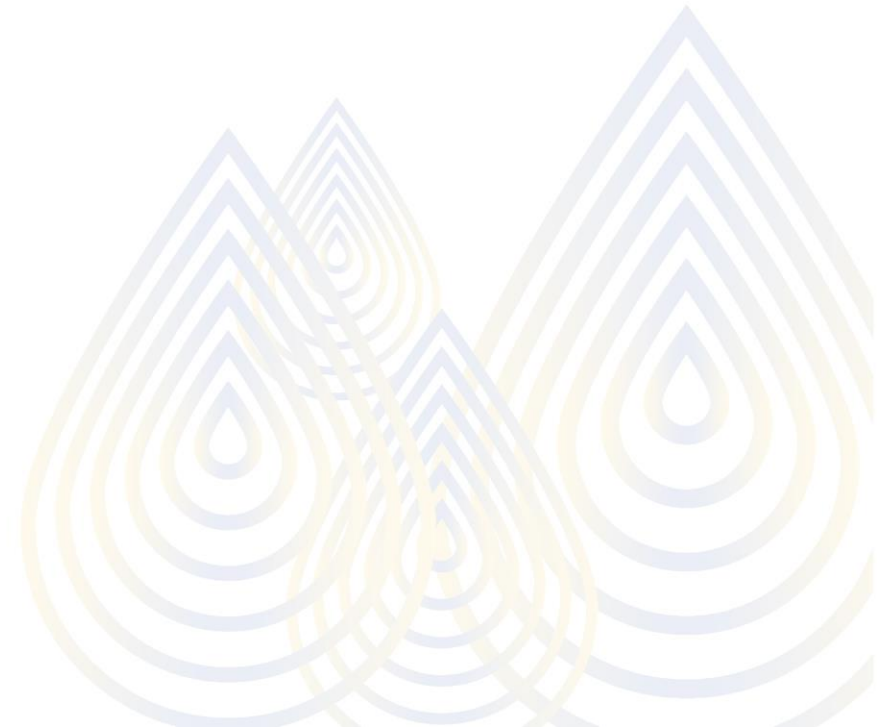
- Spark functions get triggered through the driver and run on a cluster.
- Lazy evaluation only triggers a computation when necessary
- Saves trip between Driver and Cluster – speed the execution

## Optimization:

- The number of queries being run is very low at a given time -- optimized

# Spark Core: Transformations

- Transformations do not execute immediately (lazy)
- A transformation is a single pipeline for execution later (Action)
- Keep track of all process
- The process is linear




# Transformation types

- **Narrow Transformations**

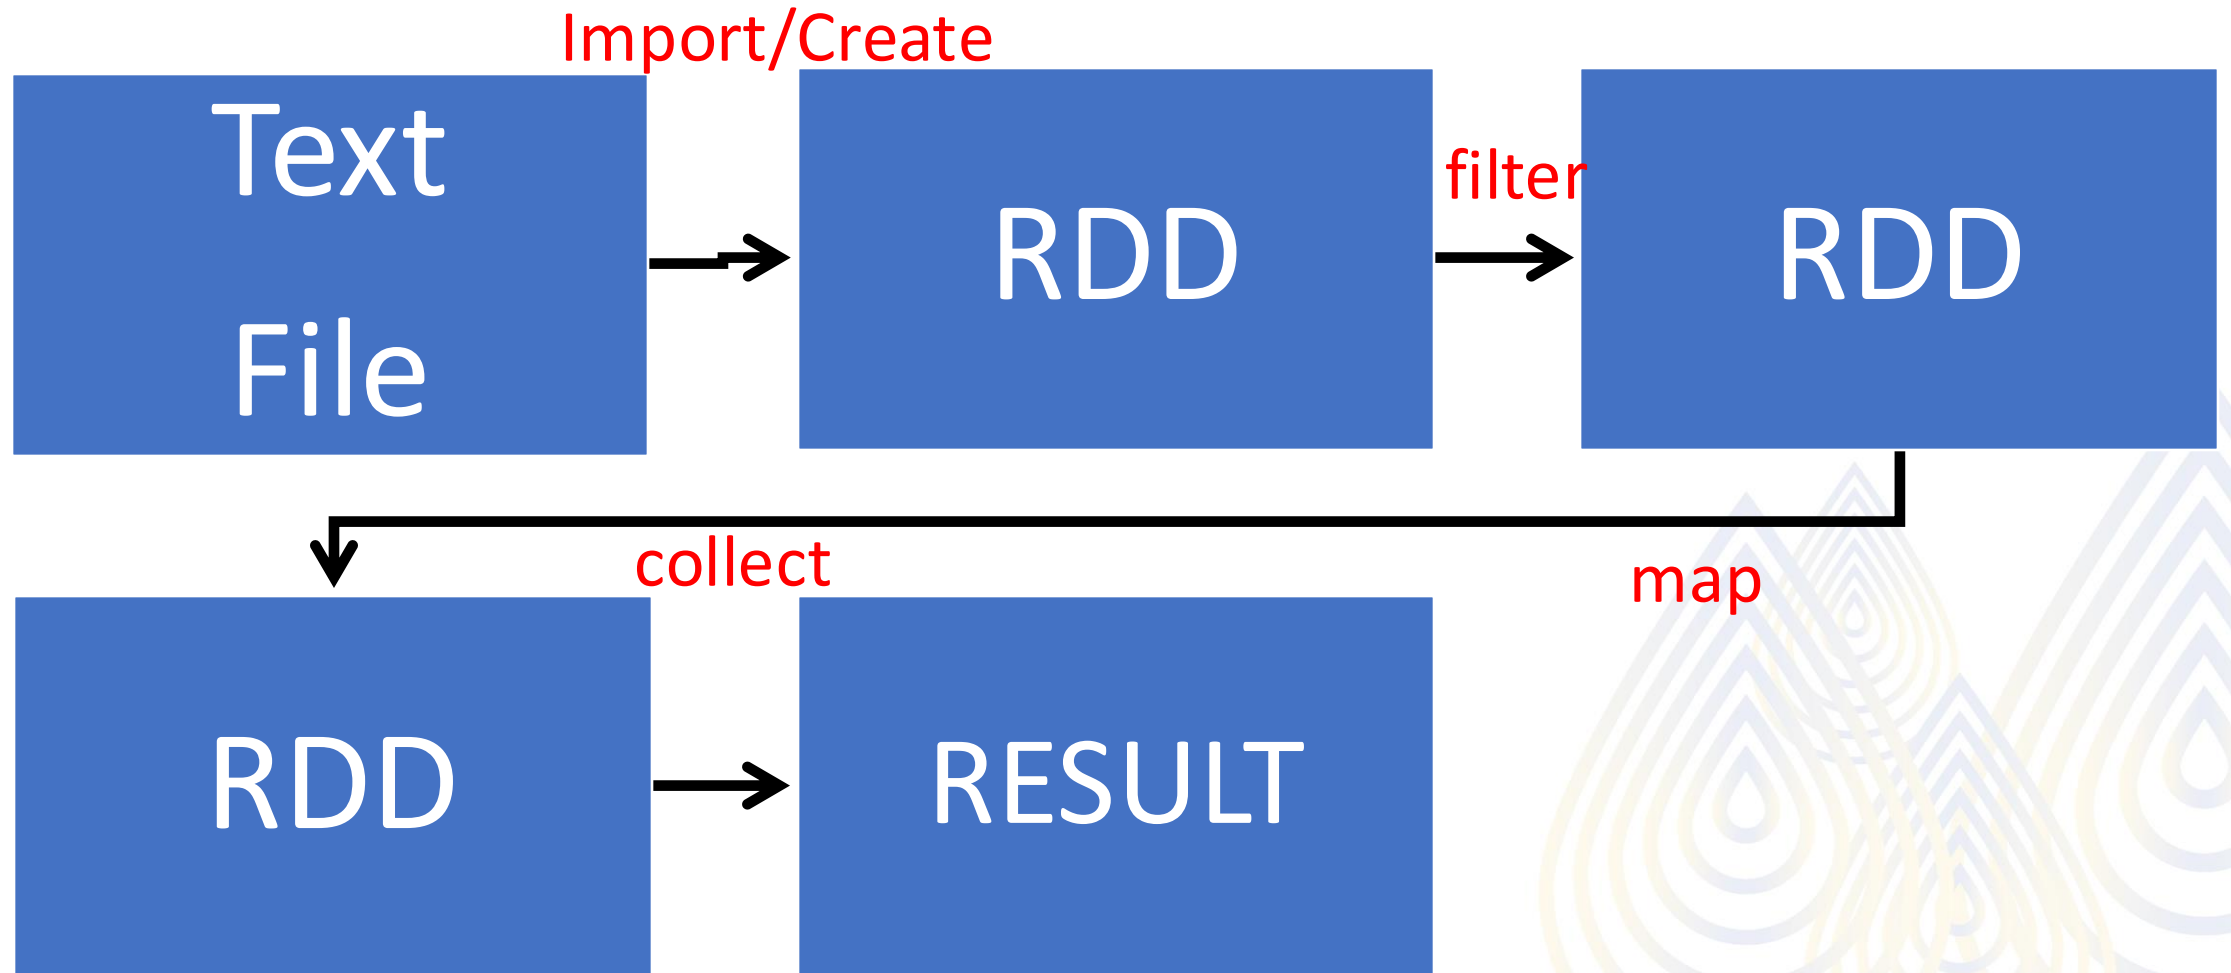
- Do not require data shuffling across partitions.
- **map(), flatmap(), filter() , coalesce()**

- **Wide Transformations**

- **Require** data shuffling through partitions
- **GroupByKey, ReduceByKey**

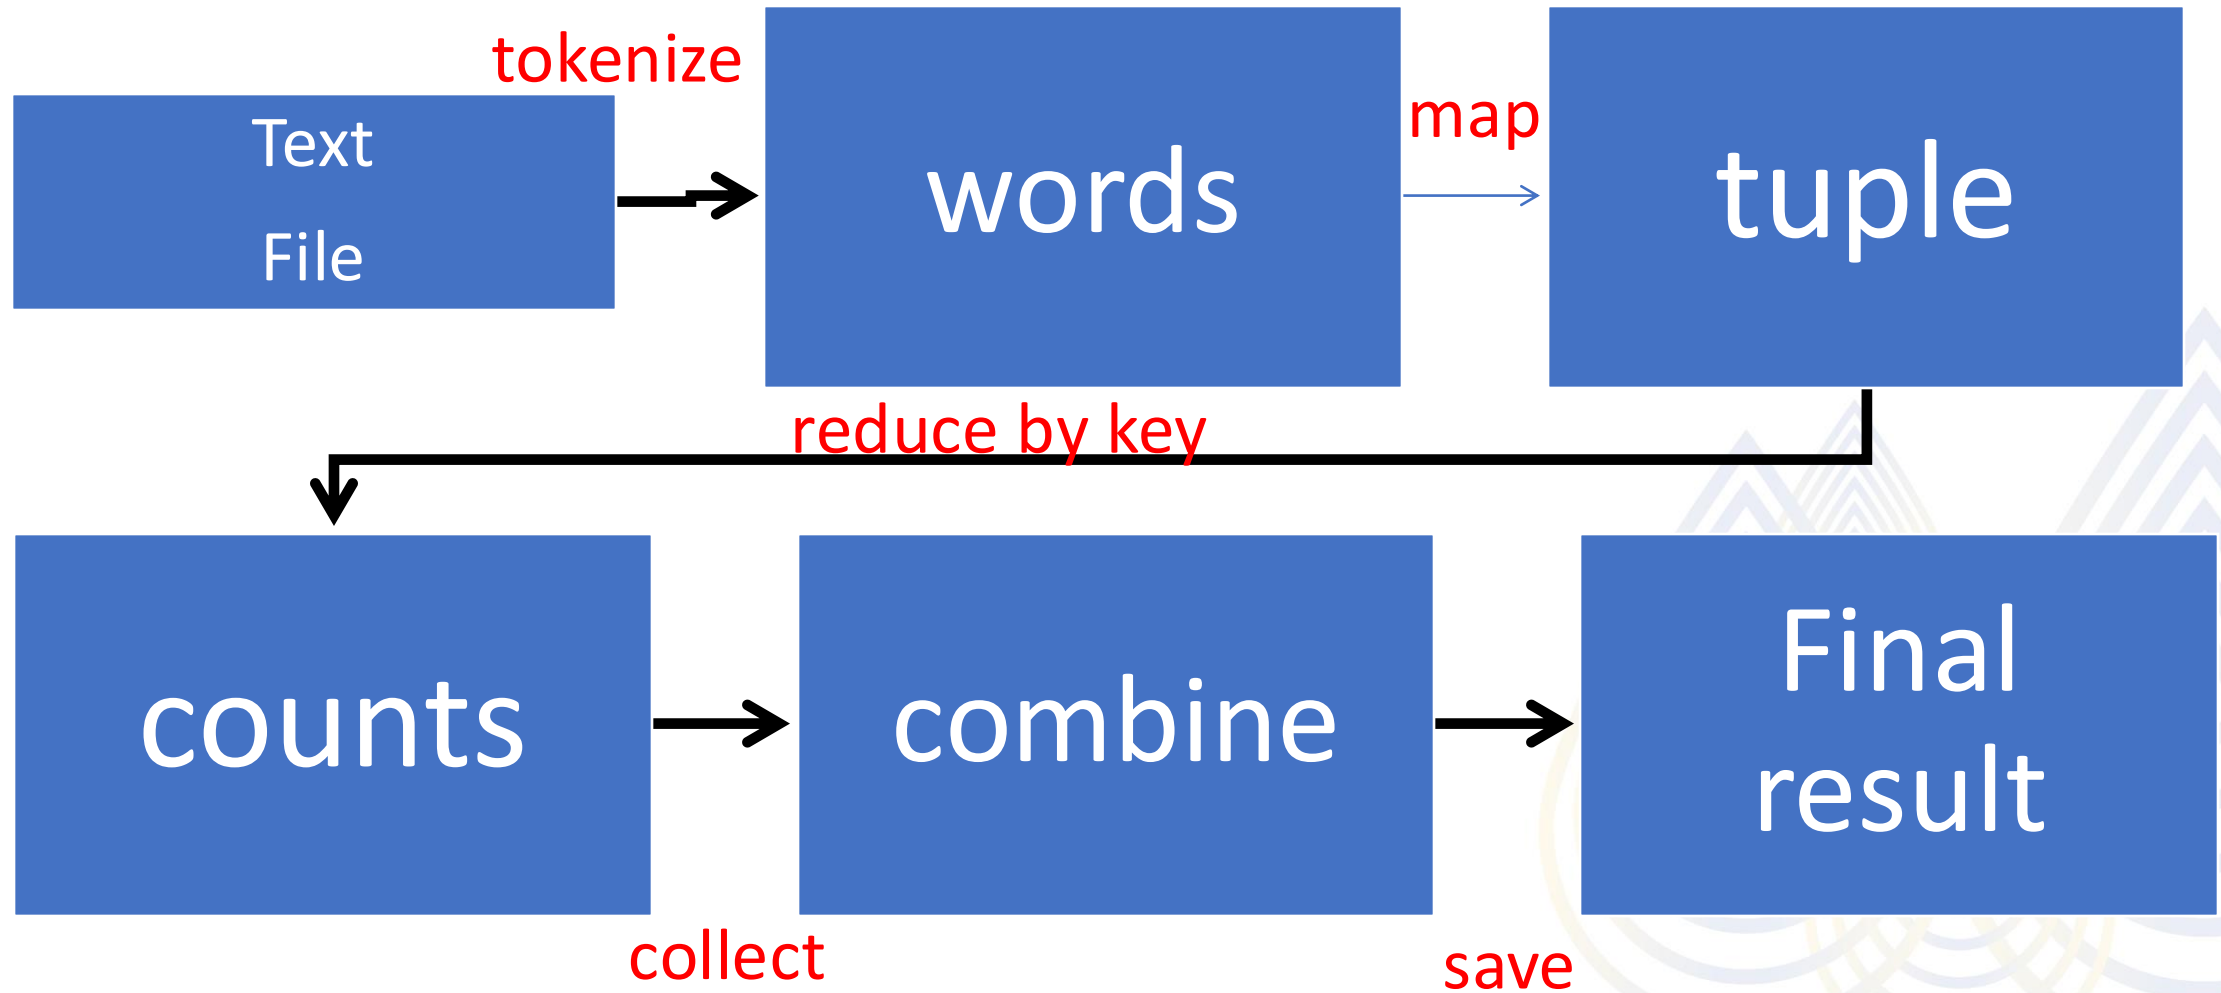


# Processing RDD : Example





# Processing RDD : word count




# Word count Example

- Import from input files
  - **line** = `sc.textfile("hdfs:/filename")`
- Transformation (Tokenized)
  - `words = lines.flatMap(lambda line : line.split(" "))`
- Transformation (count)
  - `tuples = words.map(lambda word : (word, 1))`
- Transformation (Reduce)
  - `counts = tuples.reduceByKey(lambda a, b: (a + b))`
- Save
  - `counts.coalesce(1).saveAsTextFile('hdfs:/user/cloudera/wordcount/outDir')`

# map transformation

- 1:1 operation
- apply to each element of RDD

```
def lower(line):  
    return line.lower()  
  
lower_text_RDD=text_RDD.map(lower)
```

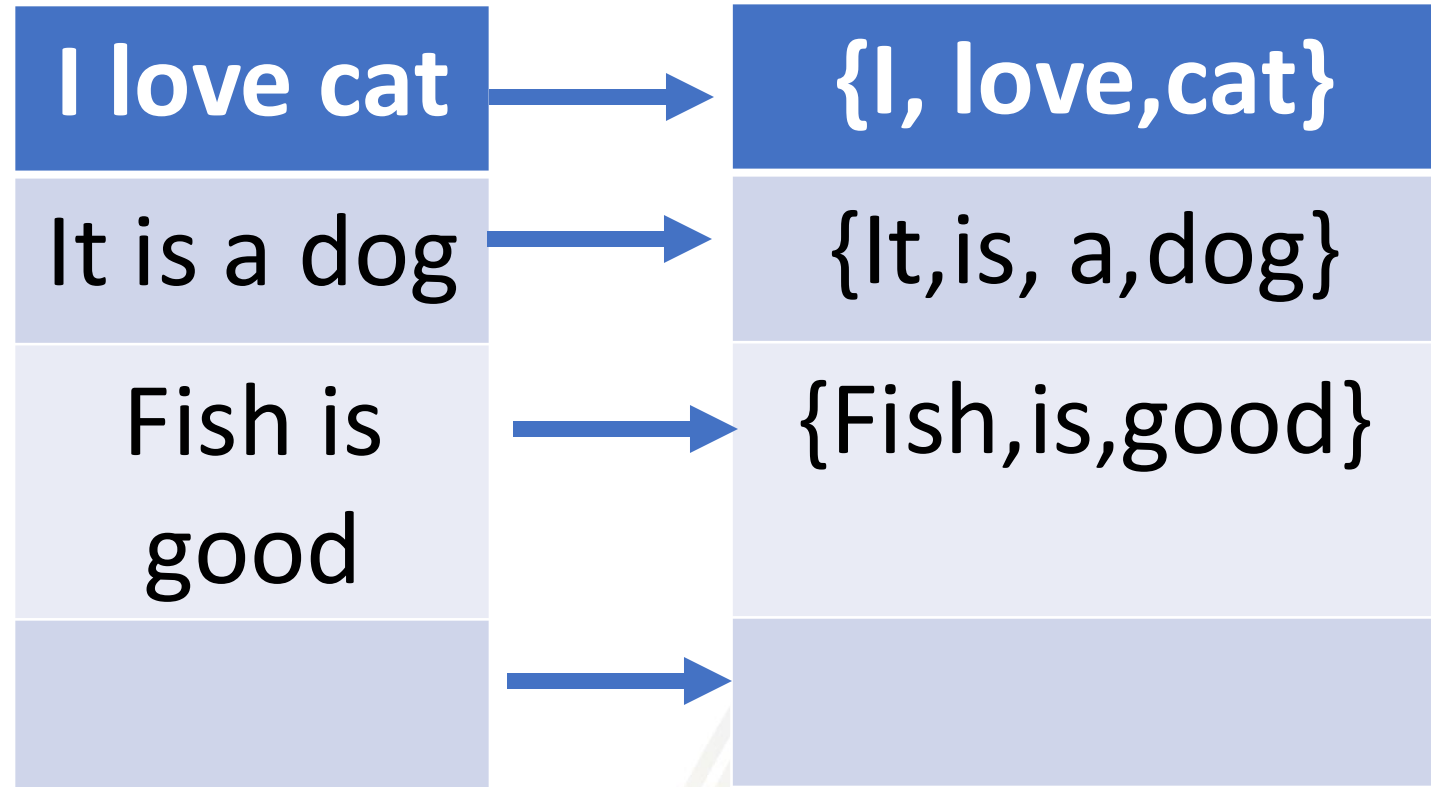


# flatMap transformation

- map then flatten output

```
def split_words(line):  
    return line.split()
```

```
words_RDD=text_RDD.flatMap(split_words)  
words_RDD.collect()
```



# Transformation:filter

## Example functions: filter()

using python:


```
hadoopbooksRDD=booksRDD.filter(lambda x: "hadoop" in x)
```

```
kafkabookRDD=booksRDD.filter(lambda x: "kafka" in x)
```

```
def starts_with_a(word)
```


```
    return word.lower().startswith("a")
```

```
words_RDD.filter(strats_with_a).collect
```



# Transformation:coalesce

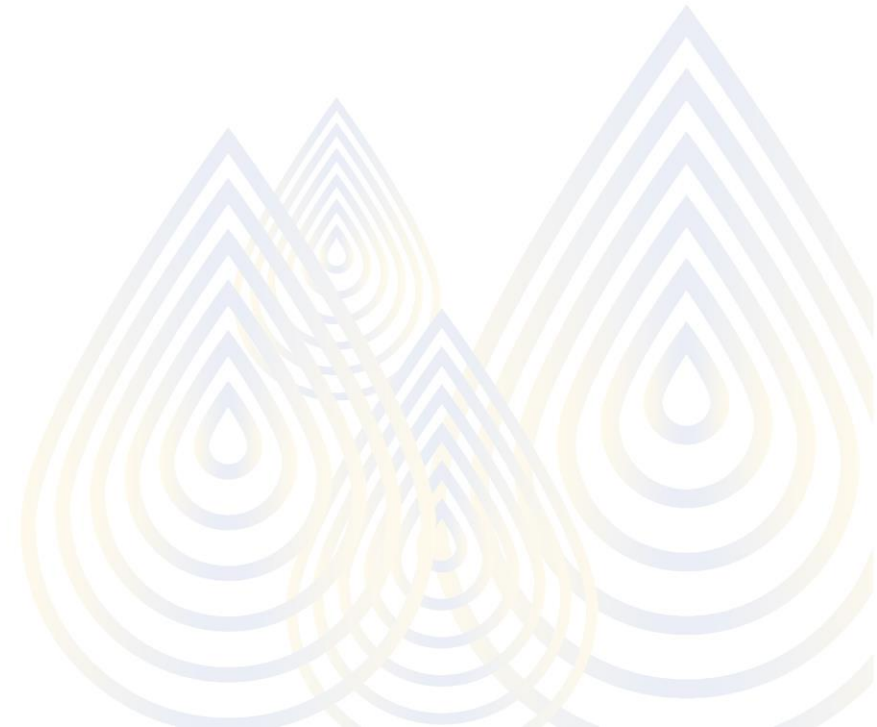
- Reduce the number of partition






# Wide transformation

- Require shuffle
- Group by partition
- `reduceByKey`, `groupByKey`




# groupByKey

(Key,value) → (key , [list of all value])




# groupByKey + reduce



# reduceByKey

- Reduce = sum for this example

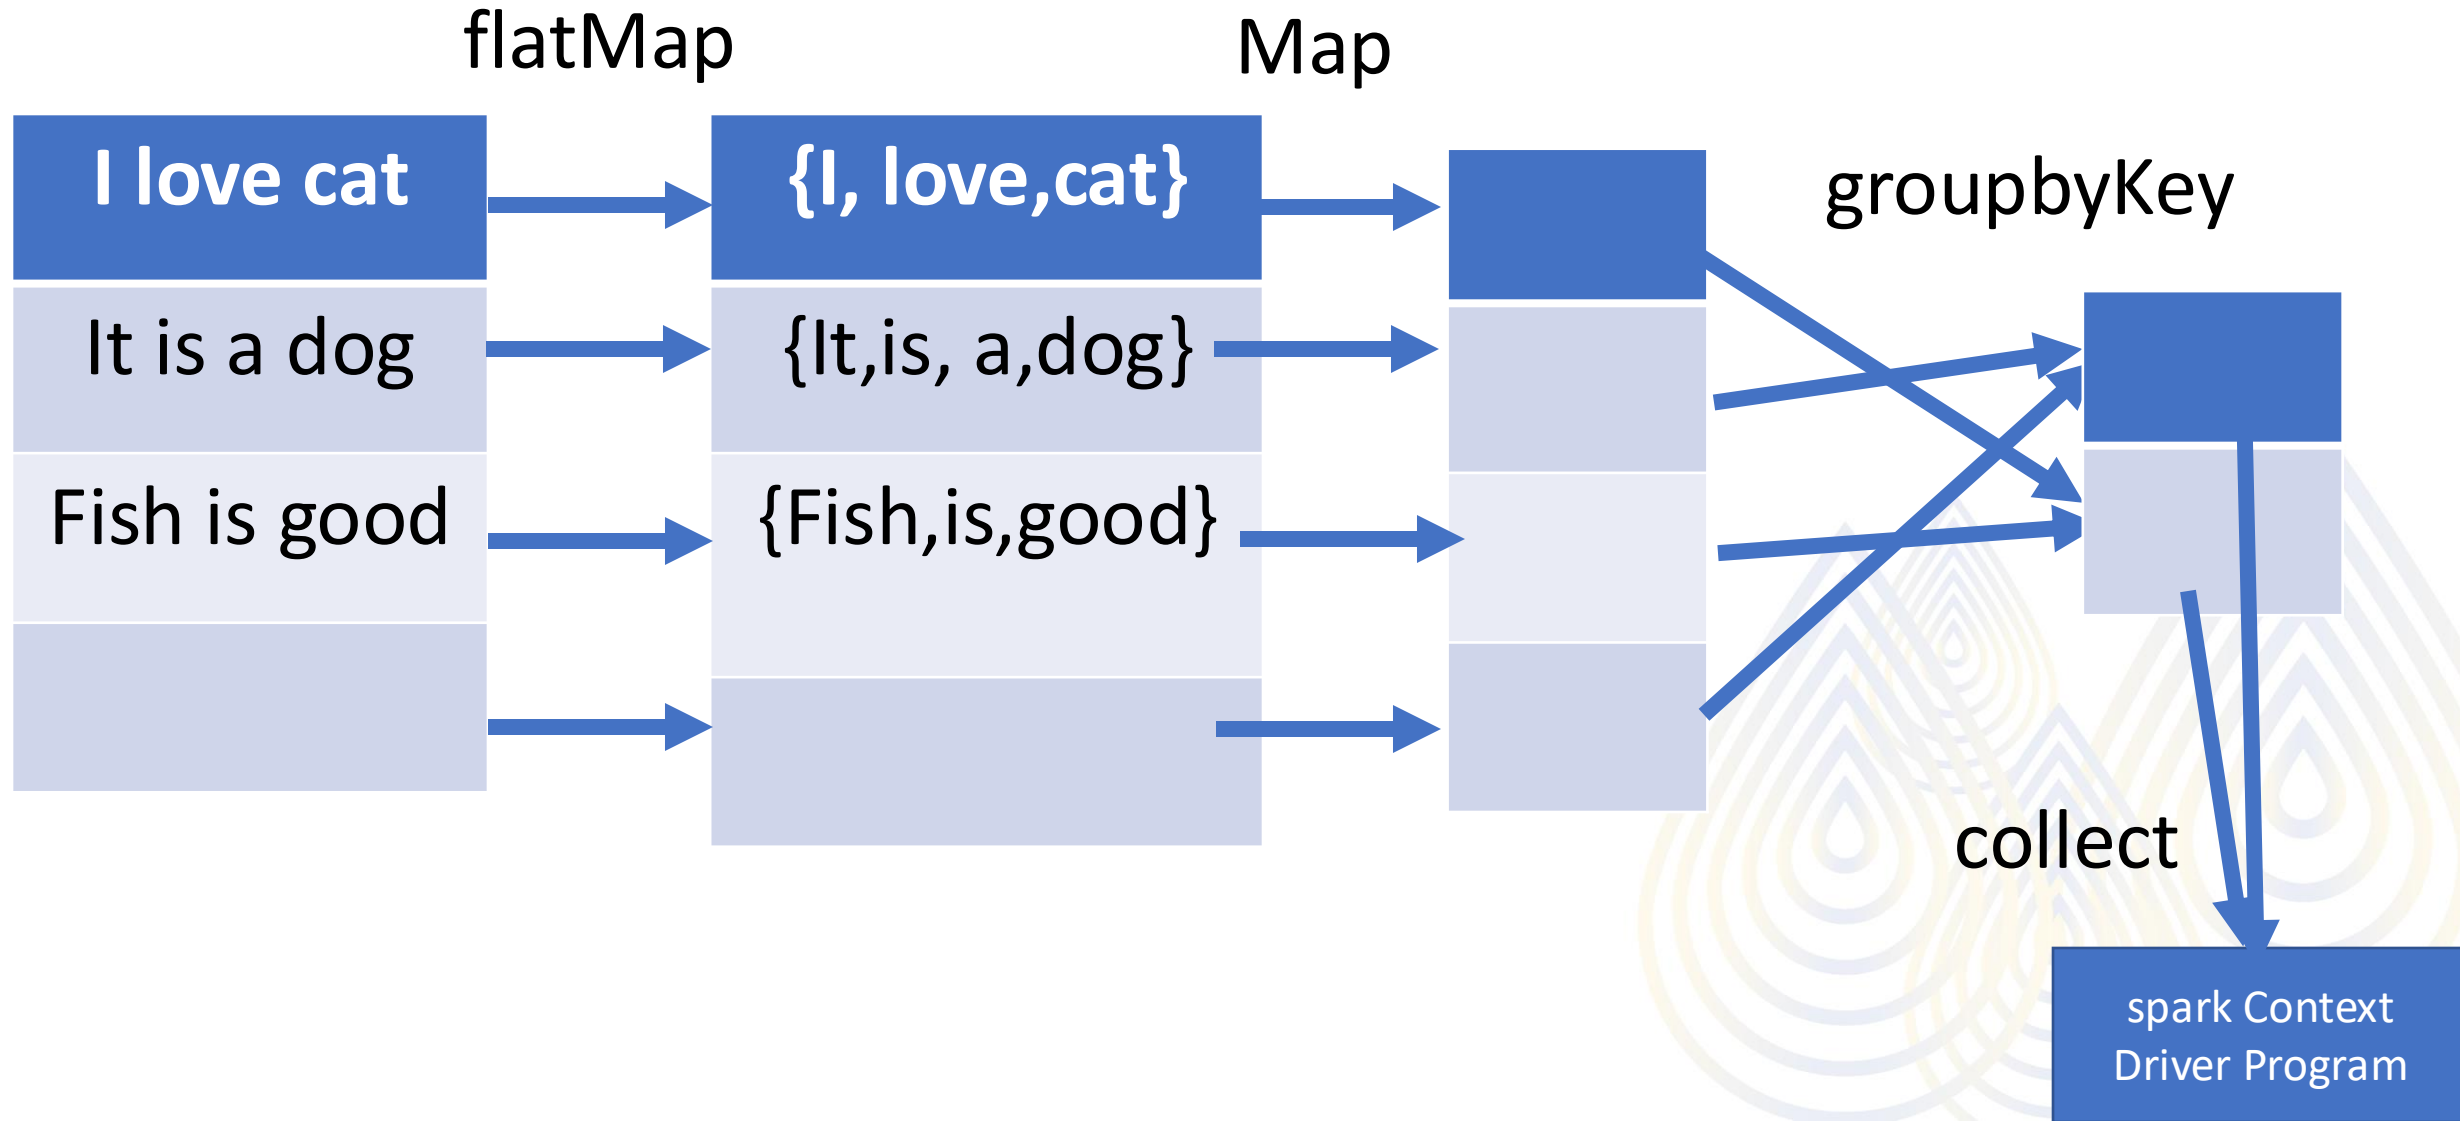


# Spark core: Actions

- Last step in a Spark pipeline
- Actions operations trigger the evaluation of transformation pipeline
- Return the final result to the driver program or save the results to a persistent storage.
- **Collect** is the most common action



# Action: word count



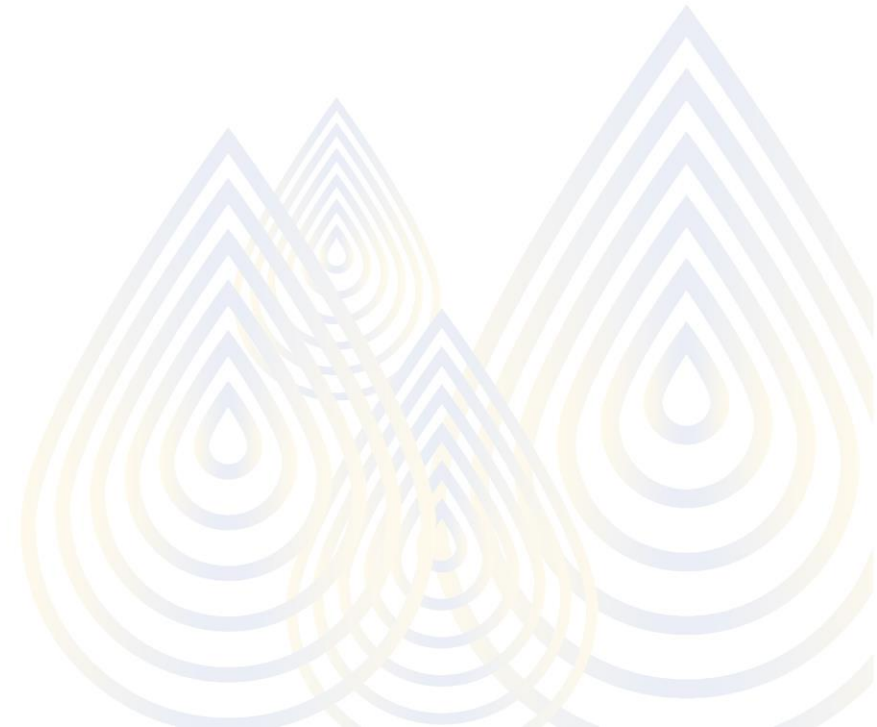
# The most common Actions

Action	Usage
<code>Collect ()</code>	Copy all element to driver
<code>Take (n)</code>	Copy first <b><i>n</i></b> element
<code>Reduce (func)</code>	Aggregate elements with <b><i>func</i></b> (take 2 elements, return 1)
<code>Count ()</code>	Count number of data
<code>saveAsTextFile (Filename)</code>	Save to a local or HDFS file



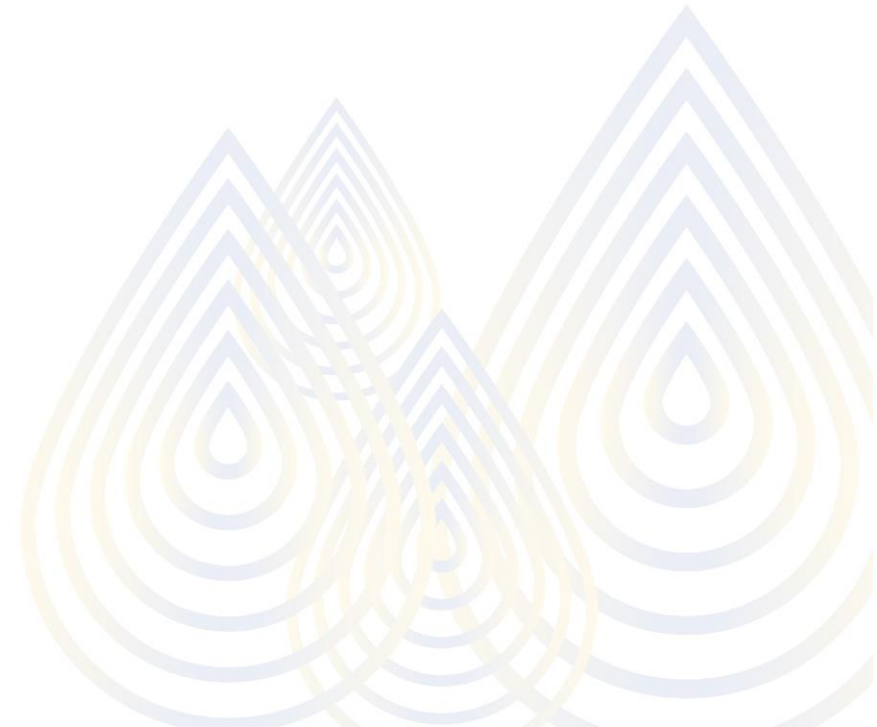
## Example 2

- Letter Count



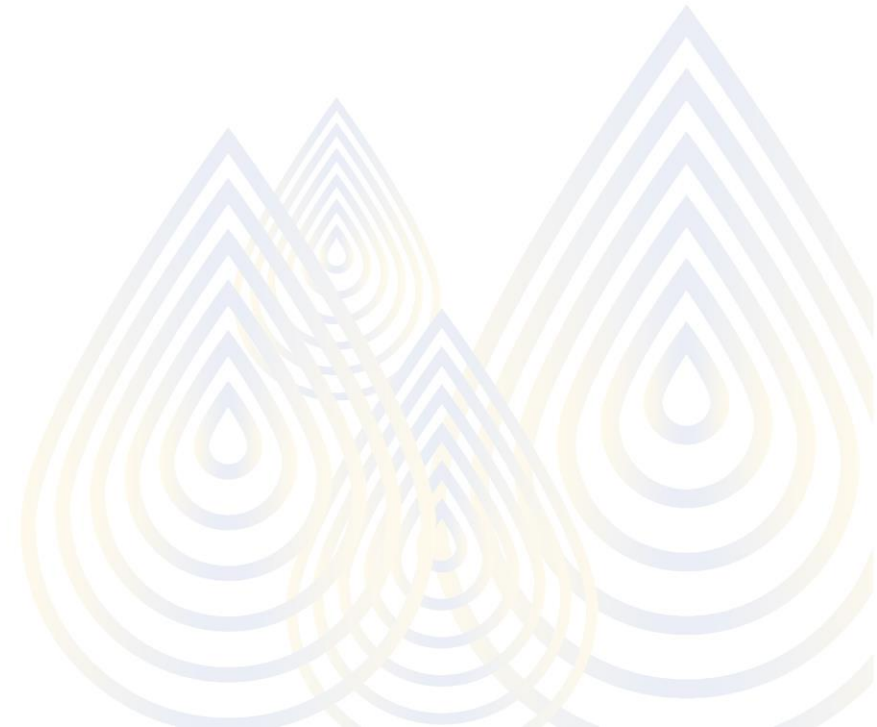
# More references

<https://spark.apache.org/docs/latest/rdd-programming-guide.html#transformations>



# Experiment on cloud

- Use pySpark as submit job
- Use Notebook to run the job



Reference: <https://spark.apache.org/docs/latest/rdd-programming-guide.html>